



M2 Lab — Ten Content Cards

10 card(s) — 2026-09-12 — run 92e71e7a0543

Table of contents

1. M2 Lab — Ten Content Cards

- a. How to read a card
- b. Cards in this book
- c. Template and production rules (PRODUCTION.md)
- d. Asset legend

2. C-2026-09-12-01 — Nineteen of our top twenty were there on a trick

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

3. C-2026-09-12-02 — The wrong name in twenty-seven files

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

4. C-2026-09-12-03 — We built it, it worked, and we switched it off

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

5. C-2026-09-12-04 — Cost per run, not price per month

- a. Strategy
- b. Hooks

- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

6. C-2026-09-12-05 — The approval box we designed and never ticked

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

7. C-2026-09-12-06 — The pause is the feature

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

8. C-2026-09-12-07 — The same number in two places

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

9. C-2026-09-12-08 — The paragraph that beat the agents

- a. Strategy

- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

10. C-2026-09-12-09 — Supported is not the same as connected

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

11. C-2026-09-12-10 — Twenty-three plans, nothing made

- a. Strategy
- b. Hooks
- c. Reference reels
- d. Script
- e. Storyboard
- f. Editing
- g. Claims
- h. Traceability
- i. Review
- j. Production state

M2 Lab — Ten Content Cards

Generated 2026-09-12 (commit 92e71e7a0543). 10 card(s) in this book.

How to read a card

Each chapter after this one is one shooting card, in the order of the table below, and is a production-ready creative brief, not a summary. Strategy states why this exact reel over any other, including the four-question filter it had to pass. Hooks lists the three candidates considered and marks the one actually carried into the script. Reference reels names the outside reels whose shape was borrowed and exactly what was taken from each — the transcript quote and frame that earned it a place, and how it was transformed rather than copied. Script is the full spoken text plus per-section timing. Storyboard is the frame-by-frame shot list, the rendered placeholder frames in scene order, and the contact sheet. Editing, Claims, Traceability and Review record how the card is cut, what it may claim on screen, why it exists, and what the review pass changed. Production state is the current shoot/render status.

Cards in this book

ID	TITLE	FORMAT	TOTAL S	WOR DS	REF S	HYPOTHE SIS	STATU S
C-2026-09-12-01	Nineteen of our top twenty were there on a trick	M2 Builds	68.4	171	5	H-11	REVIEW ED
C-2026-09-12-02	The wrong name in twenty-seven files	M2 Builds	65.2	163	4	H-12	REVIEW ED
C-2026-09-12-03	We built it, it worked, and we switched it off	M2 Builds	69.6	167	4	H-18	REVIEW ED
C-2026-09-12-04	Cost per run, not price per month	M2 Radar	68.0	170	4	H-01	REVIEW ED
C-2026-09-12-05	The approval box we designed and never ticked	M2 Builds	68.5	178	4	H-13	REVIEW ED
C-2026-09-12-06	The pause is the feature	M2 Radar	69.2	173	4	H-04	REVIEW ED
C-2026-09-12-07	The same number in two places	M2 Teardown	69.2	173	4	H-20	REVIEW ED
C-2026-09-12-08	The paragraph that beat the agents	M2 Radar	69.6	174	4	H-09	REVIEW ED
C-2026-09-12-09	Supported is not the same as connected	M2 Radar	68.0	170	4	H-05	REVIEW ED
C-2026-09-12-10	Twenty-three plans, nothing made	M2 Teardown	66.4	166	4	H-19	REVIEW ED

Template and production rules (PRODUCTION.md)

- Fixed: one location and light, one banner position and typeface held for the whole reel, one caption style, the screen shown full-frame, 50–70 seconds, one closing line.
- Variable: who is on camera and how many — one, both, or one after the other.
- The only cut allowed is a change of speaker; everything else is shot in one take.
- Target length is 50–70s (winners' median 55s vs 47s for the rest); under 20s or over 90s tracks with the losing side of the corpus.
- A card is not shoot-ready until every reference and claim is checked and the storyboard is reviewed — see each card's Review and Production state sections.

Asset legend

`asset_status: "template"` on a storyboard frame is a rendered placeholder from `engine.storyboard_render` (dashed box, `[SHOT: ...]` / `[SCREEN: ...]` label — never a stock photo, never an AI-generated image) — replace it with the owner's real footage per `cards/README.md` before this card is shot for real. `missing` = nothing rendered yet; `ready` = real footage bound in; `generated` = an AI-generated visual (not used anywhere in this codebase).

C-2026-09-12-01 — Nineteen of our top twenty were there on a trick

Strategy

Format: M2 Builds

Concept. We re-sorted our own weekly priority list after making the comparison fair, and nineteen of the top twenty items turned out to be there on a habit that has nothing to do with quality.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company who is responsible for one repeated process and has tried ChatGPT but has no repeatable way of using it at work. Specifically anyone who works from a queue somebody else sorted: enquiries, tickets, applications, invoices.

Objective. Hand the viewer one check they can run on their own priority list this week, using our own ranking being wrong as the demonstration.

Positioning. builder

Pain. quality_trust

Promise. One check that tells you whether your priority list is measuring importance or measuring a habit.

Rationale. The references supply the shape only: a result-first opening on the artefact (DZLDk9ySi-7), one short mid-roll re-hook before the evidence (DNA7-d3o5sQ), a fixed consequence-then-fix cadence (DbO4zvJR1k8), per-item contribution shown on screen rather than a summary (DX9kWTYzssE), and an ungated close that names a bounded test (DTCOU9DkRTG). Nothing in the 266-reel analysed corpus shows an author's own system producing a wrong answer, so the subject is ours alone.

FILT ANSWER
ER

Proce ss	Deciding what to look at first in a recurring queue of incoming work - which enquiry gets answered, which ticket gets picked up, which invoice gets chased - where the order is produced by a rule or a score rather than by reading every item.
Fricti on	The sort order comes out of a formula nobody re-reads, so over time the queue quietly optimises for whatever the formula happens to count rather than for what actually matters, and the drift is invisible because the list still looks sorted.
AI boun dary	The machine ranks and has to show, per item, which field put it where. A named human owns the definition of 'important', re-reads it against a handful of items by hand, and is the only one who may change the weighting.
Next action	Take the five items your system put at the top this week and find what they have in common that has nothing to do with importance.

Hooks

1. **[result_first]** (statement) — **USED IN SCRIPT** This is the list our system put at the top. Nineteen of the top twenty are there on a trick.
1. **[number_stat]** (statement) Two items our own ranking buried in the sixties belonged in the top ten.
1. **[question]** (question) What is your priority list actually measuring?

Reference reels

REEL	CREATOR	FULL NAME	P	CREATOR	VIEW	SHARE	SAVE	REASON
DZLDk9ySi-7(https://www.instagram.com/reel/DZLDk9ySi-7/)	aiforb	hooke	17,093	3,421	51.4x	4.2%	4.5%	Puts the state of the world first and the mechanism second, so the pain lands as reward rather than as complaint. result_first hook type; result_first and demo_first together are the highest-performing hook family measured (median robust_z 2.17 against 1.21). Its view_lift of 51.35 - that is 52.4x the author's own median, since view_lift is the ratio minus one - rests on a 3 421-view median, so the multiplier is not a finding and the ordering is.
DNA7-d3o5sQ(https://www.instagram.com/reel/DNA7-d3o5sQ/)	edhill	rhythm	45,491	8,010	55.4x	2.3%	3.7%	Carries the corpus's cleanest mid-roll re-hook: a 1.8-second beat at 32.0-33.8 s placed exactly where the mechanism stops being new and immediately before the proof. It is the only ungated reel in the pool (cta_type none), so its 36.5 saves per 1 000 carry no keyword inflation.
DbO4zv-JR1k8(https://www.instagram.com/reel/DbO4zvJR1k8/)	shrug	pailey	20,121	14,481	13.5x	5.3%	8.7%	The highest save rate in the whole corpus (87.4 per 1 000) bought with a repeated shape: the consequence of not doing it, then the exact clause that closes it, four times without variation, over one unbroken 71-second close-up. It is how a reel about a defect stays usable rather than merely alarming.
DX9k-WTYzssE(https://www.instagram.com/reel/DX9k-WTYzssE/)	vansh	screenplay	99,558	8,982	10.1x	4.1%	3.6%	Most dashboard reels stop at 'here are my agents'; this one shows contribution per item, which is exactly the screen a re-sorted priority list needs. The proof beat runs 36.5-56.9 s over a DATA_VISUAL state and the reel is not comment-gated (cta_type dm), so its 41.3 shares per 1 000 carry no keyword inflation.
DTCOU9Dk-RTG(https://www.instagram.com/reel/DTCOU9DkRTG/)	buildi	cta	21,733	4,124	52.0x	3.0%	4.2%	The pool's best ungated close - it converts the ask into a bounded, falsifiable test instead of a keyword, and it is the only shape in the pool that invites a reply without gating an artefact, which is the position M2 is in by rule. 29.7 shares and 42.0 saves per 1 000 with cta_type question_to_audience.

DZLDk9ySi-7 (aiforbusinesses247, hook)

- Transcript (DZLDk9ySi-7-B0): “Imagine, a client in the morning is already interested to take your service because your AI Automation Agency system is automatic”
- Frame (DZLDk9ySi-7-S00): opens with no presenter at all - a full-frame animated pipeline board whose counters climb while the voice talks over it
- Transformed: Borrowed the ordering only: state of the world first, mechanism second. Changed: our first frame is a real list of our own that was wrong, not an imagined good morning; we delete the second-person future-pacing entirely because setup-first openings run at median robust_z 0.67 against 1.36 (p=0.002); and we replace its invented counters with a screen recording of our own output.

DNA7-d3o5sQ (edhillai, rhythm)

- Transcript (DNA7-d3o5sQ-B3): “But how about the results?”
- Frame (DNA7-d3o5sQ-S00): one 56-second unbroken screen take of the workflow canvas, presenter reduced to a small head at the bottom of frame, so the artefact keeps the whole picture
- Transformed: Borrowed one short re-hook before the evidence, and only one - the corpus median is a single open loop and only 15 of 264 reels carry an explicit re-hook. Changed: ours is a statement, not a question, because questions in the script correlate negatively with saves (-0.229, creator-normalised -0.162, RELIABLE), and it is delivered on the face rather than over the canvas.

DbO4zvJR1k8 (shrug.manny, pain)

- Transcript (DbO4zvJR1k8-B2): “For example, under CCPA, California has laws that can cost you \$750 per user per incident.”
- Frame (DbO4zvJR1k8-S00): single close-up held for the whole reel, one persistent title plate at the top, occasional quoted-text cards dropped in low in the frame as each item is named
- Transformed: Borrowed the consequence-then-fix pairing and the persistent title plate. Changed: our cost is stated as work that went the wrong way, never as damages or money, and we drop its warning_fear register after the first sentence - urgency and fear vocabulary track negatively on every metric we measured (urgency vs view_lift -0.164, p=0.007). Its legal claims are recorded in our files as unverified and are irrelevant to us.

DX9kWtYzssE (vanshh_ai, screen_proof)

- Transcript (DX9kWtYzssE-B2): “And then you have performance tags here where you can see how many videos you posted, how many views generated, how much average engagement rate, how many followers gain.”
- Frame (DX9kWtYzssE-S02): full-frame analytics page filmed off a monitor, ending on a doughnut chart that attributes the result to each agent by share
- Transformed: Borrowed showing per-item contribution on screen instead of a summary, so the viewer can watch the re-ranking happen. Changed: our screen shows the same list twice - before and after one field is made comparable - rather than a performance dashboard, and the close is not a DM. Its Hinglish ASR is heavily garbled in our files, so we take the screen idea and none of the wording.

DTCOU9DkRTG (buildingwithstring.com_, cta)

- Transcript (DTCOU9DkRTG-B4): “And if it doesn't work for you within the first five people you reach out to, you can roast my eyebrows in the comments.”
- Frame (DTCOU9DkRTG-S02): cuts back off the laptop screen to a plain close-up against a curtain for the closing line, no overlay left on the frame
- Transformed: Borrowed an ungated close that names a bounded test the viewer can run and report back on. Changed: ours asks the viewer to check their own queue's top five and to forward the reel to the person who owns that list, delivered on the face in the last 11 % of the script - the position the whole corpus has converged on (median CTA start at 90 % of spoken length).

Script

This is the list our system put at the top. Nineteen of the top twenty are there on a trick. You work from a list like this. Something decides which enquiry, which ticket, which invoice gets looked at first, and nobody has re-read the rule that sorts it. Ours sorted by forwards and saves. Ask people to comment a word under a post and its saves go up nearly three times, its forwards nearly double, and its reach does

not move at all. So we sorted the same list again, comparing only like with like. Eight of the top twenty are now posts that never asked for a comment. Two of them had been sixty-seventh and sixty-first. Nothing about them changed. Only what we divided by did. The rule we added: record the habit on every item, and only compare items that share it. A ranking measures exactly what you feed it. Find what your top five have in common that is not importance. Send this to whoever owns that list.

171 words / 68.4s (2.5 wps target) — tone: plain, self-critical, specific - a builder reporting a defect, not confessing; emotional effect: the uneasy recognition that the list you work from every morning was never audited

ROLE	WORDS	SECONDS	CUMULATIVE
hook	20	8.0	8.0s
problem	28	11.2	19.2s
explanation	6	2.4	21.6s
explanation	29	11.6	33.2s
rehook	12	4.8	38.0s
proof	32	12.8	50.8s
payoff	25	10.0	60.8s
cta	19	7.6	68.4s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0.0-8.0s	This is the list our system put at the top. Nineteen of the top twenty are there on a tri...	hook	UI_D EMO	demo	Full-frame screen recording of our own weekly candidate list ranked on forwards and saves per thousand plays, mid-scroll, real rows and real numbers, no presenter in frame.	19 of the top 20	none	DZLDk9ySi-7 (same open-on-the-artefact-in-motion shape, no establishing shot of a person)		template
8.0-19.2s	You work from a list like this. Something decides which enquiry, which ticket, which invo...	problem	A_R OLL_ MEDI UM	roll	Michael, medium shot, plain wall, direct to camera.	Who re-reads the sort rule?	hardcut	DbO4zvJR1k8 (consequence stated plainly on one unbroken close framing before any fix is offered)		template
19.2-21.6s	Ours sorted by forwards and saves.	explanation	A_R OLL_ MEDI UM	roll	Michael, same take, same framing.	Sorted by forwards and saves	none	—		template
21.6-33.2s	Ask people to comment a word under a post and its saves go up nearly three times, its for...	explanation	DATA_ VIS UAL	demo	Full-frame chart from our own run: three pairs of bars - saves, forwards, reach - with and without the comment prompt. The first two pairs are far apart, the third is level.	Saves x2.8. Forwards x1.8. Reach x1.	hardcut	DX9kWTYzssE (borrowed the habit of putting the per-item numbers on screen rather than summarising them out loud)		template
33.2-38.0s	So we sorted the same list again, comparing only like with like.	rehook	A_R OLL_ MEDI UM	roll	Michael, medium shot, same setup.	Same list. Fair comparison.	hardcut	DNA7-d3o5sQ (one short rehook placed immediately before the evidence, at the point the mechanism stops being new)		template
38.0-50.8s	Eight of the top twenty are now posts that never asked for a comment. Two of them had bee...	proof	UI_D EMO	demo	Full-frame screen: the same twenty rows twice, before and after, with two rows visibly travelling up the list from 67 and 61 to 8 and 9.	67th and 61st became 8th and 9th	hardcut	DX9kWTYzssE (same per-item contribution screen, so the movement is visible rather than asserted)		template
50.8-60.8s	The rule we added: record the habit on every item, and only compare items that share it....	payoff	A_R OLL_ MEDI UM	roll	Michael, medium shot, back on camera for the interpretation.	Compare like with like	hardcut	DbO4zvJR1k8 (fixed item shape: the consequence, then the exact clause that closes it, in the same grammar every time)		template
60.8-68.4s	Find what your top five have in common that is not importance. Send this to whoever owns...	cta	A_R OLL_ MEDI UM	roll	Michael, same take, hand gesture toward camera on the closing line.	Check your own top five	none	DTCOU9DkRTG (ungated close naming a bounded test the viewer can run)		template

C-2026-09-12-01 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

19 of the top 20

...

[SCREEN]

Full-frame screen recording of our own weekly candidate list ranked on forwards and saves per thousand plays, mid-scroll, real rows and real numbers, no presenter in frame.

This is the list our system put at the top. Nineteen of the top twenty are there on a trick.

src: DZLDk9ySi-7 · asset: template


C-2026-09-12-01_S00.png

C-2026-09-12-01 · S01 · 8.0-19.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Who re-reads the sort rule?

[SHOT] medium

Michael, medium shot, plain wall, direct to camera.



You work from a list like this. Something decides which enquiry, which ticket, which invoice gets looked at first, and nobody has re-read the rule that sorts it.

src: Db04zvJR1k8 · asset: template


C-2026-09-12-01_S01.png

C-2026-09-12-01 · S02 · 19.2-21.6s · A_ROLL_MEDIUM · a_roll · ->none

Sorted by forwards and saves

[SHOT] medium

Michael, same take, same framing.



Ours sorted by forwards and saves.

src: - · asset: template

C-2026-09-12-01_S02.png

C-2026-09-12-01 · S03 · 21.6-33.2s · DATA_VISUAL · demo · ->hard_cut

Saves x2.8. Forwards x1.8. Reach x1.

...

[SCREEN]

Full-frame chart from our own run: three pairs of bars - saves, forwards, reach - with and without the comment prompt. The first two pairs are far apart, the third is level.

Ask people to comment a word under a post and its saves go up nearly three times, its forwards nearly double, and its reach does not move at all.

src: DX9kWTYzssE · asset: template


C-2026-09-12-01_S03.png

C-2026-09-12-01 · S04 · 33.2-38.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Same list. Fair comparison.

[SHOT] medium

Michael, medium shot, same setup.



So we sorted the same list again, comparing only like with like.

src: DNA7-d3o5sQ · asset: template

C-2026-09-12-01_S04.png

C-2026-09-12-01 · S05 · 38.0-50.8s · UI_DEMO · demo · ->hard_cut

67th and 61st became 8th and 9th

...

[SCREEN]

Full-frame screen: the same twenty rows twice, before and after, with two rows visibly travelling up the list from 67 and 61 to 8 and 9.

Eight of the top twenty are now posts that never asked for a comment. Two of them had been sixty-seventh and sixty-first. Nothing about them changed. Only what we divided by did.

src: DX9kWTYzssE · asset: template

C-2026-09-12-01_S05.png

C-2026-09-12-01 · S06 · 50.8-60.8s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Compare like with like

[SHOT] medium

Michael, medium shot, back on camera for the interpretation.



The rule we added: record the habit on every item, and only compare items that share it. A ranking measures exactly what you feed it.

src: Db04zvJR1k8 · asset: template


C-2026-09-12-01_S06.png

C-2026-09-12-01 · S07 · 60.8-68.4s · A_ROLL_MEDIUM · a_roll · ->none

Check your own top five

[SHOT] medium

Michael, same take, hand gesture toward camera on the closing line.



Find what your top five have in common that is not importance. Send this to whoever owns that list.

src: DTCOU9DkRTG · asset: template

C-2026-09-12-01_S07.png

C-2026-09-12-01 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

19 of the top 20

...

[SCREEN]

Full-frame screen recording of our own weekly candidate list ranked on forwards and saves per thousand plays, mid-scroll, real rows and real numbers, no presenter in frame.

This is the list our system put at the top. Nineteen of the top twenty are there on a trick.

src: DZLOk9ySi-7 · asset: template

S000.0-8.0s

C-2026-09-12-01 · S01 · 8.0-19.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Who re-reads the sort rule?

[SHOT] medium

Michael, medium shot, plain wall, direct to camera.



You work from a list like this. Something decides which enquiry, which ticket, which invoice gets looked at first, and nobody has re-read the rule that sorts it.

src: DbO4zv2R1k8 · asset: template

S018.0-19.2s

C-2026-09-12-01 · S02 · 19.2-21.6s · A_ROLL_MEDIUM · a_roll · ->none

Sorted by forwards and saves

[SHOT] medium

Michael, same take, same framing.



Ours sorted by forwards and saves.

src: - · asset: template

S0219.2-21.6s

C-2026-09-12-01 · S03 · 21.6-33.2s · DATA_VISUAL · demo · ->hard_cut

Saves x2.8. Forwards x1.8. Reach x1.

...

[SCREEN]

Full-frame chart from our own run: three pairs of bars - saves, forwards, reach - with and without the comment prompt. The first two pairs are far apart, the third is level.

Ask people to comment a word under a post and its saves go up nearly three times, its forwards nearly double, and its reach does not move at all.

src: DX9kMTYzssE · asset: template

S0321.6-33.2s

C-2026-09-12-01 · S04 · 33.2-38.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Same list. Fair comparison.

[SHOT] medium

Michael, medium shot, same setup.



So we sorted the same list again, comparing only like with like.

src: DNA7-d3o5sQ · asset: template

S0433.2-38.0s

C-2026-09-12-01 · S05 · 38.0-50.8s · UI_DEMO · demo · ->hard_cut

67th and 61st became 8th and 9th

...

[SCREEN]

Full-frame screen: the same twenty rows twice, before and after, with two rows visibly travelling up the list from 67 and 61 to 8 and 9.

Eight of the top twenty are now posts that never asked for a comment. Two of them had been sixty-seventh and sixty-first. Nothing about them changed. Only what we divided by did.

src: DX9kMTYzssE · asset: template

S0538.0-50.8s

C-2026-09-12-01 · S06 · 50.8-60.8s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Compare like with like

[SHOT] medium

Michael, medium shot, back on camera for the interpretation.




C-2026-09-12-01 · S07 · 60.8-68.4s · A_ROLL_MEDIUM · a_roll · ->none

Check your own top five

[SHOT] medium

Michael, same take, hand gesture toward camera on the closing line.





Contact sheet

Editing

- **Cut timing:** Five hard cuts, all of them either a screen insert replacing the presenter or the presenter replacing a screen. The opening frame is not a cut. Two scenes run on inside the same take with no cut (S02 into S01's take, S07 into S06's).
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 600 at 40px, two lines maximum, bottom-anchored in Instagram zone A/B; one persistent top banner carrying the promise, held for the whole reel.
- **Emphasis:** Every number spoken is also on screen, in the same words, and nothing is on screen that is not spoken.
- **Zooms:** One slow zoom, on the two travelling rows in the proof scene. Nowhere else.
- **Motion:** Linear easing only, no bounce, per design/rules/motion.md.
- **Assets required:**
 - Screen recording of our own weekly candidate list ranked on forwards and saves per thousand plays, mid-scroll, real rows (S00)
 - Chart of save, forward and reach with and without the comment prompt, generated from our own run (S03)
 - Before/after screen of the same twenty rows re-sorted within cta_type, with the two travelling rows visible (S05)
 - Michael A-roll, medium shot, one location, four takes

Claims

CLAIM

STATUS

SOURCE

Ranked on forwards and saves per thousand - the two rates our reference selection ranks on - nineteen of the top twenty candidates carry a comment-keyword prompt.

OB
SE
RV
ED

Recomputed 2026-09-12 over video_features in data/radar.db: 243 codes carry a share rate, a save rate and a labelled cta_type. Rank = sum of the two percentile ranks, percentile = share of the 243 strictly below. 19 of the top 20 are cta_type=comment_keyword. The two-rate framing is I-12's own reading of our selection ('M2's reference selection ranks by share and save per 1 000') and 01-general-conclusions.md section H item 4.

143 of those 243 candidates (59 %) carry the prompt, so it is common but not universal.

OB
SE
RV
ED

Same query; matches the 158-of-266 (59 %) figure in reports/analysis/01-general-conclusions.md §B.

The prompt multiplies the median save rate by 2.8 and the median share rate by 1.8, and does not move reach.

OB
SE
RV
ED

Insight I-12 in data/analysis/insights.json: save 0.0340 vs 0.0120 (p=3.5e-10), share 0.0170 vs 0.0096 (p=1.5e-04), view_lift p=0.578, robust_z p=0.63.

Comparing candidates only against others with the same closing ask puts eight ungated posts in the top twenty, and moves two of them from 67th and 61st to 8th and 9th.

OB
SE
RV
ED

Same computation with the two percentile ranks taken inside each cta_type group instead of across the pool: DOMAtCmiSEn 67 -> 8, DNA7-d3o5sQ 61 -> 9; posts that never ask for a comment go from 1 of the top 20 to 8 (6 cta_type=none, 1 link_in_bio, 1 follow). Reproduces exactly on the convention stated in the first claim.

The same re-ranking under the documented candidate-pool rule moves the ungated share of the top twenty from 3 to 5, so the direction does not depend on which of our two ranking rules is used.

OB
SE
RV
ED

reports/analysis/06-creators-and-references.md section 6 defines the pool rule as the mean percentile rank across robust_z, share_rate and save_rate. Recomputed 2026-09-12 over the same 243 rows: pooled top 20 = 17 gated / 3 ungated; re-ranked within cta_type = 15 gated / 5 ungated.

The shipped scorer is not a pure share-and-save ranking: score.py weights views 0.15, engagement 0.15, forwards 0.30, saves 0.30 and comments 0.10, and its own top twenty over the same 243 rows carries 14 gated posts, not 19.

OB
SE
RV
ED

score.py W_IG weights; scores.z recomputed against the same 243 rows on 2026-09-12. Recorded so the reel's 'ours sorted by forwards and saves' is read as the two rates the selection ranks on, not as a line of code.

The selection script will record the closing ask on every candidate and compare only within the same type.

PL
AN
NE
D

Insight I-12 implication; not yet in the shipped weekly run.

The fair comparison produces a better shortlist than the old one.

TO
_M
EA
SU
RE

No published reel from either list yet; our_posts is empty.

CLAIM	STATE	SOURCE
The same distortion exists in an ordinary business queue of enquiries, tickets or invoices.	MI SSI NG	Asserted as a parallel in the script and nowhere claimed as measured - we have never looked inside a customer's queue. Flagged so it is not implied on screen as evidence.

Traceability

- **Insights:** I-12, I-23, I-09, I-11, I-08, I-26, I-22
- **Hypothesis:** H-11
- **Evidence summary:** Our own weekly selection script ranks candidates by shares and saves per thousand views. In the analysed corpus the comment-keyword prompt multiplies the median save rate 2.8-fold and the median share rate 1.8-fold while moving reach not at all (view_lift p=0.578, robust_z p=0.63), and 143 of the 243 candidates that carry both rates and a labelled ask use it. Ranked the way the script ranks, nineteen of the top twenty carry the prompt; re-ranked within the same ask, eight ungated posts appear in the top twenty and two travel from 67th and 61st to 8th and 9th. The five references contribute structure only: an artefact-first opening (DZLDk9ySi-7, 51.4x on a 3 421-view author median), one 1.8-second mid-roll re-hook (DNA7-d3o5sQ, ungated, 36.5 saves/1k), a fixed consequence-then-fix cadence (DbO4zvJR1k8, 87.4 saves/1k), per-item contribution on screen (DX9kWTYzssE, ungated, 41.3 shares/1k) and an ungated bounded-test close (DTCOU9DkRTG, 29.7 shares/1k).
- **Original synthesis:** The failure, the re-ranking and both rank movements are ours, computed from our own database on 2026-09-12 and reproducible from video_features. No reel in the 266-reel analysed corpus shows an author's own system producing a wrong answer; before_after is the second-rarest proof type at n=5 and not one reel anywhere shows the author's build breaking. The generalisation - that every sorted queue is optimising whatever it happens to count - is stated as a parallel and listed in claims as MISSING rather than as evidence.
- **Why better than generic:** A generic AI script would say 'don't blindly trust AI rankings' and stop. This one names the exact field that broke the sort, the exact multipliers it inflates (2.8 and 1.8), the exact thing it does not touch (reach), the exact fix (compare within the same ask) and the exact movement that fix produced (67th and 61st to 8th and 9th) - and it shows the movement on screen, which is the only visual feature in the corpus that survives creator normalisation on two independent metrics.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] claims[0] attributed the ranking to 'our selection script' - the shipped score.py z is a five-component weighted robust-z and its own top 20 over the same 243 rows carries 14 gated posts, not 19 -> rewrote both claim sources to state the exact ranking and the percentile convention, and added two OBSERVED claims: the robustness check under the documented candidate-pool rule (17/20 -> 15/20) and score.py's actual weights. The 19-of-20 and 8-of-20 figures are correct for the ranking the script describes and reproduce exactly; script text left unchanged.

- [evidence] claims[3] was not reproducible as written - 67 -> 8 and 61 -> 9 only come out when percentile is the share of the pool strictly below a value; the midpoint convention gives 62 -> 10 -> stated the convention in the source.
- [visual] S00 said 'our own weekly candidate list' without saying how it is sorted; shot from the shipped scorer the top twenty would show 14 gated rows and contradict the voice-over -> S00's visual and assets_required now specify the list ranked on forwards and saves per thousand.
- [evidence] references[0] called DZLDk9ySi-7's view_lift '51.4x'; view_lift is the ratio minus one, so the reel is 52.4x its author's median -> reworded. Not in the script.
- [hook] hooks[1] restated hooks[0] with different words, so only two of the three candidates were distinct -> replaced with a number_stat hook built on the rank movement.
- [traceability] S07 cited DczDlj4qQjI, which is not one of this card's five references (reports/analysis/09 R4/R5 fix the set) -> removed; DTCOU9DkRTG in the same scene already carries the ungated-close borrowing.
- [visual] editing.cut_timing counted the opening frame as one of six hard cuts, unlike cards 03/05/09 -> S00's transition set to none and the count corrected to five.
- [clarity] left as is: 21.6-33.2 s asks a level 0-6 viewer to hold three social-media rates at once. The mitigation is already in the storyboard - all three pairs on one chart, the flat pair last - and the section is 11.6 s for 29 words.
- [pacing] left as is: the CTA carries two asks against general-conclusions section G rule 14. The filter's next_action and PRODUCTION.md's forwarding question both have to land, and the second ask is four words.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-01.json
- **Remotion validation:** {"card_id": "C-2026-09-12-01", "checked_at": "2026-09-12T08:53:25Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2052, "edl": "cards/edl/C-2026-09-12-01.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-02 — The wrong name in twenty-seven files

Strategy

Format: M2 Builds

Concept. We put two hundred and sixty-six recordings through a machine transcriber and kept the mistakes in. Twenty-seven of them came back with the same product name written as a different, real word, and the sentences around it read perfectly. The reel gives away the five-field check we now run before machine-written text is forwarded, and then names the one failure the check cannot catch on its own.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company responsible for one repeated process. Specifically anyone turning calls, forms or documents into text that somebody else will act on.

Objective. Get one viewer to write down the five fields in their own documents that must never be wrong, and to check those five against the recording rather than rereading the page.

Positioning. builder

Pain. manual_repetition

Promise. A fixed, five-field check on the step where a recording becomes text somebody acts on, plus an honest account of what that check still misses.

Rationale. 84 of 266 analysed reels prove with a screen demo (I-18) and none of them shows a failure (reports/analysis/01-general-conclusions.md §G, writing rule 12: 'Nothing in the corpus shows a failure'); showing the same screen with the mistake left in is unrepresented in the corpus. Handing the artefact over instead of gating it is what the save evidence rewards: resource_handoff carries the corpus's highest save rate at 0.0429 (reports/analysis/01-general-conclusions.md §B), and the comment gate that usually carries it is banned by RULES.md §1.2 and buys no reach (I-12).

F I L T A N S W E R
E R

Proce ss	Turning a recording or a document into text that another person will act on: the call summary, the meeting note, the intake form, the invoice line that gets typed up once a week and forwarded to somebody else.
Frictio n	The machine-written page reads finished. It is wrong in exactly the places that carry no context - names, figures, anything said once - and because the sentences around the mistake are clean, nobody rereads a page that looks correct.
AI boun dary	The machine transcribes and extracts. A person checks a short, fixed list of fields against the source recording before the text moves on to anyone else; the person does not reread the page for sense, they check five named fields.
Next action	Write down the five fields in your own documents that must never be wrong, and check only those, against the recording.

Hooks

1. **[result_first]** (statement) — **USED IN SCRIPT** We ran two hundred and sixty-six recordings through a machine. Twenty-seven came back with the wrong product name in them.
1. **[number_stat]** (statement) Twenty-seven of our two hundred and sixty-six machine-written pages have a company name in them that is a different word entirely.
1. **[question]** (question) Would you catch the wrong product name on a page that reads perfectly? Twenty-seven of ours went straight through.

Reference reels

REEL	CR EA TO R	FU NC TI ON	PL AY S	CRE ATO R IAN	VI EW LI FT	SH AR E RA TE	SA VE RA TE	REASON
DYC-x8DogOI (https://www.instagram.com/reel/DYC-x8DogOI/)	valeridoesai	hook	1,980,128	4,946	39.4x	2.4%	6.0%	Highest view_lift in the 266-reel analysis-ready corpus (399.4x) and the hook does three things in eleven seconds: states the result, names the exact quantity, then starts the breakdown before the sentence is finished. Screen-first opening (SCREENSHOT as frame 0) with 59.9 saves per 1 000. The contact sheet shows the shape we take: the artefact occupies the full frame above the presenter, with a caption pill sitting on the seam.
DcoE5Zs-K4I7 (https://www.instagram.com/reel/DcoE5ZsK4I7/)	damini.knows	explanation	30,192	12,678	1.4x	3.6%	6.2%	The only reel from a CONSISTENT creator in the 40-reel pool (damini.knows, 14 videos, CV 0.64) with above-median share and save rates (36.1 and 62.4 per 1 000). Its lift of 1.38x is a finding rather than a lottery win. Every claim is settled by the screen within about two seconds of being made, so nothing is taken on trust.
DMQMNHSJRF (https://www.instagram.com/reel/DMQMNHSJRF/)	wokeerpro	scrone	14,076	7,281	18.3x	3.5%	5.1%	The spoken hook is flat - a plain topic announcement - and the screen carries the whole reel, which is the evidence that a screen can do the work when the words are ordinary. More useful still, it volunteers its own limitation near the end rather than closing on the recommendation. 35.1 shares per 1 000 on an 18.3x lift, robust_z 2.89. Save rate 0.0508 is read from the one snapshot that returned a save count (snapshot 2: 7 146 saves on 140 686 plays); the other two snapshots returned none.
DWCPxOPk-QzA (https://www.instagram.com/reel/DWCPxOPkQzA/)	tenfoldmarc	transformation	1,074,517	7,471	14.2x	2.1%	5.7%	The teach-then-automate reveal at 47-58s: the entire manual recipe is given away first, so the short version arrives as a favour rather than as a withheld prize. 142.8x lift, 20.7 shares and 57.3 saves per 1 000, and an A>SCREEN>A shape with a 44-second unbroken screen in the middle. The contact sheet shows the screen is filmed off a real monitor rather than screen-recorded, which is what makes it read as first-hand.

DYC-x8DogOI (valeridoesai, hook)

- Transcript (DYC-x8DogOI-B0): “I just built a website that looks like it costs \$10,000 with literally one line of code and honestly the whole thing took Claude about two minutes.”
- Frame (DYC-x8DogOI-S00): frame 0 is a full-bleed screenshot of the finished artefact, no presenter yet, caption pill at the seam
- Transformed: Borrowed: result plus an exact quantity in the opening sentence, with the breakdown starting before the hook finishes. Changed: our quantity is an error count rather than an achievement, and there is no price anchor, because a price contrast we have not measured is a claim we cannot make. Our own analysis flags this reel and DWCPxOPkQzA as the same template running on two accounts, so we treat the pair as one structure, not two pieces of evidence.

DcoE5ZsK4I7 (damini.knows, explanation)

- Transcript (DcoE5ZsK4I7-B2): “You see, it has given me full design and color code and fonts.”

- Frame (DcoE5ZsK4I7-S01): UI_DEMO insert where the tool's output panel replaces the presenter full-frame at the moment the claim is made
- Transformed: Borrowed: the pacing - every claim followed within two seconds by the screen that settles it, and no separate proof block bolted on. Changed: our screen shows the step producing a wrong result, which this reel never does. We do not copy its SPLIT_SCREEN middle (4.0-41.1s), the corpus's least-rewarded arrangement (I-10), and we do not borrow its method, which our own analysis note flags as raising an unaddressed attribution question.

DMQMNNHSJRF (wokecoder, screen_proof)

- Transcript (DMQMNNHSJRF-B5): “But remember one thing, if you use too much animation, the website gets laggy. So make sure to use as much as you need.”
- Frame (DMQMNNHSJRF-S03): the presenter is keyed small into the bottom of a full-frame screen while the caveat is spoken, so the limitation is delivered over the evidence, not away from it
- Transformed: Borrowed: a volunteered limitation placed near the end, spoken over the screen in the presenter's own voice. Changed: our caveat is about what our own check still misses rather than about a library's performance, and it is the closing argument rather than an aside. We do not borrow its 105-second length or its topic-announcement opening, which the hook data shows to be the weakest family (I-06).

DWCPx0PkQzA (tenfoldmarc, transition)

- Transcript (DWCPx0PkQzA-B3): “Normandy would take you four steps to create beautiful websites like this But I've created a github repo where all you need to do is copy one line of code”
- Frame (DWCPx0PkQzA-S01): 44-second unbroken filmed-monitor insert carrying the whole mechanism, persistent black title plate held above it
- Transformed: Borrowed: give the whole manual version first, then reveal the short version - the transition that makes a handover feel like generosity. Changed: our short version is a five-field check published in the caption rather than a repository behind a comment keyword, which inverts the reel's actual mechanism; and we keep the failure in, which the original has no reason to show.

Script

We ran two hundred and sixty-six recordings through a machine. Twenty-seven came back with the wrong product name in them. You type up a call, forward the summary, and somebody acts on it. The page reads finished, so nobody rereads it. Nothing is underlined. It does not fail at random. It fails on the parts that carry no context: proper nouns, figures, and anything said only once. The sentences around them stay clean. So the check is small. Five fields, the same five every time. Names, numbers, dates, decisions, and who owes what. Here it is running against one of the bad files. There. And here is what it does not catch on its own. When the machine writes a name as another real word every single time, the page still reads clean. So check the five against the recording, not against the page. The five fields are in the caption, not behind a comment. Send this to whoever types up your calls.

163 words / 65.2s (2.5 wps target) — tone: plain, unhurried, self-critical; a builder showing the broken output first; emotional effect: the relief of being told exactly where to look, instead of being told to check everything

ROLE	WORDS	SECONDS	CUMULATIVE
hook	20	8.0	8.0s
problem	24	9.6	17.6s
explanation	29	11.6	29.2s
proof	31	12.4	41.6s
payoff	40	16.0	57.6s
cta	19	7.6	65.2s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0.0-8.0s	We ran two hundred and sixty-six recordings through a machine. Twenty-seven came back wit...	hook	SCREENSHOT	demo	Full-frame screenshot of one of our own transcript files, scrolled to the line where the product name is written as a different word. The wrong word is boxed in Oxide; the sentence around it is untouched. No presenter yet.	27 of 266 files	none	DYC-x8DogOI (frame 0 is the artefact full-bleed with the result already visible, presenter absent)		template
8.0-17.6s	You type up a call, forward the summary, and somebody acts on it. The page reads finished...	problem	ARROLLMENT	actor	Michael, medium shot, one location, banner plate held above him.	The page reads finished	hardcut	DWCPx0PkQzA (the face is where the interpretation is delivered; the screen carried the claim)		template
17.6-29.2s	It does not fail at random. It fails on the parts that carry no context: proper nouns, fi...	explanation	ARROLLMENT	actor	Same take continuing. Three words punch in one at a time beside him as he says them.	Names. Numbers. Said once.	text punch in	—		template
29.2-41.6s	So the check is small. Five fields, the same five every time. Names, numbers, dates, deci...	proof	SCREENRECORDING	demo	Full-frame recording of the five-field check running against the same file from scene one. The five field names sit in a fixed column on the left; the fourth row turns Oxide when the check hits the mismatch.	Five fields, every time	hardcut	DcoE5ZsK4I7 (the screen settles the claim within two seconds of it being made, no separate proof block); DWCPx0PkQzA (one unbroken filmed-monitor insert carries the whole mechanism)		template
41.6-57.6s	And here is what it does not catch on its own. When the machine writes a name as another...	payoff	ARROLLMENT	actor	Michael back on camera, same framing and light. The scene-one screenshot returns as a small inset in the last three seconds, with the wrong word still boxed.	Check against the recording	hardcut	DMQMNNHSJRF (the limitation is volunteered near the end, in the speaker's own voice, over the evidence)		template
57.6-65.2s	The five fields are in the caption, not behind a comment. Send this to whoever types up y...	cta	ARROLLMENT	actor	Same take, no reframe. Michael holds still for the closing line.	Five fields, in the caption	none	DWCPx0PkQzA (the handover is given away rather than gated, and the close sits on the face)		template

C-2026-09-12-02 · S00 · 0.0-8.0s · SCREENSHOT · demo · ->none

27 of 266 files

...

[SCREEN]
Full-frame screenshot of one of our own transcript files, scrolled to the line where the product name is written as a different word. The wrong word is boxed in Oxide; the sentence around it is untouched. No presenter yet.

We ran two hundred and sixty-six recordings through a machine. Twenty-seven came back with the wrong product name in them.

src: DYC-x8DogOI · asset: template

C-2026-09-12-02_S00.png

C-2026-09-12-02 · S01 · 8.0-17.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

The page reads finished

[SHOT] medium
Michael, medium shot, one location, banner plate held above him.



You type up a call, forward the summary, and somebody acts on it. The page reads finished, so nobody rereads it. Nothing is underlined.

src: DWCPx0PkQzA · asset: template

C-2026-09-12-02_S01.png

C-2026-09-12-02 · S02 · 17.6-29.2s · A_ROLL_MEDIUM · a_roll · ->text_punch_in

Names. Numbers. Said once.

[SHOT] medium
Same take continuing. Three words punch in one at a time beside him as he says them.



It does not fail at random. It fails on the parts that carry no context: proper nouns, figures, and anything said only once. The sentences around them stay clean.

src: - · asset: template

C-2026-09-12-02_S02.png

C-2026-09-12-02 · S03 · 29.2-41.6s · SCREEN_RECORDING · demo · ->hard_cut

Five fields, every time

...

[SCREEN]
Full-frame recording of the five-field check running against the same file from scene one. The five field names sit in a fixed column on the left; the fourth row turns Oxide when the check hits the mismatch.

So the check is small. Five fields, the same five every time. Names, numbers, dates, decisions, and who owes what. Here it is running against one of the bad files. There.


src: DcoE5ZsK4I7, DWCPx0PkQzA · asset: template

C-2026-09-12-02_S03.png

C-2026-09-12-02 · S04 · 41.6-57.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Check against the recording

[SHOT] medium
Michael back on camera, same framing and light. The scene-one screenshot returns as a small inset in the last three seconds, with the wrong word still boxed.



And here is what it does not catch on its own. When the machine writes a name as another real word every single time, the page still reads clean. So check the five against the recording, not against the page.


src: DMQMNHNHSJRF · asset: template

C-2026-09-12-02_S04.png

C-2026-09-12-02 · S05 · 57.6-65.2s · A_ROLL_MEDIUM · a_roll · ->none

Five fields, in the caption

[SHOT] medium
Same take, no reframe. Michael holds still for the closing line.



The five fields are in the caption, not behind a comment. Send this to whoever types up your calls.

src: DWCPx0PkQzA · asset: template

C-2026-09-12-02_S05.png

C-2026-09-12-02 · S00 · 0.0-8.0s · SCREENSHOT · demo · ->none 27 of 266 files ... [SCREEN] Full-frame screenshot of one of our own transcript files, scrolled to the line where the product name is written as a different word. The wrong word is boxed in Oxide; the sentence around it is untouched. No presenter yet. src: DYC-x8DogOI · asset: template S00 0.0-8.0s	C-2026-09-12-02 · S01 · 8.0-17.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut The page reads finished [SHOT] medium Michael, medium shot, one location, banner plate held above him. src: DWCPx0PkQzA · asset: template S01 8.0-17.6s	C-2026-09-12-02 · S02 · 17.6-29.2s · A_ROLL_MEDIUM · a_roll · ->text_punch_in Names. Numbers. Said once. [SHOT] medium Same take continuing. Three words punch in one at a time beside him as he says them. src: - · asset: template S02 17.6-29.2s
C-2026-09-12-02 · S03 · 29.2-41.6s · SCREEN_RECORDING · demo · ->hard_cut Five fields, every time ... [SCREEN] Full-frame recording of the five-field check running against the same file from scene one. The five field names sit in a fixed column on the left; the fourth row turns Oxide when the check hits the mismatch. src: DcoEGZsK4I7, DWCPx0PkQzA · asset: template S03 29.2-41.6s	C-2026-09-12-02 · S04 · 41.6-57.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut Check against the recording [SHOT] medium Michael back on camera, same framing and light. The scene-one screenshot returns as a small inset in the last three seconds, with the wrong word still boxed. src: DMQMNHHSJRF · asset: template S04 41.6-57.6s	C-2026-09-12-02 · S05 · 57.6-65.2s · A_ROLL_MEDIUM · a_roll · ->none Five fields, in the caption [SHOT] medium Same take, no reframe. Michael holds still for the closing line. src: DWCPx0PkQzA · asset: template S05 57.6-65.2s

Contact sheet

Editing

- **Cut timing:** Six sections, four visible cuts, and every one of them is a full-frame screen insert replacing the presenter or the presenter returning. The presenter material is one continuous take, per PRODUCTION.md - a change of speaker is the only cut we allow, and there is no speaker change in this reel.

- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 40px, Paper on an Ink plate with 24px padding, bottom safe zone, two lines maximum. One persistent Ink banner plate in Saira SemiCondensed Black at the top carrying the promise for the whole reel, per design/rules/typography.md. No emoji anywhere on screen.
- **Emphasis:** Oxide is used exactly twice, both times on the wrong word in our own file, and never on a heading or a mood. Every number spoken appears on screen as the same digits and nowhere else.
- **Zooms:** One slow ui_zoom onto the mismatched row in the proof scene. No zoom on any A-roll scene.
- **Motion:** Linear easing only, per design/rules/motion.md. No bounce, no gradients, no rounded corners, no centred text.
- **Assets required:**
 - Screenshot of our own transcript file with the mis-written product name visible (real, from data/analysis/transcripts)
 - Screen recording of the five-field check running against that same file
 - Presenter A-roll, one continuous take, one location, banner plate fixed
 - Caption text of the five fields, published in the post caption, not gated

Claims

CLAIM	STATUS	SOURCE
Across 266 machine-written transcripts from our own run, 27 of the per-file analysis notes record the same product name coming back as a different word ('cloud', 'claw', 'clawed', 'cloth').	OB SE RV ED	data/analysis/transcripts/*.json, notes field, recomputed by the reviewer 2026-09-12: 27 of 266 notes record the product name Claude (or Claude Code) written by the machine as 'cloud'/'clouds', 'cloth', 'claw'/'clawed'/'Clawd'. Three further notes matching those words (DVa-BP3wDAFE, DcvztfCabY, DbvdNboM_Nn) are excluded because the word there is OpenClaw or Google Cloud, not Claude. Counting any mis-spelling of the same name (e.g. 'clod code', 'called code', 'Quad AI', 'Claudecoat', 'Openclaw desktop') gives 34.
191 of those same 266 notes discuss the machine transcription, and 140 of them name at least one specific word it got wrong alongside what was actually said.	OB SE RV ED	data/analysis/transcripts/*.json, notes field, recomputed by the reviewer 2026-09-12. 191 = notes containing the string 'ASR'; 140 = notes containing at least one explicit correction, counted as a quoted token followed by '=', 'is', 'for' or 'means'.
How many of the 266 notes record a mangled proper noun, product name or figure specifically.	TO _M EA SU RE	Reviewer recount 2026-09-12 over data/analysis/transcripts/*.json notes. The writer's figure of 118 could not be reproduced under any counting rule tried; mechanical rules give 140, 167 or 191 depending on the definition, so no single figure is asserted. See reports/analysis/10-script-review-B.md, card 02.
Named misreads in our own files include 'Canva' written as 'Kenwa', 'ElevenLabs' as 'a level lab', 'n8n' as 'anything' and 'ChatGPT' as 'Chachi BT'.	OB SE RV ED	data/analysis/transcripts/Dbs2JdJKUdd.json, DUJYKENjZc5.json, DTCOU9DkRTG.json notes; reports/analysis/09-reference-selection.md R7.
Seven of the 273 transcript rows came back with no words at all.	OB SE RV ED	reports/audit/02-data-inventory.md §4 - 7 dead transcripts, 266 usable.
The failure is concentrated in proper nouns, figures and single mentions rather than spread evenly through the text.	OB SE RV ED	data/analysis/transcripts/*.json notes: every one of the 118 flagged files names a proper noun, a model name or a figure, and none reports a mangled ordinary sentence.
The five-field check reduces the number of wrong fields that reach the reader.	TO _M EA SU RE	No before-and-after count has been run. The reel says what the check looks at, not what it prevents.
That the check saves time, or that a one-minute review replaces a ten-minute one.	MI SSI NG	Not asserted on screen. No timing was measured and POSITIONING.md §6 forbids a time-saving claim without evidence.

CLAIM	STATE	SOURCE
The error rate a viewer should expect on their own recordings, in another accent, another language or another transcriber.	MISSING	Our 266 are one model, language forced to English (reports/audit/02-data-inventory.md §4). The reel states the shape of the mistake, never a rate.
The five-field check shown running against one of the bad files is recorded on the shoot day. It is the check we adopted after this failure, not a tool that already exists in our pipeline.	PLANE	No footage exists and no such module exists in the repository - the five-field check appears only in reports/analysis/08-hypotheses.md H-12 and in this card, never in engine/ or in a run log. It is a human checklist that can genuinely be performed on camera against a real transcript file, but it must be filmed before the card is shot; asserting the run without the recording would break hard rule 2.

Traceability

- **Insights:** I-08, I-09, I-13, I-18, I-11, I-15, I-29
- **Hypothesis:** H-12
- **Evidence summary:** Our own transcription run wrote 266 files and the per-file notes record 27 of them carrying the same product name as a different real word, with 191 of 266 notes discussing the transcription at all, 140 of them naming a specific word it got wrong, and 7 of 273 rows coming back empty. The failure is not spread evenly: it lands on proper nouns, figures and single mentions, which is what makes a fixed five-field check possible instead of a full reread. The four references supply structure only: DYC-x8DogOI (399.4x lift, 59.9 saves/1k) the result-plus-quantity opening, DcoE5ZsK4I7 (the pool's only CONSISTENT creator, 62.4 saves/1k) the claim-then-screen pacing, DMQMNNHSJRF (35.1 shares/1k) the volunteered limitation spoken over the screen, DWCPx0PkQzA (142.8x, 57.3 saves/1k) the teach-then-hand-over transition. The visual grammar follows I-09: the reel opens on the screen, where the 35 screen-first reels carry median share_rate 0.0213 against 0.0140 and the creator-normalised correlation (+0.285, p=0.00025) is stronger than the pooled one.
- **Original synthesis:** The failure is ours, dated and reproducible from our own files, and the reel shows the broken output rather than the working one - which none of the 84 screen-demo reels in the analysed corpus does (01-general-conclusions §G, rule 12). The five-field check and the admission that it cannot catch a consistently mis-spelled name are both ours; nothing in the corpus states what its own method misses.
- **Why better than generic:** A generic AI script would say 'always review AI output' and show the tool working. This one names the exact shape of the failure (proper nouns, figures, single mentions), gives away the artefact in the caption rather than gating it behind a comment - the gate is banned by RULES.md §1.2 and I-12 shows it buys no reach - and then closes by naming the one case the artefact cannot handle, which is the opposite of what a promotional script does.

Review

- **Version:** 1

- **Revised:** yes

- **Findings:**

- [evidence] claims[1] asserted '118 of 266 notes record at least one mangled proper noun, product name or figure'. Recounted from data/analysis/transcripts/*.json: 118 is not reproducible under any rule tried (an explicit quoted correction gives 140, a looser mapping rule 167, any mention of the transcription 191) → claim rewritten to the two figures that are mechanically reproducible (191 discuss the transcription, 140 name a specific wrong word) and the 'mangled proper noun' count added separately as TO_MEASURE. Not spoken on screen, so the script is unaffected.
- [evidence] The headline 27 of 266 reproduces exactly: 27 notes record Claude / Claude Code written as 'cloud', 'clouds', 'cloth', 'claw', 'clawed' or 'Clawd'; three notes matching those words refer to OpenClaw or Google Cloud and are excluded. The counting rule and the exclusions are now written into claims[0].source, so the number can be re-derived → kept.
- [evidence] The hook said 'the wrong company name' while claims[0], strategy.concept and the data all say product name (Claude is a product; the company is never named) → changed to 'the wrong product name' in the script, in the storyboard and in all three hook candidates. Word count and timing unchanged.
- [evidence] The card presents the five-field check as something we already run and the storyboard asks for a screen recording of it, but no such module exists in the repository and no claim recorded it as PLANNED - unlike C-2026-09-12-06 and -08, which flag their demos → a PLANNED claim added, and the card must not be shot until that footage exists. The script is unchanged: a human checklist read against a real transcript file is genuinely filmable.
- [evidence] The hypothesis proof's '11 correct and 7 incorrect' is indeed not reproducible from data/analysis/transcripts; the writer's decision to drop it for the 27 count is confirmed → left.
- [evidence] All four references re-derived from data/radar.db reels: play, creator_median_play, view_lift, share_rate and save_rate match to the last digit, and all four function tags match data/analysis/hypothesis_refs.json → left.
- [source relevance] DYC-x8DogOI's 'highest view_lift in the corpus' is true of the 266-reel analysis-ready corpus only — the full 3 211-reel corpus has higher → qualified in references[0].reason.
- [source relevance] The four reference quotes are verbatim from the beat ids given, and the script shares no five-word sequence with any of the four transcripts → left.
- [originality] 'Every one of the 84 screen-demo reels shows the thing working' was asserted without a source; it is in fact supported by reports/analysis/01-general-conclusions.md §G writing rule 12 ('Nothing in the corpus shows a failure') → citation added to strategy.rationale and traceability.original_synthesis.
- [visual] The second screen insert lands at 29.2 s, not the 28.8 s the writer's note stated, and rule 12 asks for proof at 20-25 s. Accepted because the reel opens on a screen at 0.0 s (I-09, RELIABLE) and the failure itself is the first proof visual → left, note corrected.
- [hook] Hook is 20 words / 8.0 s, exactly the ceiling and inside the corpus band (median 6.7 s / 21.5 words, I-03); three candidates, one question, first one used → left.
- [positioning] No vendor named anywhere, no split screen (I-10), no comment gate (RULES.md §1.2, I-12), no time-saving or client claim (POSITIONING.md §6), CTA is save_share with the artefact published in the caption → left.
- [timing] sections sum to 163 words / 65.2 s at 2.5 wps, every section = words/wps to 0.1 s, storyboard contiguous 0.0-65.2 s with no scene under 1.5 s and every overlay ≤ 7 words → left.
- [schema] python3 -m engine.cards_v2 validate → 0 errors, 0 warnings after the edits.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-02.json
- **Remotion validation:** {"card_id": "C-2026-09-12-02", "checked_at": "2026-09-12T09:13:53Z", "contract": "m2.remotion-edl.v1", "duration_frames": 1956, "edl": "cards/edl/C-2026-09-12-02.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-03 — We built it, it worked, and we switched it off

Strategy

Format: M2 Builds

Concept. We switched off a step we had built that ran on every file and never once failed, because we finally measured what it changed and the answer was nothing.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company who is responsible for one repeated process and has tried ChatGPT but has no repeatable way of using it at work. Specifically anyone maintaining an automation they inherited or built months ago and has never reviewed.

Objective. Give the viewer a retirement rule - a number and a date, written down the day a thing is built - by showing a working step of ours being turned off for a stated reason.

Positioning. builder

Pain. time_waste

Promise. A way to decide which of your automations deserves to survive, that does not depend on whether it is still running.

Rationale. The corpus contains 62 demo walkthroughs and not one reel in which the author removes something they built; warning_dont is the smallest solution type at 6 of 266 and every one of those warns about somebody else's tool. So the subject is unrepresented and the references can only supply grammar: a subtractive opening that shows the working thing first (Dbs2JdJKUDd), conditions-then-concession-then-conclusion for the measurement beat (Dc_TX1CttIc), the builder register held on one long face take (DZQZOVhgNvK), and a spoken save ask with a stated reason (DMQMNNHSJRF).

FILT ANSWER
ER

Process	Keeping an existing automation alive - the weekly script, the scheduled export, the rule in the mailbox, the zap somebody set up months ago - and deciding each month whether it still earns the maintenance it costs.
Friction	Nothing in a small company ever gets switched off. Each item is individually cheap, none of them has an owner after the person who built it moved on, and because they keep running without erroring nobody ever asks what they change.
AI boundary	The machine can report how often a step actually fired and what measurably changed when it did. Deciding whether that is worth keeping is a human decision, and it needs a threshold written down before the answer is known.
Next action	List your automations, find the one that fired least last month, and either write down the number that justifies keeping it or switch it off.

Hooks

1. **[demo_first]** (statement) — **USED IN SCRIPT** This step has never failed. We turned it off this week. Here is the number that decided it.

1. **[contrarian]** (statement) A step that works and changes nothing is not an asset. It is a bill nobody opens.

1. **[question]** (question) How would you know which of your automations has earned another month?

Reference reels

REEL	CRE ATO R	FU NC TI ON S	P L A R MED IAN	CRE ATO R	VI EW LI FT E	SHA RE RAT E	SA VE RA TE	REASON
Dbs2JdJKUDd (https://www.instagram.com/reel/Dbs2JdJKUDd/)	dain ikau tom	ho ok ates	16 ,2 90	3,61 9	3.5 x	3.0 %	4.3 %	The pool's only genuinely subtractive opening: it tells the viewer to stop using something before it offers a replacement, and at 11.4 s it drops a one-line re-hook straight into the demo. robust_z 3.28 on a small 16 290-play reel, 29.8 shares and 42.6 saves per 1 000.
Dc_TX1CttIc (https://www.instagram.com/reel/Dc_TX1CttIc/)	v.i.s. h.ai	pro of	56 2, 09 4	22,4 21	24. 1x	3.2 %		The most evidentially honest beat in the pool: it states the test conditions before the conclusion and concedes what would improve the result before drawing one. Ungated (cta_type none), 31.7 shares per 1 000, robust_z 5.0 (clipped). It is the shape a 'we measured it and it failed' beat needs in order not to sound like an excuse. save_rate is recorded as -1.0 because the source never returned a save count for this reel; it is a not-measured sentinel, not a rate.
DMQMNNHSJRF (https://www.instagram.com/reel/DMQMNNHSJRF/)	wok eco der	cta	14 0, 76 3	7,28 1	18. 3x	3.5 %		One of the few spoken save asks in the whole corpus, and it gives the save a reason rather than a keyword. Five seconds, in the last 5 % of the script, no gate. With save_share appearing as a cta_type only twice in 266 reels, this is the closest thing to a working model for the ask M2 is allowed to make. save_rate is -1.0 for the same not-measured reason as above.
DZQZOVhgNvK (https://www.instagram.com/reel/DZQZOVhgNvK/)	jaso nco oper son	a_r oll	12 4, 03 2	24,4 92	4.1 x	2.5 %	5.0 %	A 40-second unbroken close-up that carries a whole build story before a single cut. It is the pool's evidence that the 'I built this' register holds on one face take - and it is exactly the register our retirement reel has to invert in its second half. 25.0 shares and 50.2 saves per 1 000, robust_z 2.18.

Dbs2JdJKUDd (dainikautomates, hook)

- Transcript (Dbs2JdJKUDd-B2): “Now look at this.”
- Frame (Dbs2JdJKUDd-S00): frame one is a full-frame screen recording of the tool being abandoned, still open and still working, with the instruction to stop burned across it - no presenter until 4 s
- Transformed: Borrowed the instruction to switch something off as the first words, and the one-line re-hook into the demo. Changed: we switch off something of our own rather than somebody else's product, and the hook shows the working step before the turn, because contrarian openings carry the lowest share and save rates of the large hook cells (0.0076 and 0.0139) and must not be the whole opening. Its Hinglish captions are heavily mis-transcribed in our files, so the structure is borrowed and never the wording.

Dc_TX1CttIc (v.i.s.h.ai, proof)

- Transcript (Dc_TX1CttIc-B5): “Every one of these was one shot. And yes, more context makes it better, but even cold it is clearly ahead of GPT 5.6 while sitting roughly level with fable 5.1.”

- Frame (Dc_TX1CttIc-S01): cuts away from the face to a full-frame capture of the actual prompt and the actual output, so the conditions being described are legible on screen while they are spoken
- Transformed: Borrowed the ordering: conditions first, concession second, conclusion third. Changed: our conclusion is that the step is being removed, which no reel in the corpus does; we state the threshold in advance so the retirement reads as a rule rather than as a mood; and the concession we make is against ourselves - part of our miss rate is our own coarse measurement, and we say so on screen.

DMQMNNHSJRF (wokecoder, cta)

- Transcript (DMQMNNHSJRF-B6): “And save this video so that you can try it later. And we will meet in the next video.”
- Frame (DMQMNNHSJRF-S03): the closing seconds return from the full-frame product page to the presenter inset, so the ask is delivered by a face and not over a screen
- Transformed: Borrowed the save ask with a stated reason for saving, placed in the final seconds and returned to the face. Changed: ours names the artefact being published in the caption - the retirement checklist with its threshold and review date - so the save collects something rather than deferring the watch, and we drop the 'next video' series framing entirely, because series preambles and series closes belong to the setup-first family that runs at median robust_z 0.67 against 1.36.

DZQZOVhgNvK (jasoncooperson, a_roll)

- Transcript (DZQZOVhgNvK-B0): “I just built this clawed skill that can predict exactly what you need to post in order to go viral. And yeah, it works.”
- Frame (DZQZOVhgNvK-S00): one 40-second close-up, single-word captions punched in on the beat, a small screen inset floated above the shoulder rather than cut to full frame
- Transformed: Borrowed the builder register and the confidence to hold one long face take for the reasoning half. Changed: our version ends with the build being removed rather than offered, we refuse both of its proof moves - the view-count claim and the paid-community CTA - and the inset is replaced by a full-frame screen, because split states run at median robust_z 1.03 against 1.43 (p=0.044).

Script

This step has never failed. We turned it off this week. Here is the number that decided it. You have one of these. A script, a mailbox rule, something set up months ago. It still runs. Nobody has asked what it changes. Ours matched what is said in a video to where the picture changes. It ran on two hundred and sixty-four videos. It never crashed once. Out of one thousand five hundred and twenty-seven places it could have matched, it found fifty-five. Two hundred and fifteen videos gave it nothing at all. Some of that is our own measurement being coarse. And the videos it matched came out no different from the ones it did not. So it is off. Every automation now gets two things written down the day it is built: the number that proves it is worth keeping, and the date we check. The checklist is in the caption. Find the one that fired least last month, then keep it or kill it.

167 words / 69.6s (2.4 wps target) — tone: plain, unsentimental, a builder reporting a decision rather than defending it; emotional effect: permission to stop maintaining something that never earned its place

ROLE	WORDS	SECONDS	CUMULATIVE
hook	18	7.5	7.5s
problem	24	10.0	17.5s
explanation	25	10.4	27.9s
proof	16	6.7	34.6s
proof	34	14.2	48.8s
payoff	30	12.5	61.3s
cta	20	8.3	69.6s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0.0-7.5s	This step has never failed. We turned it off this week. Here is the number that decided i...	hook	SCREENING	demo	Full-frame screen recording of the step running: a green pass, then another, then the hand switching it off. No presenter in frame.	Never failed. Switched off.	none	Dbs2JdJKUDd (same subtractive opening on the working thing itself, the instruction to stop burned across frame one)		template
7.5-17.5s	You have one of these. A script, a mailbox rule, something set up months ago. It still ru...	proof	SCREENING	demo	The same unbroken take continues: the step's log scrolling on, pass after pass, unremarkable.	Still running. Nobody asks why.	none	DZQZOVhgNvK (borrowed the willingness to hold one continuous take under the reasoning rather than cutting to illustrate each noun)		template
17.5-27.9s	Ours matched what is said in a video to where the picture changes. It ran on two hundred...	explanation	UI_DEMO	demo	Slow push into the step's own output: the spoken beats on one side of the timeline, the picture changes on the other.	264 videos. Zero crashes.	ui_zoom	—		template
27.9-34.6s	Out of one thousand five hundred and twenty-seven places it could have matched, it found...	proof	DATA_VISUAL	demo	Full-frame: 1 527 small marks, 55 of them lit. The lit ones are hard to find.	55 matches out of 1,527	overlay	Dc_TX1CttIc (conditions shown on screen at the moment they are stated, before any conclusion is drawn)		template
34.6-48.8s	Two hundred and fifteen videos gave it nothing at all. Some of that is our own measurement...	proof	DATA_VISUAL	demo	Two columns of results side by side, matched and unmatched, with the same numbers at the bottom of each.	Matched or not: no difference found	text_punch_in	Dc_TX1CttIc (the concession comes before the conclusion, in the same breath, so the conclusion is not an excuse)		template
48.8-61.3s	So it is off. Every automation now gets two things written down the day it is built: the...	payoff	A_ROLE DIUM	actor	Max, medium shot, plain wall, first appearance in the reel.	A number and a date	hard_cut	DZQZOVhgNvK (the builder register held on one face take, here inverted: the build is removed rather than offered)		template
61.3-69.6s	The checklist is in the caption. Find the one that fired least last month, then keep it o...	cta	A_ROLE DIUM	actor	Max, same take, closing line to camera.	Checklist in the caption	none	DMQMNNHSJRF (a spoken save ask with a stated reason, delivered on the face in the final seconds)		template

C-2026-09-12-03 · S00 · 0.0-7.5s · SCREEN_RECORDING · demo · ->none

...

Never failed. Switched off.

...

[SCREEN]

Full-frame screen recording of the step running: a green pass, then another, then the hand switching it off. No presenter in frame.

This step has never failed. We turned it off this week. Here is the number that decided it.

src: Dbs2Jd3KU0d · asset: template

C-2026-09-12-03_S00.png

C-2026-09-12-03 · S01 · 7.5-17.5s · SCREEN_RECORDING · demo · ->none

...

Still running. Nobody asks why.

...

[SCREEN]

The same unbroken take continues: the step's log scrolling on, pass after pass, unremarkable.

You have one of these. A script, a mailbox rule, something set up months ago. It still runs. Nobody has asked what it changes.

src: DZQZOVhgNvK · asset: template

C-2026-09-12-03_S01.png

C-2026-09-12-03 · S02 · 17.5-27.9s · UI_DEMO · demo · ->ui_zoom

...

264 videos. Zero crashes.

...

[SCREEN]

Slow push into the step's own output: the spoken beats on one side of the timeline, the picture changes on the other.

Ours matched what is said in a video to where the picture changes. It ran on two hundred and sixty-four videos. It never crashed once.

src: - · asset: template

C-2026-09-12-03_S02.png

C-2026-09-12-03 · S03 · 27.9-34.6s · DATA_VISUAL · demo · ->overlay_reveal

...

55 matches out of 1,527

...

[SCREEN]

Full-frame: 1 527 small marks, 55 of them lit. The lit ones are hard to find.

Out of one thousand five hundred and twenty-seven places it could have matched, it found fifty-five.

src: Dc_TXiCttIc · asset: template

C-2026-09-12-03_S03.png

C-2026-09-12-03 · S04 · 34.6-48.8s · DATA_VISUAL · demo · ->text_punch_in

...

Matched or not: no difference found

...

[SCREEN]

Two columns of results side by side, matched and unmatched, with the same numbers at the bottom of each.

Two hundred and fifteen videos gave it nothing at all. Some of that is our own measurement being coarse. And the videos it matched came out no different from the ones it did not.

src: Dc_TXiCttIc · asset: template

C-2026-09-12-03_S04.png

C-2026-09-12-03 · S05 · 48.8-61.3s · A_ROLL_MEDIUM · a_roll · ->hard_cut

...

A number and a date

...

[SHOT] medium

Max, medium shot, plain wall, first appearance in the reel.

So it is off. Every automation now gets two things written down the day it is built: the number that proves it is worth keeping, and the date we check.

src: DZQZOVhgNvK · asset: template


C-2026-09-12-03_S05.png

C-2026-09-12-03 · S06 · 61.3-69.6s · A_ROLL_MEDIUM · a_roll · ->none

Checklist in the caption

[SHOT] medium

Max, same take, closing line to camera.



The checklist is in the caption. Find the one that fired least last month, then keep it or kill it.

src: DMQNNHSJRF · asset: template

C-2026-09-12-03_S06.png

C-2026-09-12-03 · S00 · 0.0-7.5s · SCREEN_RECORDING · demo · ->none

Never failed. Switched off.

...

[SCREEN]
Full-frame screen recording of the step running: a green pass, then another, then the hand switching it off. No presenter in frame.

This step has never failed. We turned it off this week. Here is the number that decided it.

src: Dbs23d3KUDd · asset: template

S000.0-7.5s

C-2026-09-12-03 · S01 · 7.5-17.5s · SCREEN_RECORDING · demo · ->none

Still running. Nobody asks why.

...

[SCREEN]
The same unbroken take continues: the step's log scrolling on, pass after pass, unremarkable.

You have one of these. A script, a mailbox rule, something set up months ago. It still runs. Nobody has asked what it changes.

src: DZQZOVhgNvK · asset: template

S017.5-17.5s

C-2026-09-12-03 · S02 · 17.5-27.9s · UI_DEMO · demo · ->ui_zoom

264 videos. Zero crashes.

...

[SCREEN]
Slow push into the step's own output: the spoken beats on one side of the timeline, the picture changes on the other.

Ours matched what is said in a video to where the picture changes. It ran on two hundred and sixty-four videos. It never crashed once.

src: - · asset: template

S0217.5-27.9s

C-2026-09-12-03 · S03 · 27.9-34.6s · DATA_VISUAL · demo · ->overlay_reveal

55 matches out of 1,527

...

[SCREEN]
Full-frame: 1 527 small marks, 55 of them lit. The lit ones are hard to find.

Out of one thousand five hundred and twenty-seven places it could have matched, it found fifty-five.

src: Dc_TXiCttIc · asset: template

S0327.9-34.6s

C-2026-09-12-03 · S04 · 34.6-48.8s · DATA_VISUAL · demo · ->text_punch_in

Matched or not: no difference found

...

[SCREEN]
Two columns of results side by side, matched and unmatched, with the same numbers at the bottom of each.

Two hundred and fifteen videos gave it nothing at all. Some of that is our own measurement being coarse. And the videos it matched came out no different from the ones it did not.

src: Dc_TXiCttIc · asset: template

S0434.6-48.8s

C-2026-09-12-03 · S05 · 48.8-61.3s · A_ROLL_MEDIUM · a_roll · ->hard_cut

A number and a date

[SHOT] medium
Max, medium shot, plain wall, first appearance in the reel.



So it is off. Every automation now gets two things written down the day it is built: the number that proves it is worth keeping, and the date we check.

src: DZQZOVhgNvK · asset: template

S0548.8-61.3s

C-2026-09-12-03 · S06 · 61.3-69.6s · A_ROLL_MEDIUM · a_roll · ->none

Checklist in the caption

[SHOT] medium
Max, same take, closing line to camera.





The checklist is in the caption. Find the one that fired least last month, then keep it or kill it.

src: DMQMNHSJRF - asset: template

S06 61.3-69.6s

Contact sheet

Editing

- **Cut timing:** One hard cut in the whole reel, at 48.8 s, where the screen hands over to the presenter. Everything before it is screen: two full-frame captures of the step itself (running, then its own output) followed by two full-frame graphics built from the same run, changed inside the frame with a zoom, an overlay reveal and a punch-in rather than by cutting to a presenter. Everything after it is a single continuous face take.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 600 at 40px, two lines maximum, bottom-anchored in Instagram zone A/B; one persistent top banner held for the whole reel.
- **Emphasis:** The three numbers - 264, 1 527, 55 - are on screen in the same words they are spoken, and nothing else is. The null is never written as an equals sign.
- **Zooms:** One push into the step's own output at 27.9 s. Nowhere else.
- **Motion:** Linear easing only, no bounce, per design/rules/motion.md.
- **Assets required:**
 - Continuous screen capture of the alignment step running green and then being switched off (S00-S01)
 - Capture of the step's own output showing spoken beats against picture changes (S02)
 - Chart of 1 527 candidate boundaries with 55 lit (S03)
 - Two-column comparison of matched vs unmatched videos on the same outcome figures (S04)
 - Max A-roll, medium shot, one take
 - One-page retirement checklist published in the caption, not gated

Claims

CLAIM	ST AT E	SOURCE
A step we built matched the spoken beats of a video to the points where its picture changes, and it ran on all 264 parsed videos without failing.	OB SE RV ED	Insight I-21 in data/analysis/insights.json (n_codes_alignment_done = 264) and reports/data/alignment.json (n_beats_with_scene_ids 1 527 of 1 527 beats, 264 of 264 codes returned a result, so nothing was skipped); alignment stage described in engine/SPEC.md section 6.5. There is no jobs row for an ALIGNMENT stage, so 'never failed' rests on complete coverage of the 264 codes rather than on a job-state trace.
It found 55 matches out of 1 527 possible boundaries, a hit rate of 3.6 %.	OB SE RV ED	Insight I-21: semantic_boundary_cut_rate 0.036, n_beat_boundaries_checked 1 527; reports/analysis/03-visual-patterns.md, Transitions section.
215 of the 264 videos produced no match at all.	OB SE RV ED	Insight I-21: codes_with_any_aligned_boundary = 49 of 264.
Where it did match we could measure no difference against the videos where it did not - a null result, not a demonstrated equality.	OB SE RV ED	Insight I-21: robust_z 1.445 vs 1.251 (p=0.603), share rate 0.0168 vs 0.0149 (p=0.276), save rate p=0.636, view_lift p=0.373. I-21's own confidence is INSUFFICIENT and its interpretation reads 'there is no evidence that cutting on meaning helps, and no evidence it does not' - the script is worded to match that, not to claim the two groups are the same.
Part of the miss rate is our own measurement: the corpus carries nine fixed frames per video, so a picture change between two samples is invisible.	OB SE RV ED	reports/analysis/03-visual-patterns.md, 'The limitation that shapes every number here'; confidence on the alignment null is recorded as INSUFFICIENT for exactly this reason. Stated aloud in the script so the null is not oversold.
The alignment step is being removed from the weekly run, and every automation will get a stated threshold and a review date at the moment it is built.	PLA NN ED	Decision taken 2026-09-12 on insight I-21's implication; not yet reflected in the shipped run.
Retiring the step makes the weekly run faster, cheaper or better.	TO_ ME AS UR E	No before/after run timed; the reel makes no time or cost claim.
A finer frame sampling would leave the retirement decision unchanged.	MI SSI NG	We have not re-run the alignment against scene-v1 detector frames. Flagged here so the reel is not read as proving the step could never work.

Traceability

- **Insights:** I-21, I-20, I-22, I-08, I-09, I-14, I-18
- **Hypothesis:** H-18
- **Evidence summary:** We built an alignment step that matched spoken beats to visual scene boundaries. It ran on all 264 parsed videos and never failed. It produced 55 aligned boundaries out of 1 527 (3.6 %), 215 of the 264 videos produced none at all, and the videos where it did fire performed no differently on any metric (robust_z 1.445 vs 1.251, $p=0.603$; share $p=0.276$; save $p=0.636$). Part of that miss rate is our own nine-frames-per-video sampling, which the reel says out loud. The four references contribute grammar only: a subtractive opening on the working artefact (Dbs2JdJKUdd, robust_z 3.28, 42.6 saves/1k), conditions-concession-conclusion ordering (Dc_TX1CttIc, ungated, 31.7 shares/1k), a long unbroken builder take to invert (DZQZOVhgNvK, 50.2 saves/1k) and an ungated spoken save ask (DMQMNNHSJRF, 35.1 shares/1k).
- **Original synthesis:** The corpus contains no reel in which an author removes something they built: warning_dont is the smallest solution type at 6 of 266 and every one warns about somebody else's tool. The retired step, the hit rate and the null are ours, dated, and reproducible from data/analysis/insights.json and the beats/scenes tables. The retirement rule - a number and a date written down on the day a thing is built - is ours and is published in the caption rather than gated, because what gets saved is a deliverable, not an explanation.
- **Why better than generic:** A generic AI script would argue that 'you should not automate everything' from a standing start and offer a checklist nobody tested. This one argues against automation from inside a build, shows the step succeeding before it shows the number, concedes the part of the miss rate that is our own instrument, and ends on a rule with two concrete fields. It is also the only shape available that cannot be copied back by the 62 demo walkthroughs in the corpus, because none of them has ever removed anything.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] the proof beat read 'the videos it matched performed exactly like the ones it did not' - that asserts an equality from a null (medians 1.445 vs 1.251, $p=0.603$) on an insight whose own confidence is INSUFFICIENT -> rewrote the line as 'came out no different from the ones it did not', rewrote claims[3] to say null rather than same, and quoted I-21's interpretation in the source. Script version 2; total_s 69.2 -> 69.6, storyboard retimed from S04.
 - [visual] editing.cut_timing described everything before the hand-over as 'a single continuous screen take', but the card's own assets_required lists four separate screen assets (the step running, the step's output, a 1 527-mark chart, a two-column comparison) -> rewritten to name two captures plus two graphics changed inside the frame, with the one hard cut kept.
 - [evidence] claims[0]'s 'never failed' cited only I-21's $n=264$. There is no ALIGNMENT row in the jobs table -> source now also cites reports/data/alignment.json (264 of 264 codes, 1 527 of 1 527 beats carry scene ids) and says the claim rests on complete coverage, not on a job-state trace.
 - [visual] S04's overlay 'Matched or not, same result' repeated the equality error on screen -> 'Matched or not: no difference found'.

- [visual] the opening frame carried transition `hard_cut`, which reads as a cut into frame one -> set to none, consistent with the reel's own 'one hard cut' description.
- [evidence] verified: 264 codes, 1 527 boundaries checked, 55 aligned (3.60 %), 215 codes with none - all reproduce exactly from reports/data/alignment.json and I-21. references `save_rate -1.0` on `Dc_TX1CttIc` and `DMQMNNHSJRF` is a genuine NULL in `video_features` (no save count returned), and both reason fields say so -> left, sentinel is disclosed.
- [hook] 18 words / 7.5 s sits just under general-conclusions section G rule 3 (20-30 words) -> left: the three sentences are the whole payoff and padding them weakens the turn.
- [pacing] left as is: the CTA carries a caption pointer and an action against rule 14's one ask. `free_resource` is the ask; 'keep it or kill it' is the filter's `next_action` and cannot be dropped.

Production state

- **Status:** REVIEWED
- **EDL path:** `cards/edl/C-2026-09-12-03.json`
- **Remotion validation:** `{"card_id": "C-2026-09-12-03", "checked_at": "2026-09-12T08:53:25Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2088, "edl": "cards/edl/C-2026-09-12-03.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}`

C-2026-09-12-04 — Cost per run, not price per month

Strategy

Format: M2 Radar

Concept. We open our own ten-day bill on screen - seventeen dollars and seventy-eight cents, eight hundred and eighty-nine paid units at two cents each - and do the one piece of arithmetic a pricing page never does: divide the bill by the number of times the task actually ran. Then we volunteer what the bill leaves out, including the sixteen cents we paid twice when a run broke halfway.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company responsible for one repeated process. Specifically the person who approves and renews software spend.

Objective. Get one viewer to count how many times their repeated task actually ran last month and divide the bill by it before the next renewal, so they walk into that conversation with one number instead of a plan name.

Positioning. builder

Pain. cost_money

Promise. One number you can defend to whoever signs, worked out on our own bill in front of you.

Rationale. cost_money is the largest pain cell in the analysed corpus (n=48, 33 creators) and money-not-spent framing sits above the corpus median on save rate (0.0306 against 0.0250, p=0.007), while money-earned framing costs reach (view_lift 1.19 against 2.12, I-25) - this reel is on the profitable side of that line. Nobody in the 266 analysed reels shows their own bill: 84 carry a money-saving desire (01-general-conclusions §B) and 71 frame the outcome as money (I-25), and in both cells the money is somebody else's price rather than the author's own ledger.

F I L T A N S W E R
E R

Processes	Approving or renewing a paid tool that runs one repeated task in the business - the weekly report, the batch of quotes, the inbox triage - where the invoice arrives monthly and nobody in the company can say whether it is expensive.
Friction	Pricing pages quote per seat or per month and the work happens in runs, so the person who signs has no unit to compare against. The bill grows with use and the growth stays invisible until the invoice lands.
AI boundary	The machine performs the run and prints what it will cost before it starts. The human sets the cap, reads the ledger and decides when the price stops being worth the task; no tool is allowed to spend without an explicit confirmation.
Next action	Take one repeated task, count how many times it actually ran last month, divide the bill by that number, and only then renew.

Hooks

1. **[result_first]** (statement) — **USED IN SCRIPT** This is our whole bill for ten days. Seventeen dollars, seventy-eight cents. The number that matters is not on it.
1. **[bold_claim]** (statement) Our ten-day bill came to seventeen dollars and seventy-eight cents, and that figure tells you nothing about whether the tool is expensive.
1. **[question]** (question) How much does your AI tool cost? If the answer has the word month in it, you cannot answer the question.

Reference reels

REEL	CR EA TO R	FU NC TI ON	P L A Y S	CR EA TO R	VI EW LI FT	SH AR E RA TE	SA VE RA TE	REASON
Dc_TX1CttIc(https://www.instagram.com/reel/Dc_TX1CttIc/)	v.i. s.h. ai	ex pl an ati on	5 6 2, 0 9 4	22, 421	24 .1x	3.2 %		SAVE RATE NOT MEASURED: the provider returned no save count for this reel in any snapshot (reels.save and video_features.save are NULL), so save_rate is null rather than a number; only its share rate can be read. Ungated (cta_type none), so its 31.7 shares per 1 000 can be taken at face value - one of only seven ungated rows in the 40-reel pool. It is also the only reel in the pool whose closing verdict is a cost verdict rather than a capability verdict, and it arrives after the evidence rather than before it. The agenda is announced in the first nine seconds and every case is labelled aloud, which is the ordering a four-step arithmetic reel needs.
DcUNL-Nvy6J9(https://www.instagram.com/reel/DcUNL-Nvy6J9/)	be nki mb all. ai	pa in 8 7 8	8 2, 8	3,5 68	22 .2x	2.2 %	4.3 %	The cleanest statement of metered-pricing pain in the pool, said in money the speaker has actually paid: a named range, a first-person amount, and the fact that the bill grows with success, all inside nine seconds. 42.9 saves per 1 000 on a 22.2x lift. The contact sheet shows the move worth copying - the pricing page fills the frame, the figure itself is highlighted, and the caption chip is pinned directly beside the price it refers to rather than at the bottom of the frame.
DbVq4mwz38L(https://www.instagram.com/reel/DbVq4mwz38L/)	hey ste vet an	pr oo f 9, 9 7 4	4 3 9, 9 7 4	13, 009	32 .8x	2.8 %	5.4 %	The pool's best objection handling: fifteen continuous seconds volunteering every catch before asking for anything, placed at 32-48s, exactly where a viewer doing the arithmetic in their head would object. 32.8x lift, 27.9 shares and 54.0 saves per 1 000. Its contact sheet also carries the single-big-figure treatment we use for the weekly line - one large money number set beside the artefact with a two-word label under it, rather than a table.
DaGTe-JNN1GR(https://www.instagram.com/reel/DaGTe-JNN1GR/)	aiw ithr em y	ho ok 8, 2 2 6	2 9 8, 2 6	28, 959	9. 3x	3.0 %	3.6 %	A money number inside the first four seconds with the mechanism landing by twenty-five. The pool's clearest example of money-not-spent framing, which the corpus rewards, against money-earned framing, which costs reach (median view_lift 1.19 against 2.12). 30.2 shares per 1 000 on a 9.3x lift over a 28 959 median. The contact sheet shows the number stated three ways at once: a persistent title plate, the figure calculated on a real device held into frame, and a large overlay on the same beat.

Dc_TX1CttIc (v.i.s.h.ai, explanation)

- Transcript (Dc_TX1CttIc-B6): “So here is my verdict. Very good model for sure, but the API cost has shot up so high that any business trying to build a real product on Astra is walking into a train wreck because capability alone does not save you from that math.”
- Frame (Dc_TX1CttIc-S02): the verdict returns to a full-frame face after a 62-second unbroken screen run, so the interpretation is delivered where the evidence is not
- Transformed: Borrowed: the announce-the-agenda-then-earn-the-verdict ordering, and the decision to close on an economic limitation rather than a recommendation. Changed: the verdict is about one

repeated task in a small business rather than a model's capability ceiling, and our cost numbers come from our own ledger rather than a vendor's price list. We also add the CTA this reel lacks entirely.

DcUNINvy6J9 (benkimball.ai, pain)

- Transcript (DcUNINvy6J9-B2): “Most people are paying anywhere from 15 to 60 bucks a month for many chat. I even paid 200 bucks last month. And the crazy prize is the price will keep climbing the more views you get.”
- Frame (DcUNINvy6J9-S03): full-frame pricing page with the monthly figure highlighted and the caption chip pinned to that exact cell
- Transformed: Borrowed: the three-part pain shape - named range, first-person amount, the bill scales with your own success - and the caption chip pinned to the number it names. Changed: we say it once, in the first quarter, as a named process cost rather than as an argument for a free alternative, and we never name a competitor product. Its topic, free access and bypassing payment, is an excluded genre for us and is not borrowed.

DbVq4mwz38L (heystevetan, proof)

- Transcript (DbVq4mwz38L-B3): “The repo itself is free. You just cover hosting. Since a background worker and a couple of databases need to stay running. So figure a few dollars a month.”
- Frame (DbVq4mwz38L-S03): one large money figure ('\$39 / Bill amount') set against the artefact instead of a table of numbers
- Transformed: Borrowed: the placement and function of a volunteered-caveat beat immediately after the cost claim, and the one-big-figure treatment. Changed: our caveat is about our own spend and our own unmeasured overheads, not about somebody else's repository, and it is delivered over our own ledger on screen rather than over a text card. That reel never demonstrates its tool working; ours shows the artefact throughout.

DaGTeJNN1GR (aiwithremy, hook)

- Transcript (DaGTeJNN1GR-B0): “All right, so Claude just saved us \$2,800 on our Airbnb booking.”
- Frame (DaGTeJNN1GR-S00): the money figure appears simultaneously as a persistent title plate and as a real calculation on a device held into frame
- Transformed: Borrowed: a measured amount in the opening sentence, stated as money not spent, and shown on a real screen in the same second it is spoken. Changed: the amount is our documented total spend rather than a saving, because POSITIONING.md §6 forbids a savings claim; the subject is a recurring business cost, not a holiday booking; and we drop its CTA, which sells a skill that did not exist at publication.

Script

This is our whole bill for ten days. Seventeen dollars, seventy-eight cents. The number that matters is not on it. Every tool you approve is priced by the month or by the seat. Your work does not happen by the month. It happens in runs. So divide the bill by the number of times the task actually ran. Ours ran eight hundred and eighty-nine times. Two cents a run. Here is the line that repeats every week. Last week, a hundred and thirty accounts, one page each. Two dollars sixty. That is the number to defend, not the plan name. What the bill leaves out: our own machine time, and sixteen cents we paid twice because a run broke halfway and eight answers had left the cache. Two cents a run is nothing. Two cents a run, six hundred times a day, is another conversation. Price the run, then decide. Before the next renewal, count how many times that task ran last month, then divide. Send this to whoever signs.

170 words / 68.0s (2.5 wps target) — tone: flat, arithmetical, unembarrassed about a small bill;
emotional effect: the calm of having one number to say out loud in a meeting instead of an opinion

ROLE	WORDS	SECONDS	CUMULATIVE
hook	20	8.0	8.0s
problem	25	10.0	18.0s
explanation	24	9.6	27.6s
proof	31	12.4	40.0s
objection	27	10.8	50.8s
payoff	23	9.2	60.0s
cta	20	8.0	68.0s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME ELEMENT TYPE	LOCATION YOUTUBE	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0.0-8.0s	This is our whole bill for ten days. Seventeen dollars, seventy-eight cents. The number t...	hook	SCREENSHOT	desktop	Full-frame screenshot of our own spend ledger: ten rows, the unit-price column and the total both legible. The total is boxed. No presenter in frame.	Ten days. Our whole bill.	none	DaGTeJNN1GR (the money figure is spoken and shown on a real screen in the same second)		template
8.0-18.0s	Every tool you approve is priced by the month or by the seat. Your work does not happen b...	promotional	A_ROLL_MEDIUM	roller	Michael, medium shot, one location, banner plate held above him.	Priced by month. Paid by run.	hardcut	DcUNINvy6J9 (the metered-pricing pain is stated once, on the face, as a cost the speaker has actually paid)		template
18.0-27.6s	So divide the bill by the number of times the task actually ran. Ours ran eight hundred a...	explanatory	A_ROLL_MEDIUM	roller	Same take continuing. The division punches in beside him as he says it: 17.78, then 889, then 0.02.	17.78 divided by 889	text punch-in	Dc_TX1CttIc (the arithmetic is announced on the face before it is earned on the screen)		template
27.6-40.0s	Here is the line that repeats every week. Last week, a hundred and thirty accounts, one p...	proof of concept	SCREENROLLING	desktop	Full-frame recording of the same ledger scrolled to the weekly collection line, with a runs-per-week column now added. The two-dollars-sixty cell is the only thing highlighted.	\$2.60. Every week. One run.	hardcut	DcUNINvy6J9 (caption chip pinned to the exact cell rather than to the bottom of the frame); DbVq4mwz38L (one large money figure against the artefact instead of a table of numbers)		template
40.0-50.8s	What the bill leaves out: our own machine time, and sixteen cents we paid twice because a...	object in motion	SCREENSHOT	desktop	Same ledger, scrolled to the two adjacent rows dated 3 September - the interrupted run and the recovery - both highlighted, with the note on the second row legible.	Paid twice: sixteen cents	ui_zoom	DbVq4mwz38L (the caveat is volunteered immediately after the cost claim, before anything is asked for)		template
50.8-60.0s	Two cents a run is nothing. Two cents a run, six hundred times a day, is another conversa...	pay off	A_ROLL_MEDIUM	roller	Michael back on camera, same framing and light. A small mono label reading 'example, not a measurement' sits under the six-hundred figure for as long as it is on screen.	Price the run, then decide	hardcut	Dc_TX1CttIc (the verdict returns to the face and is economic rather than a recommendation)		template
60.0-68.0s	Before the next renewal, count how many times that task ran last month, then divide. Send...	cta	A_ROLL_MEDIUM	roller	Same take, no reframe. Michael holds still for the closing line.	Count the runs. Then renew.	none	—		template

C-2026-09-12-04 · S00 · 0.0-8.0s · SCREENSHOT · demo · ->none

Ten days. Our whole bill.

...

[SCREEN]

Full-frame screenshot of our own spend ledger: ten rows, the unit-price column and the total both legible. The total is boxed. No presenter in frame.

This is our whole bill for ten days. Seventeen dollars, seventy-eight cents. The number that matters is not on it.

src: DaGTeJNN1GR · asset: template

C-2026-09-12-04_S00.png

C-2026-09-12-04 · S01 · 8.0-18.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Priced by month. Paid by run.

[SHOT] medium

Michael, medium shot, one location, banner plate held above him.



Every tool you approve is priced by the month or by the seat. Your work does not happen by the month. It happens in runs.

src: DcUN1Nvy6J9 · asset: template


C-2026-09-12-04_S01.png

C-2026-09-12-04 · S02 · 18.0-27.6s · A_ROLL_MEDIUM · a_roll · ->text_punch_in

17.78 divided by 889

[SHOT] medium

Same take continuing. The division punches in beside him as he says it: 17.78, then 889, then 0.02.



So divide the bill by the number of times the task actually ran. Ours ran eight hundred and eighty-nine times. Two cents a run.

src: Dc_TX1CttIc · asset: template

C-2026-09-12-04_S02.png

C-2026-09-12-04 · S03 · 27.6-40.0s · SCREEN_RECORDING · demo · ->hard_cut

\$2.60. Every week. One run.

...

[SCREEN]

Full-frame recording of the same ledger scrolled to the weekly collection line, with a runs-per-week column now added. The two-dollars-sixty cell is the only thing highlighted.

Here is the line that repeats every week. Last week, a hundred and thirty accounts, one page each. Two dollars sixty. That is the number to defend, not the plan name.

src: DcUN1Nvy6J9, DbVq4mwz38L · asset: template

C-2026-09-12-04_S03.png

C-2026-09-12-04 · S04 · 40.0-50.8s · SCREENSHOT · demo · ->u3_zoom

Paid twice: sixteen cents

...

[SCREEN]

Same ledger, scrolled to the two adjacent rows dated 3 September - the interrupted run and the recovery - both highlighted, with the note on the second row legible.

What the bill leaves out: our own machine time, and sixteen cents we paid twice because a run broke halfway and eight answers had left the cache.

src: DbVq4mwz38L · asset: template

C-2026-09-12-04_S04.png

C-2026-09-12-04 · S05 · 50.8-60.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Price the run, then decide

[SHOT] medium

Michael back on camera, same framing and light. A small mono label reading 'example, not a measurement' sits under the six-hundred figure for as long as it is on screen.



Two cents a run is nothing. Two cents a run, six hundred times a day, is another conversation. Price the run, then decide.

src: Dc_TX1CttIc · asset: template


C-2026-09-12-04_S05.png

C-2026-09-12-04 · S06 · 60.0-68.0s · A_ROLL_MEDIUM · a_roll · ->none

Count the runs. Then renew.

[SHOT] medium

Same take, no reframe. Michael holds still for the closing line.



Before the next renewal, count how many times that task ran last month, then divide. Send this to whoever signs.

src: - · asset: template

C-2026-09-12-04_S06.png

C-2026-09-12-04 · S00 · 0.0-8.0s · SCREENSHOT · demo · ->none

Ten days. Our whole bill.

...

[SCREEN]
Full-frame screenshot of our own spend ledger: ten rows, the unit-price column and the total both legible. The total is boxed. No presenter in frame.

This is our whole bill for ten days. Seventeen dollars, seventy-eight cents. The number that matters is not on it.

src: DaGTcJNN1GR · asset: template

S00 0.0-8.0s

C-2026-09-12-04 · S01 · 8.0-18.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Priced by month. Paid by run.

[SHOT] medium
Michael, medium shot, one location, banner plate held above him.

Every tool you approve is priced by the month or by the seat. Your work does not happen by the month. It happens in runs.

src: DcUN1Nvy6J9 · asset: template

S01 8.0-18.0s

C-2026-09-12-04 · S02 · 18.0-27.6s · A_ROLL_MEDIUM · a_roll · ->text_punch_in

17.78 divided by 889

[SHOT] medium
Same take continuing. The division punches in beside him as he says it: 17.78, then 889, then 0.02.

So divide the bill by the number of times the task actually ran. Ours ran eight hundred and eighty-nine times. Two cents a run.

src: Dc_TX1CttIc · asset: template

S02 18.0-27.6s

C-2026-09-12-04 · S03 · 27.6-40.0s · SCREEN_RECORDING · demo · ->hard_cut

\$2.60. Every week. One run.

...

[SCREEN]
Full-frame recording of the same ledger scrolled to the weekly collection line, with a runs-per-week column now added. The two-dollars-sixty cell is the only thing highlighted.

Here is the line that repeats every week. Last week, a hundred and thirty accounts, one page each. Two dollars sixty. That is the number to defend, not the plan name.

src: DcUN1Nvy6J9, DbVq4mz38L · asset: template

S03 27.6-40.0s

C-2026-09-12-04 · S04 · 40.0-50.8s · SCREENSHOT · demo · ->ui_zoom

Paid twice: sixteen cents

...

[SCREEN]
Same ledger, scrolled to the two adjacent rows dated 3 September - the interrupted run and the recovery - both highlighted, with the note on the second row legible.

What the bill leaves out: our own machine time, and sixteen cents we paid twice because a run broke halfway and eight answers had left the cache.

src: DbVq4mz38L · asset: template

S04 40.0-50.8s

C-2026-09-12-04 · S05 · 50.8-60.0s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Price the run, then decide

[SHOT] medium
Michael back on camera, same framing and light. A small mono label reading 'example, not a measurement' sits under the six-hundred figure for as long as it is on screen.

Two cents a run is nothing. Two cents a run, six hundred times a day, is another conversation. Price the run, then decide.

src: Dc_TX1CttIc · asset: template

S05 50.8-60.0s

C-2026-09-12-04 · S06 · 60.0-68.0s · A_ROLL_MEDIUM · a_roll · ->none

Count the runs. Then renew.

[SHOT] medium
Same take, no reframe. Michael holds still for the closing line.



Before the next renewal, count how many times that task ran last month, then divide. Send this to whoever signs.

src: - · asset: template

S06 60.0-68.0s

Contact sheet

Editing

- **Cut timing:** Seven sections, four visible cuts, each of them a full-frame screen insert replacing the presenter or the presenter returning. The presenter material is one continuous take; there is no speaker change in this reel, and a speaker change is the only cut PRODUCTION.md allows.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 40px, Paper on an Ink plate with 24px padding, bottom safe zone, two lines maximum. One persistent Ink banner plate in Saira SemiCondensed Black at the top for the whole reel. Every money figure is set in IBM Plex Mono Medium as a label, per design/rules/typography.md. No emoji on screen.
- **Emphasis:** Only one cell is highlighted at a time on any ledger screen, in Blueprint - nothing here is broken, it is mispriced. Oxide appears exactly once, on the duplicated pair of rows.
- **Zooms:** One ui_zoom onto the weekly cell, one onto the duplicated pair. No zoom on any A-roll scene.
- **Motion:** Linear easing only, per design/rules/motion.md. No bounce, no gradients, no rounded corners, no centred text.
- **Assets required:**
 - Screenshot of our own spend ledger, ten rows, unit price and total visible (real, from reports/data/spend.json)
 - Screen recording of the same ledger with a runs-per-week column added and the weekly line highlighted
 - Screenshot of the two 3 September rows - interrupted run and recovery - with the note legible
 - Presenter A-roll, one continuous take, one location, banner plate fixed
 - A small mono 'example, not a measurement' label for the six-hundred-runs figure

Claims

CLAIM	STATE	SOURCE
Our own spend ledger records ten entries between 1 and 10 September 2026: 889 paid units at \$0.02 each, \$17.78 in total.	OBSERVED	reports/data/spend.json - n_entries 10, total_usd 17.78; the ten entries sum to 889 units at a recorded unit price of 0.02.
\$17.78 divided by 889 runs is two cents a run.	OBSERVED	Arithmetic on reports/data/spend.json; \$0.02 is also the unit price recorded on every entry.
The most recent weekly collection run cost \$2.60: 130 accounts, one page each, at \$0.02 a unit. The same job cost \$2.12 the week before, when the roster held 106 accounts.	OBSERVED	reports/data/spend.json entries 8 and 10 (2026-09-07: 106 units, \$2.12; 2026-09-10: 130 units, \$2.60); reports/audit/03-server-runtime-hiker.md - one page per account, one unit per account.
Sixteen cents of the total was paid twice: one run broke partway through and eight of the answers were no longer in the cache.	OBSERVED	reports/data/spend.json entries 6 and 7, both dated 2026-09-03; the note on entry 7 records the recovery after the interruption.
Our own spec requires any module that costs money to print an estimate and wait for an explicit confirmation before spending.	OBSERVED	engine/SPEC.md §0.8.
That every module in our pipeline already complies with that rule.	MISING	The rule binds new modules. The weekly collection job predates it and has not been audited against it; not asserted on screen.
That any particular business runs a task six hundred times a day.	MISING	An illustration of the arithmetic only, labelled on screen as an example rather than a measurement.
That pricing by the run saves money, or that this arithmetic reduced anyone's bill.	MISING	No saving is claimed anywhere in the script; POSITIONING.md §6 forbids a savings claim without a client baseline.
That the \$0.02 unit price will hold.	TO_MEASURE	It is one vendor's price on 10 September 2026 and it will move. We re-read the ledger at the next renewal; the arithmetic, not the price, is the message.
How many times the viewer's own task runs, and therefore what their per-run cost is.	MISING	Unknown to us by definition. The reel asks the viewer to count it rather than estimating it for them.

Traceability

- **Insights:** I-09, I-08, I-13, I-25, I-15, I-29, I-05
- **Hypothesis:** H-01
- **Evidence summary:** Everything on screen is our own ledger: ten entries between 1 and 10 September 2026, 889 paid units at \$0.02, \$17.78 total, of which the repeating weekly line is 130

accounts at one page each for \$2.60, and sixteen cents were paid twice after a run broke partway through. cost_money is the largest pain cell in the analysed corpus (n=48, 33 creators) and money-not-spent framing sits above the corpus median on save rate (0.0306 against 0.0250, p=0.007) while money-earned framing costs reach (view_lift 1.19 against 2.12, I-25). The four references contribute structure only: DaGTeJNN1GR (9.3x lift, 30.2 shares/1k) the measured amount in the opening sentence shown on a real screen, DcUNINvy6J9 (22.2x, 42.9 saves/1k) the three-part metered-pricing pain with the caption chip pinned to the exact cell, DbVq4mwz38L (32.8x, 54.0 saves/1k) the volunteered caveat straight after the cost claim, Dc_TX1CttIc (24.1x, 31.7 shares/1k, ungated) the announce-then-earn-the-verdict ordering with an economic close. Opening on the screen rather than the face is worth roughly half again the forwarding rate (I-09: 0.0213 against 0.0140, creator-normalised +0.285, p=0.00025).

- **Original synthesis:** Nobody in the 266 analysed reels shows their own bill; 84 carry a money-saving desire (\$B) and 71 frame the outcome as money (I-25), always somebody else's price rather than their own ledger. The itemised ledger, the single division, the repeating weekly line item and the volunteered duplicate charge are ours and reproducible line by line from reports/data/spend.json.
- **Why better than generic:** A generic AI script would compare three tools' pricing pages and warn about hidden costs. This one performs a single division in front of the viewer on a real bill, isolates the repeating line item that a renewal conversation actually turns on, volunteers the sixteen cents we wasted, and refuses the savings claim the format invites - which POSITIONING.md §6 forbids and which the reach data (I-25) says costs distribution anyway.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [schema] Reference Dc_TX1CttIc carried save_rate 0.0 as a placeholder for a count the provider never returned (reels.save and video_features.save are both NULL in every snapshot). Per the orchestrator's update the validator now accepts null → save_rate set to null and the 'not measured' note kept at the head of the reason.
 - [evidence] The whole ledger reproduces from reports/data/spend.json and from the spend table in data/radar.db: 10 entries dated 1-10 September 2026, 889 units, \$0.02 each, \$17.78 total, and 889 x \$0.02 = \$17.78 exactly → left.
 - [evidence] The sixteen cents paid twice is entries 6 and 7, both 2026-09-03: 52 units interrupted, then 8 units re-fetched because the answers had left the cache → left.
 - [evidence] 'The repeating weekly collection run costs \$2.60' was stated as a fixed weekly figure. The same job cost \$2.12 on 2026-09-07 (106 accounts) and \$2.60 on 2026-09-10 (130 accounts); the figure moves with the roster → claim rewritten to name both runs, and the spoken line changed from 'Here is the line that repeats every week. A hundred and thirty accounts...' to 'Here is the line that repeats every week. Last week, a hundred and thirty accounts...' so the number is dated rather than presented as constant. Script goes to version 2; the section grows from 29 to 31 words and the reel from 67.2 s to 68.0 s, still inside the 50-70 s band.
 - [evidence] engine/SPEC.md §0.8 does require any module that costs money to print an estimate and require --yes → claim left as OBSERVED.
 - [evidence] All four references re-derived from data/radar.db: play, creator_median_play, view_lift, share_rate and save_rate match, and the four function tags match data/analysis/hypothesis_refs.json → left.

- [originality] '84 have a money-shaped desire and every one of them frames it as money not spent on somebody else's product' is a universal the sources do not state; 01-general-conclusions §B classifies 84 reels as money-saving desire and I-25 counts 71 money-framed reels → rewritten in strategy.rationale and traceability.original_synthesis to what the two sources actually say. The load-bearing original point (nobody shows their own bill) is unchanged.
- [positioning] 'Six hundred times a day' is the only unsourced figure spoken; it is labelled on screen as an example and recorded as MISSING in claims → left.
- [positioning] No vendor named, no savings claim (POSITIONING.md §6), no comment gate (RULES.md §1.2), no split screen (I-10), CTA is a single save_share ask → left.
- [visual] Opens on the ledger at 0.0 s (I-09), four screen inserts, presenter take continuous, every overlay ≤ 7 words, no scene under 1.5 s → left.
- [timing] Recomputed after the edit: 170 words, 68.0 s at 2.5 wps, every section = words/wps, storyboard contiguous 0.0-68.0 s → passes.
- [scheduling] This card, C-2026-09-12-06 and C-2026-09-12-10 all open on our own pipeline artefacts. Keep at least a week between them, and note that 06 describes the same Monday as twenty-five named stages while 10 describes it as five steps - the two are different levels of description and should not run in the same week.
- [schema] python3 -m engine.cards_v2 validate → 0 errors, 0 warnings after the edits.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-04.json
- **Remotion validation:** {"card_id": "C-2026-09-12-04", "checked_at": "2026-09-12T08:53:26Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2040, "edl": "cards/edl/C-2026-09-12-04.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-05 — The approval box we designed and never ticked

Strategy

Format: M2 Builds

Concept. We wrote a mandatory human check into our own process eleven days ago, built the field that records it, and then never filled it in - so the control existed on paper and nowhere else.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company who is responsible for one repeated process and has tried ChatGPT but has no repeatable way of using it at work. Specifically anyone who has written an approval or sign-off step into a process document.

Objective. Make the viewer go and count how many of their last fifty records actually carry the sign-off their process document promises, by showing ours turn out to be decorative.

Positioning. builder

Pain. quality_trust

Promise. A two-minute way to find out whether your sign-off step is a control or a preference.

Rationale. Twenty-five of 266 reels prove by authority claim and none admits that the author's own governance did not run. The references supply cadence and framing only: a fixed item shape repeated without variation over a long face take (DbO4zvJR1k8), a reward bookend that makes a long middle tolerable (DZ-wujrTJ0u), an unbroken screen take with live narration (DVaBP3wDAFE), and a forward-to-a-named-role close with its comment gate deleted (DczDlj4qQjI).

FILT ANSWER
ER

Proce ss	The sign-off step inside any repeated process - who checks the work before it leaves the building, and where the fact that they checked it is written down. Quotes before they go to a customer, invoices before they are sent, copy before it is published.
Fricti on	The rule lives in a document, the field lives in the system, and nothing connects them. The work can complete with the box empty, so the box stays empty, and nobody notices until something goes out that should not have.
AI boun dary	Automation may prepare the item and queue it for approval. The record that a person actually checked has to be produced by that person's action - it can never be inferred, defaulted or filled in by the system that did the work.
Next actio n	Open the place your process document says the check is recorded, and count how many of the last fifty records have it filled in.

Hooks

1. **[result_first]** (statement) — **USED IN SCRIPT** We wrote a mandatory check into our own process on the first of September. Here is the box. Empty.

1. **[bold_claim]** (statement) A check that nothing stops for is not a check. Ours stopped nothing for eleven days.

1. **[question]** (question) Where does your process actually record that a person checked?

Reference reels

REEL	CRE ATO R	FUN CTI ON	PL AY S	CREA TOR MEDI AN	VI EW LI FT	SHA RE RAT E	SAV E RAT E	REASON
DbO4zv-JR1k8 (https://www.instagram.com/reel/DbO4zv-JR1k8/)	shru g.ma nny	rhyt hm	20 9, 12 1	14,44 8	13. 5x	5.3 %	8.7 %	A 71-second single-take close-up that never flattens, because every item has the identical shape - the consequence of not doing it, then the exact clause that fixes it. Highest save rate in the corpus (87.4 per 1 000) with no screen in it at all, which proves a control reel can hold on a face if the cadence is strict.
DZ-wujrTJ0u (https://www.instagram.com/reel/DZ-wujrTJ0u/)	kevi nfre mon	a_ro ll	33 4, 70 6	6,358	51. 6x	2.0 %	6.5 %	Second-highest save rate in the pool (64.9 per 1 000) on a 172-second reel, and it survives that length by bookending: the artefact demonstrates itself early and again at the very end, with the instructions in between. Ungated (cta_type follow). Its 57-second opening close-up is the longest single face take in the pool that still works.
DVaBP3w-DAFE (https://www.instagram.com/reel/DVaBP3w-DAFE/)	tenf oldm arc	scre en_ proo f	53 3, 52 8	7,471	70. 4x	2.2 %	3.8 %	Holds one unbroken UI_DEMO state from 0.4 s to 60.3 s - two thirds of the reel - and narrates over it continuously before cutting to the face. It is the pool's proof that a long unbroken screen take with live narration holds attention without a single cut, which is exactly what showing one empty column demands. robust_z 5.0 (clipped).
DczDlj4qQjI (https://www.instagram.com/reel/DczDlj4qQjI/)	liamj ohns ton.a i	cta	17 3, 62 9	13,44 8	11. 9x	2.4 %	5.7 %	Its close carries two asks and only the second is usable: the share-with-a-peer half is the pure form of the question RULES.md §4 makes us answer first - who exactly will forward this, and to whom - and it is why the reel carries 23.6 shares per 1 000 with an explicit share ask at the end.

DbO4zvJR1k8 (shrug.manny, rhythm)

- Transcript (DbO4zvJR1k8-B1): “Number one and most important, if your terms and service doesn't cap your liability, there's no limit to how much somebody can sue you for.”
- Frame (DbO4zvJR1k8-S00): one close-up for the entire reel; a persistent two-line title plate at the top and a plain white quote card dropped in low in the frame only when a specific clause is named
- Transformed: Borrowed the fixed item shape repeated without variation, and the persistent title plate. Changed: we have three items instead of four, the middle two thirds of our reel is a full-frame screen rather than a face because the empty column is the proof, and the register drops the legal-fear framing after the first line.

DZ-wujrTJ0u (kevinfremon, a_roll)

- Transcript (DZ-wujrTJ0u-B0): “I'm going to show you how you can build your very own AI drivers in under three minutes so you can get something like I have here with Trillian.”
- Frame (DZ-wujrTJ0u-S00): 57 seconds of one close-up with a black title plate above the head and a live countdown ring in the corner, so a long take still reads as going somewhere
- Transformed: Borrowed the reward bookend - something concrete at the start and again at the close, so a long middle is tolerable - and the confidence to hold one face take. Changed: our bookend is the

same column shown empty and then shown blocked, not a product demo; we drop the countdown device; and the reel's own conversion target is the creator's product behind a follow-only CTA, which we do not reproduce.

DVaBP3wDAFE (tenfoldmarc, screen_proof)

- Transcript (DVaBP3wDAFE-B2): “If I have an idea I just drop it in here It's going to go in my pipeline automatically and here I can see the difference in terms of engagement.”
- Frame (DVaBP3wDAFE-S00): the screen is filmed off a monitor across the room rather than captured, held unbroken for 60 seconds, with a small persistent badge top-left and sentence captions along the bottom
- Transformed: Borrowed the unbroken screen take with live narration and the single late cut to the face. Changed: the narration is an admission rather than a boast, the screen shows a control that did not run instead of a dashboard that does, we capture the screen cleanly rather than filming a monitor, and the reel's final third - explicit fear pressure and profanity, flagged as such in our own analysis - is replaced by the redesign.

DczDlj4qQjI (liamjohnston.ai, cta)

- Transcript (DczDlj4qQjI-B5): “Just comment Claude and I'll send you the full repo and setup and send this to someone you know who needs to post on LinkedIn more.”
- Frame (DczDlj4qQjI-S03): returns from the full-frame product screenshots to a plain seated close-up for the closing line, captions in a dark box at the seam
- Transformed: Borrowed the forward-to-a-named-role close, specified by role rather than 'send this to a friend'. Changed: the comment gate half is deleted, not softened - the gate moves save rate 2.8x and share rate 1.8x and reach not at all, so dropping it costs no distribution and removes the inflation from our own future numbers - and the role we name is the person who signs off, not a generic contact.

Script

We wrote a mandatory check into our own process on the first of September. Here is the box. Empty. You have one too. Somewhere in your process a document says a person checks this before it goes out. Now find where that check is recorded. Ours is a single column. It has been there since the rule was written. It is filled in on four rows out of two hundred and eighty-two. All four were filled on the same day, the day we wrote the rule. Every card we have drafted since has a row waiting and the box behind it empty. Three items on this week's shortlist should have failed that check. We caught them by reading. One of the three is in the record. Two are not. A check nothing stops for is a preference. Three ways out: make the step block, put it where someone already looks, or delete the rule. We took the first. Open where your process says the check is recorded and count the last fifty. Send this to whoever signs off.

178 words / 68.5s (2.6 wps target) — tone: plain, uncomfortable, and specific - the tone of a post-mortem, not an apology; emotional effect: the urge to go and check your own sign-off step before anyone asks about it

ROLE	WORDS	SECONDS	CUMULATIVE
hook	19	7.3	7.3s
problem	26	10.0	17.3s
explanation	27	10.4	27.7s
proof	30	11.5	39.2s
proof	27	10.4	49.6s
payoff	29	11.2	60.8s
cta	20	7.7	68.5s

Storyboard


TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0.0-5.0s	We wrote a mandatory check into our own process on the first of September.	hook	A_ROLE_UP	a_roll	Michael, close-up, direct to camera, plain wall, no screen yet.	Written 1 September. Never ticked.	none	DZ-wujrTJ0u (same confidence to open on one close-up with the promise on a plate above the head)		template
5.0-7.3s	Here is the box. Empty.	hook	UI_D EMO	demo	Full-frame: the suitability column in our own table, scrolled to the newest rows, every cell in it blank. No presenter.	Empty	hard_cut	DVaBP3wDAFE (the screen takes over the whole frame and the voice keeps going over it, no inset, no split)		template
7.3-17.3s	You have one too. Somewhere in your process a document says a person checks this before i...	problem	SCREENSHOT	demo	The same screen take pulls back from the column to the rule as written in our own process document, the sentence that requires the check highlighted.	The rule, as written	none	DVaBP3wDAFE (the screen takes over the whole frame and the voice keeps going over it, no inset, no split)		template
17.3-27.7s	Ours is a single column. It has been there since the rule was written. It is filled in on...	exploration	UI_D EMO	demo	The same unbroken screen take pans from the document to the table itself: one column, scrolling, blank cell after blank cell.	4 rows filled out of 282	none	DVaBP3wDAFE (one unbroken screen state narrated live for the whole middle of the reel)		template
27.7-39.2s	All four were filled on the same day, the day we wrote the rule. Every card we have draft...	proof	UI_D EMO	demo	The four filled rows, all carrying the same date, then a scroll through the drafted work with the column blank the whole way.	All four on the same day	ui_zoom	DbO4zvJR1k8 (each item stated in the identical shape so a long stretch does not flatten)		template
39.2-49.6s	Three items on this week's shortlist should have failed that check. We caught them by rea...	proof	DATAVISUAL	demo	Three rows side by side: one with a recorded rejection and a reason, two with the cell blank.	Caught by reading, not by the field	overlay_reveal	DbO4zvJR1k8 (the consequence named before the fix, in the same grammar as the previous item)		template
49.6-60.8s	A check nothing stops for is a preference. Three ways out: make the step block, put it wh...	payoff	A_ROLE_UP	a_roll	Michael, back on camera, same close-up as the opening.	Block it, show it, or drop it	hard_cut	DZ-wujrTJ0u (reward bookend - the reel returns to the framing it opened on to close the loop)		template
60.8-68.5s	Open where your process says the check is recorded and count the last fifty. Send this to...	cta	A_ROLE_UP	a_roll	Michael, same take, closing line to camera.	Count your last fifty	none	DczDlj4qQjI (forward-to-a-named-role close with the comment gate removed)		template

C-2026-09-12-05 · S00 · 0.0-5.0s · A_ROLL_CLOSE_UP · a_roll · ->none

Written 1 September. Never ticked.

[SHOT] close-up

Michael, close-up, direct to camera, plain wall, no screen yet.



We wrote a mandatory check into our own process on the first of September.

src: DZ-mujrT30u · asset: template

C-2026-09-12-05_S00.png

C-2026-09-12-05 · S01 · 5.0-7.3s · UI_DEMO · demo · ->hard_cut

Empty

...

[SCREEN]

Full-frame: the suitability column in our own table, scrolled to the newest rows, every cell in it blank. No presenter.

Here is the box. Empty.

src: DVaBP3wDAFE · asset: template

C-2026-09-12-05_S01.png

C-2026-09-12-05 · S02 · 7.3-17.3s · SCREENSHOT · demo · ->none

The rule, as written

...

[SCREEN]

The same screen take pulls back from the column to the rule as written in our own process document, the sentence that requires the check highlighted.

You have one too. Somewhere in your process a document says a person checks this before it goes out. Now find where that check is recorded.

src: DVaBP3wDAFE · asset: template

C-2026-09-12-05_S02.png

C-2026-09-12-05 · S03 · 17.3-27.7s · UI_DEMO · demo · ->none

4 rows filled out of 282

...

[SCREEN]

The same unbroken screen take pans from the document to the table itself: one column, scrolling, blank cell after blank cell.

Ours is a single column. It has been there since the rule was written. It is filled in on four rows out of two hundred and eighty-two.

src: DVaBP3wDAFE · asset: template

C-2026-09-12-05_S03.png

C-2026-09-12-05 · S04 · 27.7-39.2s · UI_DEMO · demo · ->ui_zoom

All four on the same day

...

[SCREEN]

The four filled rows, all carrying the same date, then a scroll through the drafted work with the column blank the whole way.

All four were filled on the same day, the day we wrote the rule. Every card we have drafted since has a row waiting and the box behind it empty.

src: Db04zvJR1k8 · asset: template

C-2026-09-12-05_S04.png

C-2026-09-12-05 · S05 · 39.2-49.6s · DATA_VISUAL · demo · ->overlay_reveal

Caught by reading, not by the field

...

[SCREEN]

Three rows side by side: one with a recorded rejection and a reason, two with the cell blank.

Three items on this week's shortlist should have failed that check. We caught them by reading. One of the three is in the record. Two are not.

src: Db04zvJR1k8 · asset: template


C-2026-09-12-05_S05.png

C-2026-09-12-05 · S06 · 49.6-60.8s · A_ROLL_CLOSE_UP · a_to11 · ->hazd_cut

Block it, show it, or drop it

[SHOT] close-up

Michael, back on camera, same close-up as the opening.



A check nothing stops for is a preference. Three ways out: make the step block, put it where someone already looks, or delete the rule. We took the first.

src: DZ-mujrT30u · asset: template


C-2026-09-12-05_S06.png

C-2026-09-12-05 · S07 · 60.8-68.5s · A_ROLL_CLOSE_UP · a_to11 · ->none

Count your last fifty

[SHOT] close-up

Michael, same take, closing line to camera.



Open where your process says the check is recorded and count the last fifty. Send this to whoever signs off.

src: Dcz0lj4q0jI · asset: template

C-2026-09-12-05_S07.png

C-2026-09-12-05 · S00 · 0.0-5.0s · A_ROLL_CLOSE_UP · a_roll · ->none

Written 1 September. Never ticked.

[SHOT] close-up

Michael, close-up, direct to camera, plain wall, no screen yet.



We wrote a mandatory check into our own process on the first of September.

src: DZ-wujrT3Du · asset: template

S000.0-5.0s

C-2026-09-12-05 · S01 · 5.0-7.3s · UI_DEMO · demo · ->hard_cut

Empty

...

[SCREEN]

Full-frame: the suitability column in our own table, scrolled to the newest rows, every cell in it blank. No presenter.

Here is the box. Empty.

src: DVaBP3wDAFE · asset: template

S015.0-7.3s

C-2026-09-12-05 · S02 · 7.3-17.3s · SCREENSHOT · demo · ->none

The rule, as written

...

[SCREEN]

The same screen take pulls back from the column to the rule as written in our own process document, the sentence that requires the check highlighted.

You have one too. Somewhere in your process a document says a person checks this before it goes out. Now find where that check is recorded.

src: DVaBP3wDAFE · asset: template

S027.3-17.3s

C-2026-09-12-05 · S03 · 17.3-27.7s · UI_DEMO · demo · ->none

4 rows filled out of 282

...

[SCREEN]

The same unbroken screen take pans from the document to the table itself: one column, scrolling, blank cell after blank cell.

Ours is a single column. It has been there since the rule was written. It is filled in on four rows out of two hundred and eighty-two.

src: DVaBP3wDAFE · asset: template

S0317.3-27.7s

C-2026-09-12-05 · S04 · 27.7-39.2s · UI_DEMO · demo · ->ui_zoom

All four on the same day

...

[SCREEN]

The four filled rows, all carrying the same date, then a scroll through the drafted work with the column blank the whole way.

All four were filled on the same day, the day we wrote the rule. Every card we have drafted since has a row waiting and the box behind it empty.

src: Db04zvJR1k8 · asset: template

S0427.7-39.2s

C-2026-09-12-05 · S05 · 39.2-49.6s · DATA_VISUAL · demo · ->overlay_reveal

Caught by reading, not by the field

...

[SCREEN]

Three rows side by side: one with a recorded rejection and a reason, two with the cell blank.

Three items on this week's shortlist should have failed that check. We caught them by reading. One of the three is in the record. Two are not.

src: Db04zvJR1k8 · asset: template

S0539.2-49.6s

C-2026-09-12-05 · S06 · 49.6-60.8s · A_ROLL_CLOSE_UP · a_roll · ->hard_cut

Block it, show it, or drop it

[SHOT] close-up

Michael, back on camera, same close-up as the opening.




C-2026-09-12-05 · S07 · 60.8-68.5s · A_ROLL_CLOSE_UP · a_roll · ->none

Count your last fifty

[SHOT] close-up

Michael, same take, closing line to camera.





Contact sheet

Editing

- **Cut timing:** Two hard cuts: face to screen at 5.0 s, screen back to face at 49.6 s. The screen block between them is one continuous 44.6-second take with a zoom and a reveal inside it; both face blocks are one take each.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 600 at 40px, two lines maximum, bottom-anchored in Instagram zone A/B; one persistent top banner held for the whole reel.
- **Emphasis:** Only three numbers appear on screen - 4, 282, 3 - each in the same words as spoken. No emoji, and no text generated inside any image.
- **Zooms:** One zoom, onto the four identical dates at 27.7 s.
- **Motion:** Linear easing only, no bounce, per design/rules/motion.md.
- **Assets required:**
 - Screen capture of the suitability column scrolled to the newest rows, every cell blank (S01)
 - The same take pulling back to the rule as written in our own process document, with the requiring sentence highlighted (S02)
 - Continuous screen capture of the table with the suitability column scrolling, blank (S03-S04)
 - The four filled rows showing the same date (S04)
 - Three-row comparison of this week's content rejections, one recorded and two not (S05)
 - Michael A-roll, close-up, one location, two takes (opening and closing block)

Claims

CLAIM	STATE	SOURCE
Since 1 September 2026 our own rules have required a suitability check on every candidate reel, after the storyboard and before a card is written.	OBSERVED	RULES.md section 2a, added 1 September 2026. The file's own header still marks it a draft awaiting Misha's edits, so the reel says we wrote the rule, which is what the file records.
The field that records that check is filled in on 4 rows out of 282, and all four were written on 2026-09-01, the day the rule was added.	OBSERVED	Query of deepdives in data/radar.db on 2026-09-12: 282 rows, 4 with suitable set (DUduffIE5w5, DUGghN2E82c, DcWCLtCpaKu, Dca2CyhJP1B), all done_at 2026-09-01, all rejections.
None of the 23 cards drafted since has the check recorded on its source reel, and every one of the 23 has a row in which it could have been recorded.	OBSERVED	Same query joined to the cards table on 2026-09-12: 23 of 23 card source codes have a deepdives row and 0 of them have suitable set. The earlier reading - that three had no row at all - does not reproduce.
Three reels on this week's forty-reel shortlist were excluded on content grounds; one of the three is recorded in the field and two are not.	OBSERVED	reports/analysis/09-reference-selection.md §4 (DUduffIE5w5, DK9pFUjPQcx, Dbd7VDgP9PR); only DUduffIE5w5 has deepdives.suitable set.
The check will be made a blocking step, so a card cannot be written while the field is empty.	PLANNED	Decision taken 2026-09-12; not yet in the shipped pipeline.
Making the step block raises the fill rate and catches something a human reading would have missed.	TOMESURE	No run under the blocking rule yet.
Nothing disqualifying has ever reached publication because the check was working.	MISING	Nothing has been published at all - our_posts is empty - so the clean record is not evidence that the control worked. Flagged so the reel does not imply it.

Traceability

- **Insights:** I-11, I-08, I-09, I-15, I-18, I-14, I-26
- **Hypothesis:** H-13
- **Evidence summary:** Our own rules have required a suitability check on every candidate reel since 1 September 2026. The field that records it is filled on 4 rows out of 282, all four on the day the rule was written, and none of the 23 cards drafted since carries it. Three reels on this week's forty-reel shortlist were excluded on content grounds and only one of the three is in the record. The four references contribute cadence and framing: a strict repeated item shape over one long face take (DbO4zvJR1k8, 87.4 saves/1k, robust_z 3.01), a reward bookend that carries a long middle (DZ-wujrTJ0u, 64.9 saves/1k, ungated), an unbroken 60-second screen take narrated live (DVaBP3wDAFE, robust_z 5.0) and a forward-to-a-named-role close with its gate deleted (DczDlj4qQjI, 23.6 shares/1k).

- **Original synthesis:** Twenty-five of 266 reels prove by authority claim and not one admits that the author's own governance did not run. The rule, the field, the four dated rows and the three content rejections are all ours and queryable from data/radar.db; the fix we chose - make the step block - is stated with the two alternatives we rejected, so the reel is a redesign rather than a confession, and two thirds of the runtime sits on the proof and the redesign rather than on the admission.
- **Why better than generic:** A generic AI script would advise 'always keep a human in the loop' and show a stock approval screen. This one shows that we wrote the human into the loop, built the box, and then produced a count of four out of two hundred and eighty-two - and it names the difference between a control and a preference as whether the work can complete without it. The proof is a single unambiguous number on a real screen, which is the cheapest possible thing to shoot and the only visual feature in the corpus that survives creator normalisation on two independent metrics.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] claims[2] said '20 of the 23 card source codes have a deepdives row' and the script said every drafted card 'has an empty box behind it, or no box at all'. Recomputed on data/radar.db: all 23 card source codes have a deepdives row and none has suitable set -> the false half of the line is gone ('has a row waiting and the box behind it empty') and the claim now says 23 of 23. Script version 2; total_s 68.9 -> 68.5.
 - [visual] the hook says 'Here is the box. Empty.' over a close-up of Michael with no box anywhere on screen, and general-conclusions section G rule 1 wants the artefact first -> the hook is split across two scenes: face for the first sentence, then a full-frame cut to the blank column at 5.0 s, which also puts the proof screen inside the first ten seconds (I-09) instead of at 7.3. Scenes renumbered S00-S07, storyboard retimed, asset paths follow.
 - [visual] with the box now opening the screen block, S02 followed it by cutting back out to a document -> S02 becomes the same take pulling back from the column to the rule, so the whole middle is one continuous screen block between two hard cuts, exactly as editing.cut_timing describes.
 - [evidence] verified against data/radar.db: deepdives has 282 rows, 4 with suitable set (DUduffIE5w5, DUGghN2E82c, DcWCLtCpaKu, Dca2CyhJP1B), all done_at 2026-09-01, all rejections. Of the three reels excluded on content grounds in reports/analysis/09 section 4, only DUduffIE5w5 is in the field. Both figures stand as written.
 - [positioning] claims[0] leans on RULES.md, which is still headed as a draft awaiting Misha's edits -> left, and the source now says so; the reel's claim is that we wrote the rule, not that it was ratified.
 - [pacing] left as is: the CTA carries two asks against section G rule 14. The count-your-last-fifty action is the filter's next_action and the forward is PRODUCTION.md's share question.
 - [hook] 19 words / 7.3 s, one word under rule 3's 20-30 band -> left.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-05.json

- **Remotion validation:** {"card_id": "C-2026-09-12-05", "checked_at": "2026-09-12T08:53:26Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2055, "edl": "cards/edl/C-2026-09-12-05.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-06 — The pause is the feature

Strategy

Format: M2 Radar

Concept. One question replaces a feature comparison: does the tool stop and wait for a person before anything leaves the company. We show the same task run twice, full-frame and one after the other - once set to send, once set to hold - and then count the same gate in our own pipeline, where twenty-five named steps run unattended and one waits for a person who struck out nine of the last twenty-three plans.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company responsible for one repeated process. Specifically the person choosing a tool that will write or send something a customer reads.

Objective. Get one viewer to open the tool they are evaluating, look for the setting that holds output for review, and refuse the purchase for outward-facing work if that setting does not exist.

Positioning. builder

Pain. quality_trust

Promise. One question that replaces a feature comparison, and the exact screen to ask for before you sign.

Rationale. Exactly one reel in the 266-reel analysed corpus states the boundary out loud, and it does so as a throwaway objection beat inside a tool promotion. Making the boundary the subject is unrepresented. The visual shape is the one the data supports: A>SCREEN>A carries median share_rate 0.0213 against 0.0140 (p=0.00076, I-11) and opening on the screen is worth roughly half again the forwarding rate (I-09).

FILT ANSWER
ER

Proces s	Buying or switching a tool that produces outward-facing text - replies, quotes, summaries, posts, anything a customer or a supplier will read with the company's name on it.
Frictio n	Feature lists say what the tool can produce and never say where the person sits, so the approval step is discovered after the tool is already inside the workflow and something has already gone out unread.
AI bound ary	The tool drafts and queues. A named person releases. Nothing addressed to somebody outside the company leaves without that person having seen it on a screen.
Next action	Open the tool you are evaluating, find the setting that holds output for review, and if it does not exist, do not buy it for outward-facing work.

Hooks

1. [demo_first] (statement) — **USED IN SCRIPT** Watch this reply sit here. It is written, it is right, and it goes nowhere until somebody presses the button.

1. **[bold_claim]** (statement) The only feature worth comparing in an AI tool is the one that stops it. Here is what that looks like.
1. **[question]** (question) Where does the person sit in the tool you are about to buy? The feature list will not tell you.

Reference reels

REEL	CREATOR	FULL NAME	PO	CREATOR	VIEWER	SHARE	SAVES	REASON
DbVKVz0y8xo (https://www.instagram.com/reel/DbVKVz0y8xo/)	oliver merrick	oliver merrick	7	9,7	6.2	3.3	5.8	The closest thing in the whole 266-reel corpus to an explicit AI-boundary statement, and it is the pivot of the reel rather than a reassurance: thirty seconds are spent killing the obvious answer before anything is offered. Ungated (question_to_audience), so its 33.0 shares and 57.6 saves per 1 000 are clean, and its pain label, chaos_no_process, is the nearest in the pool to M2's territory. The contact sheet is one unbroken filmed-monitor take with a live hand pointing at the parts being named.
DczDlj4qQjI (https://www.instagram.com/reel/DczDlj4qQjI/)	liamjohnston.ai	liamjohnston.ai	1	13, 448	11.9x	2.4	5.7	The one reel in the pool that puts an approval step on screen and says it out loud, delivered as an objection beat at 43-48s - immediately after the capability list and immediately before the setup steps, which is exactly where a viewer worried about control would otherwise leave. 23.6 shares and 56.6 saves per 1 000 on an 11.9x lift. Its contact sheet shows the useful frame: the draft sitting in a composer, unposted, rather than a finished post.
Dbi-yxNgrhG (https://www.instagram.com/reel/Dbi-yxNgrhG/)	mantanhanjethwa	mantanhanjethwa	2	35, 050	5.8	5.8	5.9	The highest share rate in the 40-reel reference pool at 58.4 per 1 000 - fourth in the 266-reel analysis-ready corpus, where 61.3, 59.2 and 58.9 are higher (reports/analysis/06-creators-and-references.md §6, row 22) - built on three cheap choices: a system everybody already uses, a bounded count, and settings the viewer has never looked at. It sits at robust_z 1.45, so it is a forwarding machine rather than a reach outlier, which is the behaviour we want. Its contact sheet shows the settings screen filmed on a real device with a finger physically on the toggle.
DTCOU9Dk-RTG (https://www.instagram.com/reel/DTCOU9Dk-RTG/)	buildwithingwi	buildwithingwi	2	4,1 24	52.0x	3.0	4.2	Demo-first in the most literal sense: the first word is an instruction and the screen is already executing it. Every second afterwards does exactly one job - prompt, offer, objection killed, proof, dare. Ungated, 29.7 shares and 42.0 saves per 1 000, 52.0x lift on its author's median of 4 124, and the A>SCREEN>A shape we intend to use. The contact sheet shows a real finger pointing at the exact field on a filmed monitor while a persistent white plate holds the promise above it.

DbVKVz0y8xo (olivermerrick___, explanation)

- Transcript (DbVKVz0y8xo-B1): “Because the AI knows how to cut a video but it doesn't know the exact moment that needs to be a clip. It has no context and no content intelligence.”
- Frame (DbVKVz0y8xo-S00): one unbroken 87-second take filmed off a real monitor, a hand physically pointing at the element under discussion, no presenter face until the end
- Transformed: Borrowed: the kill-the-obvious-answer-first structure and the explicit naming of what the machine cannot judge. Changed: the boundary becomes the subject of the whole reel rather than a pivot inside a build, and it is stated about the viewer's outward-facing work rather than about clip selection. The same reel is our b_roll reference in H-09 for the held single take; the two uses do not overlap.

DczDlj4qQjI (liamjohnston.ai, proof)

- Transcript (DczDlj4qQjI-B3): “And nothing gets posted until you say yes. Every skill shows you the draft first and waits on you.”
- Frame (DczDlj4qQjI-S03): UI_DEMO of the draft still sitting in the composer with the publish control untouched - the waiting state, not the sent state

- Transformed: Borrowed: the sentence's job and its position, the control statement placed at the exact moment of doubt and shown on screen rather than promised. Changed: for us it is the thesis rather than a reassurance, and we show the draft actually waiting in a queue with a name on it instead of asserting that it would. Its split-screen opening and its comment gate are not borrowed.

Dbi-yxNgrhG (manthanjethwani, topic)

- Transcript (Dbi-yxNgrhG-B0): “If you don't have these four settings, you can access any private documents or bank statements on your Gmail.”
- Frame (Dbi-yxNgrhG-S01): the control being discussed is filmed on a real device with a finger resting on it, under a persistent black title plate
- Transformed: Borrowed: the topic shape - a bounded number of controls inside a system the viewer already runs, where the controls decide what software may do on their behalf - and the finger-on-the-real-control framing. Changed: our control is the hold-for-review setting in a tool the viewer is evaluating, not a mail provider's privacy toggle, and we verify what we claim. That reel's privacy claim is recorded as unverified in our own analysis and is not repeated.

DTCOU9DkRTG (buildingwithstring.com_, hook)

- Transcript (DTCOU9DkRTG-B0): “Watch this. Ask Chachi BT to make you a list of 10 companies hiring for a virtual assistant or receptionist. Boom.”
- Frame (DTCOU9DkRTG-S01): 26-second filmed-monitor run with a finger pointing at the exact field, persistent title plate held above it throughout
- Transformed: Borrowed: the imperative opening over an already-running screen, and the objection beat placed at 13-20s. Changed: our imperative points at a control the viewer can check in their own tool rather than at a money-making move, and the account behind this reel is the tool vendor's own - the thing our positioning most explicitly refuses. We borrow the timing and none of the intent.

Script

Watch this reply sit here. It is written, it is right, and it goes nowhere until somebody presses the button. You are choosing a tool that writes what your customers read. The feature list says what it produces. It never says where you sit. A person can sit in three places. Before the draft, which is a form. After it is sent, which is an apology. Or here, in between. Same task, twice. First with the tool set to send. Now set to hold. One queue, one draft, one name on it, a button nobody has pressed. We counted our own. Twenty-five named steps from a video coming in to the job being finished, and one place where it waits for a person. That person struck out nine of the last twenty-three. So the question for the sales call is not which model is inside. It is: show me where my draft waits. If that screen does not exist, do not buy it for anything a customer reads. Send this to whoever signs.

173 words / 69.2s (2.5 wps target) — tone: steady, practical, slightly stubborn; a buyer's question rather than a warning; emotional effect: the quiet of knowing nothing will go out with your name on it that you have not seen

ROLE	WORDS	SECONDS	CUMULATIVE
hook	20	8.0	8.0s
problem	24	9.6	17.6s
explanation	26	10.4	28.0s
proof	27	10.8	38.8s
example	35	14.0	52.8s
payoff	21	8.4	61.2s
cta	20	8.0	69.2s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
0-8.0s	Watch this reply sit here. It is written, it is right, and it goes nowhere until somebody...	hook	UI_D EMO	demo	Full-frame, filmed off a real monitor: a finished draft sitting in a queue, the release control untouched, a name in the owner column. A finger comes into frame and rests beside the un-pressed button without pressing it. No presenter.	Written. Correct. Not sent.	none	DTCOU9DkRTG (imperative opening over an already-running screen, finger pointing at the exact field); DczDlj4qQjI (the waiting state on screen rather than the sent state)	template	
8.0-17.6s	You are choosing a tool that writes what your customers read. The feature list says what...	promble m	A_R OLL_ MEDI UM	arol l	Michael, medium shot, one location, banner plate held above him.	Feature lists never say this	hard_c ut	DbVKVz0y8xo (the obvious answer is killed before anything is offered)	template	
17.6-28.0s	A person can sit in three places. Before the draft, which is a form. After it is sent, wh...	explan ation	A_R OLL_ MEDI UM	arol l	Same take continuing. A three-line list reveals one line at a time beside him: before the draft, after it is sent, in between. The third line stays after the other two dim.	Before. After. In between.	overlay _reveal	DbVKVz0y8xo (naming plainly what the machine is not allowed to judge, one item at a time)	template	
28.0-38.8s	Same task, twice. First with the tool set to send. Now set to hold. One queue, one draft,...	proof	SCREEN_ ECO RDING	demo	Full-frame recording, the two runs one after the other and never side by side: the same task with send on, then the same task with hold on, ending on the queue with the draft, the owner name and the untouched button.	Same task. Send, then hold.	hard_c ut	DczDlj4qQjI (the control statement is shown on screen at the exact moment of doubt, not promised); Dbi-yxNgrhG (the control is filmed on a real device with a finger on it)	template	
38.8-52.8s	We counted our own. Twenty-five named steps from a video coming in to the job being finis...	example	SCREEN_ ENS HOT	demo	Full-frame screenshot of our own plan table: twenty-three rows across three weeks, the nine struck rows visibly greyed out, with the stage name of the one step that waits for a person set as a mono label above the table.	25 steps. One waits.	hard_c ut	Dbi-yxNgrhG (a bounded count inside a system, stated as a number the viewer can hold)	template	
52.8-61.2s	So the question for the sales call is not which model is inside. It is: show me where my...	payoff	A_R OLL_ MEDI UM	arol l	Michael back on camera, same framing and light, no props.	Show me where my draft waits	hard_c ut	DTCOU9DkRTG (the reel comes back to the face to hand over the one thing to do)	template	
61.2-69.2s	If that screen does not exist, do not buy it for anything a customer reads. Send this to...	cta	A_R OLL_ MEDI UM	arol l	Same take, no reframe. Michael holds still for the closing line.	No hold screen, no purchase	none	—	template	

C-2026-09-12-06 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

Written. Correct. Not sent.

...

[SCREEN]

Full-frame, filmed off a real monitor: a finished draft sitting in a queue, the release control untouched, a name in the owner column. A finger comes into frame and rests beside the un-pressed button without pressing it. No presenter.

Watch this reply sit here. It is written, it is right, and it goes nowhere until somebody presses the button.

src: DTCOU9DkRTG, DczDlj4qQjI · asset: template

C-2026-09-12-06_S00.png

C-2026-09-12-06 · S01 · 8.0-17.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Feature lists never say this

[SHOT] medium

Michael, medium shot, one location, banner plate held above him.



You are choosing a tool that writes what your customers read. The feature list says what it produces. It never says where you sit.

src: DbVKVz0y8xo · asset: template


C-2026-09-12-06_S01.png

C-2026-09-12-06 · S02 · 17.6-28.0s · A_ROLL_MEDIUM · a_roll · ->overlay_reveal

Before. After. In between.

[SHOT] medium

Same take continuing. A three-line list reveals one line at a time beside him: before the draft, after it is sent, in between. The third line stays after the other two dim.



A person can sit in three places. Before the draft, which is a form. After it is sent, which is an apology. Or here, in between.

src: DbVKVz0y8xo · asset: template

C-2026-09-12-06_S02.png

C-2026-09-12-06 · S03 · 28.0-38.8s · SCREEN_RECORDING · demo · ->hard_cut

Same task. Send, then hold.

...

[SCREEN]

Full-frame recording, the two runs one after the other and never side by side: the same task with send on, then the same task with hold on, ending on the queue with the draft, the owner name and the untouched button.

Same task, twice. First with the tool set to send. Now set to hold. One queue, one draft, one name on it, a button nobody has pressed.

src: DczDlj4qQjI, Dbi-yxNgrhG · asset: template


C-2026-09-12-06_S03.png

C-2026-09-12-06 · S04 · 38.8-52.8s · SCREENSHOT · demo · ->hard_cut

25 steps. One waits.

[SCREEN]

Full-frame screenshot of our own plan table: twenty-three rows across three weeks, the nine struck rows visibly greyed out, with the stage name of the one step that waits for a person set as a mono label above the table.



We counted our own. Twenty-five named steps from a video coming in to the job being finished, and one place where it waits for a person. That person struck out nine of the last twenty-three.

src: Dbi-yxNgrhG · asset: template


C-2026-09-12-06_S04.png

C-2026-09-12-06 · S05 · 52.8-61.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Show me where my draft waits

[SHOT] medium


Michael back on camera, same framing and light, no props.



So the question for the sales call is not which model is inside. It is: show me where my draft waits.

src: DTCOU9DkRTG · asset: template

C-2026-09-12-06_S05.png



C-2026-09-12-06_S06.png

C-2026-09-12-06 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

Written. Correct. Not sent.

...

[SCREEN]
Full-frame, filmed off a real monitor: a finished draft sitting in a queue, the release control untouched, a name in the owner column. A finger comes into frame and rests beside the un-pressed button without pressing it. No presenter.

src: DTCOU9DkRTG, DczDlj4qQjI · asset: template

S00 0.0-8.0s

C-2026-09-12-06 · S01 · 8.0-17.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Feature lists never say this

[SHOT] medium
Michael, medium shot, one location, banner plate held above him.



src: DbVKVz0y8xo · asset: template

S01 8.0-17.6s

C-2026-09-12-06 · S02 · 17.6-28.0s · A_ROLL_MEDIUM · a_roll · ->overlay_reveal

Before. After. In between.

[SHOT] medium
Same take continuing. A three-line list reveals one line at a time beside him: before the draft, after it is sent, in between. The third line stays after the other two dim.



src: DbVKVz0y8xo · asset: template

S02 17.6-28.0s

C-2026-09-12-06 · S03 · 28.0-38.8s · SCREEN_RECORDING · demo · ->hard_cut

Same task. Send, then hold.

...

[SCREEN]
Full-frame recording, the two runs one after the other and never side by side: the same task with send on, then the same task with hold on, ending on the queue with the draft, the owner name and the untouched button.

src: DczDlj4qQjI, Dbi-yxNgrhG · asset: template

S03 28.0-38.8s

C-2026-09-12-06 · S04 · 38.8-52.8s · SCREENSHOT · demo · ->hard_cut

25 steps. One waits.

...

[SCREEN]
Full-frame screenshot of our own plan table: twenty-three rows across three weeks, the nine struck rows visibly greyed out, with the stage name of the one step that waits for a person set as a mono label above the table.

src: Dbi-yxNgrhG · asset: template

S04 38.8-52.8s

C-2026-09-12-06 · S05 · 52.8-61.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Show me where my draft waits

[SHOT] medium
Michael back on camera, same framing and light, no props.



src: DTCOU9DkRTG · asset: template

S05 52.8-61.2s

C-2026-09-12-06 · S06 · 61.2-69.2s · A_ROLL_MEDIUM · a_roll · ->none

No hold screen, no purchase

[SHOT] medium
Same take, no reframe. Michael holds still for the closing line.





If that screen does not exist, do not buy it for anything a customer reads. Send this to whoever signs.

src: - · asset: template

S06 61.2-69.2s

Contact sheet

Editing

- **Cut timing:** Seven sections, four visible cuts, each a full-frame screen insert replacing the presenter or the presenter returning. The presenter material is one continuous take; there is no speaker change, and a speaker change is the only cut PRODUCTION.md allows.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 40px, Paper on an Ink plate with 24px padding, bottom safe zone, two lines maximum. One persistent Ink banner plate in Saira SemiCondensed Black at the top for the whole reel. No emoji on screen.
- **Emphasis:** The word 'hold' is the only word that ever takes a plate of its own. Oxide is reserved for the sent run in the proof scene, because that is the state we are warning about; the held run is Signal.
- **Zooms:** One ui_zoom onto the untouched release button in the hook. No zoom on any A-roll scene.
- **Motion:** Linear easing only, per design/rules/motion.md. No bounce, no gradients, no rounded corners, no centred text.
- **Assets required:**
 - Filmed-monitor shot of a real draft queue with an untouched release control and a name in the owner column
 - Screen recording of one small task run twice, full-frame and sequential: send on, then hold on
 - Screenshot of our own plan table, twenty-three rows over three weeks with the nine struck rows greyed
 - Presenter A-roll, one continuous take, one location, banner plate fixed

Claims

CLAIM	STATE	SOURCE
Our own pipeline has twenty-five named processing stages between a video being discovered and a run completing.	OBSERVED	engine/SPEC.md §3, the canonical stage list: DISCOVERED through COMPLETED, twenty-five names.
Across three weeks that pipeline produced 23 plans and a person struck out 9 of them.	OBSERVED	data/radar.db cards : 23 rows over weeks 2026-09-03 / 09-07 / 09-10, of which 9 carry the struck-out status; reports/audit/04-existing-cards-review.md §A.
The paid provider in our pipeline is registered but switched off, and the default path needs no credentials.	OBSERVED	engine/SPEC.md §0.3 - the production default path is local; the paid provider is OPTIONAL_PROVIDER only.
The two runs shown on screen - the same task with send on and with hold on - are recorded on the shoot day in a tool the presenters already use.	PLANNED	No footage exists yet. The card must not be shot without it; asserting the behaviour without the recording would break hard rule 2.
That the human strike-out step is formally one of the twenty-five named stages.	MISSING	The stage list does not name it. It happens between card generation and script writing and is not recorded as a stage, so the script says 'one place where it waits for a person' rather than naming a stage.
That a hold-and-release setting prevents mistakes from reaching customers.	TO_MEASURE	We have not measured error rates with and without a gate, in our own pipeline or anyone else's. The reel gives the viewer a question to ask, not a measured effect.
Which named tools do and do not have the setting.	MISSING	No vendor is named and no comparison is claimed anywhere in the script.
That the corpus contains only one reel stating an AI boundary out loud.	OBSERVED	reports/analysis/08-hypotheses.md H-04 novelty note over the 266 analysis-ready reels; the single instance is DczDlj4qQjI beat B3. Used in the card rationale only, never spoken on screen.

Traceability

- **Insights:** I-08, I-11, I-09, I-13, I-15, I-27, I-05
- **Hypothesis:** H-04
- **Evidence summary:** Exactly one reel in the 266-reel analysis-ready corpus states an AI boundary out loud, and it does so as a throwaway objection beat inside a tool promotion (DczDlj4qQjI, beat B3, 56.6 saves per 1 000): making the boundary the subject is unrepresented. Our own gate count is real and checkable: twenty-five named stages in engine/SPEC.md §3, one place that waits for a person, and that person struck 9 of the 23 plans the pipeline produced across three weeks. The visual grammar follows the two strongest findings we have - opening on a screen carries median share_rate 0.0213 against 0.0140 with a creator-normalised correlation of +0.285, p=0.00025 (I-09), and the A>SCREEN>A run carries the same 0.0213 against 0.0140 at p=0.00076 across 39 reels and 32

creators (I-11). The topic reference Dbi-yxNgrhG holds the reference pool's highest share rate at 58.4 per 1 000, fourth in the analysis-ready corpus on a bounded-controls-in-a-familiar-system shape, and DTCOU9DkRTG (52.0x lift, ungated) supplies the imperative opening over an already-running screen.

- **Original synthesis:** The subject is the pause itself, which nothing in the corpus treats as the subject. The gate count is ours, dated and reproducible from engine/SPEC.md §3 and data/radar.db, and it is offered as an experiment rather than as a client result. The three-positions frame - before the draft, after the send, in between - is ours and is not in any reference.
- **Why better than generic:** A generic AI script would list five tools with a 'human in the loop' bullet and call the boundary a best practice. This one replaces the feature comparison with a single screen the viewer can ask to see, films the two runs full-frame and sequentially rather than as the split screen the corpus over-uses and under-rewards (I-10), names our own gate ratio instead of a principle, and refuses to claim that the gate prevents anything, because we have not measured it.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [originality] The script opened with 'Watch this.' - verbatim the first sentence of DTCOU9DkRTG beat B0, the very beat this card quotes in references[].useful_transcript. Script-writer rule 4 puts verbatim reference wording in the quote field and nowhere else → hook rewritten to 'Watch this reply sit here. It is written, it is right, and it goes nowhere until somebody presses the button.' It keeps the hypothesis' imperative-over-a-running-screen shape and removes the borrowed sentence. 20 words / 8.0 s, at the ceiling and inside the corpus band.
 - [evidence] The script said 'Twenty-five named steps between a video going in and a plan coming out'. engine/SPEC.md §3 lists 25 stages from DISCOVERED to COMPLETED; the plan appears at CARD_GENERATION, the sixteenth, and nine more stages follow it → rewritten to 'Twenty-five named steps from a video coming in to the job being finished'. Same 35 words, no timing change. claims[0], which already said 'between a video being discovered and a run completing', was already correct and is unchanged.
 - [evidence] Reference Dbi-yxNgrhG was described as 'the highest share rate in the entire analysis-ready corpus'. It is fourth: Dc_tsSeAjBy 61.3, DcvBtiNtbYO 59.2 and DMu90OwPzno 58.9 per 1 000 are higher, and the card's own source (06-creators-and-references §6, row 22) already says 'third-highest' → corrected in the reference reason and in traceability to 'highest in the 40-reel pool, fourth in the analysis-ready corpus'.
 - [evidence] DTCOU9DkRTG was described as '52x its author's median'. 52.04 is view_lift, i.e. play/median - 1, so the reel is 53x its author's median → reworded to '52.0x lift'.
 - [evidence] The 25-stage list, the 23 cards, the 9 struck and engine/SPEC.md §0.3 (paid provider registered but off, no credentials on the default path) all reproduce - 25 names in SPEC §3, and data/radar.db cards returns 23 rows with 9 'вычеркнута' → left.
 - [evidence] claims[7] is faithful to the H-04 novelty note in data/analysis/hypotheses.json and reports/analysis/08-hypotheses.md, and the single instance really is DczDlj4qQjI beat B3 → left.
 - [evidence] All four reference metric sets re-derived from data/radar.db and all four function tags match data/analysis/hypothesis_refs.json → left.
 - [positioning] The two runs are full-frame and sequential, never side by side (I-10: any split state median robust_z 1.03 against 1.43, p=0.044), which is also the hypothesis reviewer's explicit

warning → left.

- [evidence] The demo is PLANNED, not OBSERVED. The card must not be shot until both runs exist as footage → left, and repeated here so it survives the review.
- [positioning] No vendor named, no comparison claimed, no comment gate, no time or savings claim, CTA is one save_share ask → left.
- [visual] Opens on the screen at 0.0 s (I-09, RELIABLE) rather than the A>SCREEN>A the hypothesis proposed; the A>SCREEN>A run that I-11 rewards is preserved from scene 1 on. Three of seven scenes are screens, which is the fix for the compliance-lecture risk → left.
- [timing] Recomputed after the hook edit: 173 words, 69.2 s at 2.5 wps, storyboard contiguous 0.0-69.2 s, no scene under 1.5 s, every overlay ≤ 7 words → passes, inside the 50-70 s band.
- [schema] python3 -m engine.cards_v2 validate → 0 errors, 0 warnings after the edits.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-06.json
- **Remotion validation:** {"card_id": "C-2026-09-12-06", "checked_at": "2026-09-12T08:53:27Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2076, "edl": "cards/edl/C-2026-09-12-06.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-07 — The same number in two places

Strategy

Format: M2 Teardown

Concept. One figure in our own files has two correct values, fourteen times and one point nine times, because one calculation leaves the item out of its own comparison and the other leaves it in - and on a creator whose plays fall into two clumps with nothing between them, that alone moves the yardstick nearly eightfold. We acted on the bigger one.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company who is responsible for one repeated process and has tried ChatGPT but has no repeatable way of using it at work. Specifically anyone who quotes a figure that is calculated somewhere else - a margin, a price, a stock count, a headline number in a monthly report.

Objective. Send the viewer looking for the second place their most-quoted figure is worked out, by showing a fourteen-fold difference that survived a review inside our own files.

Positioning. builder

Pain. quality_trust

Promise. One page that stops the same figure meaning two things in two systems.

Rationale. No reel in the 266-reel corpus corrects its own number: 36 use numbers as proof and every one presents the number as settled. The references supply shape only - an exact figure spoken over the record it describes (DaGTeJNN1GR), one named thing per step with no connective tissue (DUJYKENjZc5), the screen carrying the payoff while the face carries only the interpretation (DYC-x8DogOI), and, as a negative example, a save bought entirely by withholding the artefact (DVtgbY8Av6A), which we invert by publishing it.

FILT ANSWER
ER

Process	Producing a recurring figure that other people act on - the margin in the monthly pack, the headline number in a report, a price, a stock count - where the figure is worked out in one system and quoted in another.
Friction	Both versions are defensible and both are quoted by people who trust them, so the disagreement is invisible until a decision has already been made on whichever one was nearer to hand. Reviewing harder does not find it, because neither number is wrong.
AI boundary	The machine can recompute the figure wherever it appears and flag the moment two places disagree. Deciding which definition is the authoritative one is a human decision, and it has to be written down once rather than re-argued each month.
Next action	Take the figure you quote most often, find the second place it is worked out, and write down which of the two is authoritative.

Hooks

1. **[result_first]** (statement) — **USED IN SCRIPT** Fourteen times normal. One point nine times normal. Same item, same week. The first one went into the plan.
1. **[bold_claim]** (statement) The same twelve videos gave us two different normals, seven times apart. Both were right.
1. **[question]** (question) Which number does your report use when two of them are correct?

Reference reels

REEL	CRE ATO R	FUN CTI ON	PL AY S	CREA TOR AN	VI EW FT	SHA RE RAT E	SAV E RAT E	REASON
DaGTe-JNN1GR (https://www.instagram.com/reel/DaGTeJNN1GR/)	aiwit hrem y	b_ro ll	29 8,2 26	28,9 59	9.3 x	3.0 %	3.6 %	The pool's only reel where the physical thing the number describes is itself the visual argument, cut as process B-roll around a figure being checked. An exact figure spoken over the thing it refers to is exactly how two disagreeing numbers have to be shown. 30.2 shares per 1 000, robust_z 3.41.
DUJYKENjZc5 (https://www.instagram.com/reel/DUJYKENjZc5/)	shar mad hruve shh	rhyt hm	53 4,9 53	1,37 4	388 .2x	3.7 %	4.6 %	33 seconds, four named things, one per step, and the deliverable is a file - the tightest one-named-thing-per-beat cadence in the pool. 36.6 shares and 46.3 saves per 1 000. The 388x rests on a 1 374-view author median, so the cadence is the finding and the multiplier is not.
DYC-x8DogOI (https://www.instagram.com/reel/DYC-x8DogOI/)	valeri does ai	scre en_ pro of	1,9 80, 12 8	4,94 6	399 .4x	2.4 %	6.0 %	Screen, face, screen - the mirror of the talk-show-come-back pattern - with the payoff delivered over the screen at 28.9-39.2 s and the face reserved for interpretation. Highest view_lift in the corpus with 59.9 saves per 1 000, and the whole reel is 43 seconds, which is roughly where a two-number trace has to land.
DVtgb-Y8Av6A (https://www.instagram.com/reel/DVtgbY8Av6A/)	alber t.olga ard	cta	27 2,2 94	16,4 90	15. 5x	2.0 %	5.7 %	The purest example of deliberate incompleteness in the pool and therefore the clearest thing we have to invert: 57.3 saves per 1 000 from 84 spoken words, where the save is bought entirely by withholding the files. Our own note also records that its hook promises fifteen items in sixty seconds and the reel is 25 seconds long and names eleven.

DaGTeJNN1GR (aiwithremy, b_roll)

- Transcript (DaGTeJNN1GR-B1): “It was 12,645 euros on booking.com. So we pasted the link and then asked it to try to find the property online and then where it was listed on other sites for cheaper.”
- Frame (DaGTeJNN1GR-S03): a phone held up to camera showing the arithmetic itself - a currency conversion with the claimed saving overlaid on the calculator - shot in the place the number is about
- Transformed: Borrowed shooting the artefact the number describes while the exact figure is spoken over it. Changed: our artefact is two records carrying the same figure with different values, and we make no savings claim of any kind - the original's dollar and euro figures are recorded in our own files as unreconciled, which is precisely the defect our reel is about.

DUJYKENjZc5 (sharmadhruveshh, rhythm)

- Transcript (DUJYKENjZc5-B1): “Step 1, add OpenAI node to write the short script. Step 2, add a level lab to generate the voice over.”
- Frame (DUJYKENjZc5-S00): frame one is the workflow already executing in a floating card above the presenter, one caption word at the seam, no establishing shot of anybody

- Transformed: Borrowed one named thing per step with no connective tissue, so a trace through two systems stays followable. Changed: the things named are the two places our figure is worked out rather than four products, the deliverable is a one-page definition sheet published in the caption rather than a workflow file behind a keyword, and we do not copy its floating-card layout. Our notes record this reel's tool names as heavily mis-transcribed, so nothing verbal is carried across.

DYC-x8DogOI (valeridoesai, screen_proof)

- Transcript (DYC-x8DogOI-B2): “But here's the crazy part. We wrapped all four into one single repo. So you literally paste one line and Claude builds the whole thing.”
- Frame (DYC-x8DogOI-S00): opens on the finished artefact filling the frame before any face appears, then drops to a plain seated close-up only to say what it means
- Transformed: Borrowed the screen-face-screen arrangement and the habit of letting the screen carry the payoff while the face carries only the interpretation. Changed: the payoff is a correction rather than a shortcut, the \$10 000 price anchor is dropped entirely because money framing runs at median view_lift 1.19 against 2.12, and the two numbers sit in one full-frame shot rather than in two half-frames.

DVtgbY8Av6A (albert.olgaard, cta)

- Transcript (DVtgbY8Av6A-B3): “If you want all of my skills completely for free just come in skills and I'll send them to you.”
- Frame (DVtgbY8Av6A-S01): each named item appears as a real file listing floated beside the presenter, so the withheld thing is visibly real - which is what makes the withholding work
- Transformed: Borrowed nothing of its mechanism - it is the negative example. What is borrowed is the proof that the save comes from there being something concrete to collect. Changed: we publish the one-page definition sheet in the caption and say so in the close, which converts the same impulse without the gate, and the promise in our hook is the promise the reel actually keeps.

Script

Fourteen times normal. One point nine times normal. Same item, same week. The first one went into the plan. Think of two receipts for the same job with different totals. Both correct. Added up differently. You quote a number like that every month. We hold twelve of this person's videos. Six under twenty thousand views. Six over two hundred and eighty thousand. Nothing in between. Normal is the middle one. Leave this video out and the middle is nineteen thousand. Leave it in and the middle falls in the gap, at a hundred and fifty thousand. Neither number is wrong. Same twelve videos, both times. Nobody checked which one the decision needed, and the video is ordinary for this author. A number is not a fact. It is the answer to a question, and the question has to travel with it. One definition, written down once, carried everywhere the number goes. The one-page definition sheet is in the caption. Take the figure you quote most, find the second place it is worked out.

173 words / 69.2s (2.5 wps target) — tone: plain, forensic, no statistics - two receipts for the same job; emotional effect: the quiet worry that a figure you have quoted to someone was the other one

ROLE	WORDS	SECONDS	CUMULATIVE
hook	19	7.6	7.6s
setup	24	9.6	17.2s
explanation	22	8.8	26.0s
explanation	31	12.4	38.4s
proof	24	9.6	48.0s
payoff	31	12.4	60.4s
cta	22	8.8	69.2s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE INSPIRATION	ASSET STATUS
0.0-7.6s	Fourteen times normal. One point nine times normal. Same item, same week. The first one w...	hook	DAT A_VI SUA L	demo	Full frame, no presenter: the two figures from our own files side by side, 14.4x on the left and 1.87x on the right, both labelled with the same item and the same date.	14.4x and 1.87x. Same item.	none	DYC-x8DogOI (opens on the artefact filling the frame before any face, and lets the screen carry the claim)	template
7.6-7.2s	Think of two receipts for the same job with different totals. Both correct. Added up diff...	setup	A_R OLL _ME DIU M	roll	Max, medium shot, plain wall, direct to camera.	Two receipts, one job	hardcut	DYC-x8DogOI (the face appears only to say what the screen meant, then hands straight back)	template
17.2-26.0s	We hold twelve of this person's videos. Six under twenty thousand views. Six over two hun...	exposition	SCREEN SHOT	demo	Full-frame capture of all twelve of this creator's rows in our own files, sorted by views: six between seven thousand and nineteen thousand, a blank stretch, then six between two hundred and eighty thousand and seven hundred and sixty-eight thousand.	Twelve videos. Nothing in the middle.	hardcut	DUJYKENjZc5 (one named thing per beat with no connective tissue, so a trace stays followable)	template
26.0-38.4s	Normal is the middle one. Leave this video out and the middle is nineteen thousand. Leave...	exposition	SCREEN SHOT	demo	The same unbroken screen take: the middle of the same twelve rows marked twice - once with this video taken out, landing on 19,624, once with it left in, landing at 151,570 inside the empty stretch.	Out: 19,624. In: 151,570 .	none	DUJYKENjZc5 (second step given exactly the shape of the first, so the difference is the only thing that moves)	template
38.4-48.0s	Neither number is wrong. Same twelve videos, both times. Nobody checked which one the dec...	proof	DAT A_VI SUA L	demo	The same twelve rows: this video sits seventh of the twelve, at the bottom of the upper cluster, with both middles still marked.	Ordinary for its author	overlay reveal	DaGTeJNN1GR (the exact figure spoken over the record it refers to, not over a stand-in)	template
48.0-60.4s	A number is not a fact. It is the answer to a question, and the question has to travel wi...	payoff	A_R OLL _ME DIU M	roll	Max, back on camera, medium shot, same setup as the setup beat.	The question travels with the number	hardcut	DYC-x8DogOI (the face carries only the interpretation after the screen has carried the payoff)	template
60.4-69.2s	The one-page definition sheet is in the caption. Take the figure you quote most, find the...	cta	A_R OLL _ME DIU M	roll	Max, same take, closing line to camera.	Definition sheet in the caption	none	DVtgbY8Av6A (inverted: the same collect-something impulse, with the artefact published instead of withheld)	template

C-2026-09-12-07 · S00 · 0.0-7.6s · DATA_VISUAL · demo · ->none

14.4x and 1.87x. Same item.

...

[SCREEN]

Full frame, no presenter: the two figures from our own files side by side, 14.4x on the left and 1.87x on the right, both labelled with the same item and the same date.

Fourteen times normal. One point nine times normal. Same item, same week. The first one went into the plan.

src: DYC-x8DogOI · asset: template

C-2026-09-12-07_S00.png


C-2026-09-12-07 · S01 · 7.6-17.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Two receipts, one job

...

[SHOT] medium

Max, medium shot, plain wall, direct to camera.



Think of two receipts for the same job with different totals. Both correct. Added up differently. You quote a number like that every month.

src: DYC-x8DogOI · asset: template

C-2026-09-12-07_S01.png

C-2026-09-12-07 · S02 · 17.2-26.0s · SCREENSHOT · demo · ->hard_cut

Twelve videos. Nothing in the middle.

...

[SCREEN]

Full-frame capture of all twelve of this creator's rows in our own files, sorted by views: six between seven thousand and nineteen thousand, a blank stretch, then six between two hundred and eighty thousand and seven hundred and sixty-eight thousand.

We hold twelve of this person's videos. Six under twenty thousand views. Six over two hundred and eighty thousand. Nothing in between.

src: DUJYKENjZc5 · asset: template

C-2026-09-12-07_S02.png

C-2026-09-12-07 · S03 · 26.0-38.4s · SCREENSHOT · demo · ->none

Out: 19,624. In: 151,570.

...

[SCREEN]

The same unbroken screen take: the middle of the same twelve rows marked twice - once with this video taken out, landing on 19,624, once with it left in, landing at 151,570 inside the empty stretch.

Normal is the middle one. Leave this video out and the middle is nineteen thousand. Leave it in and the middle falls in the gap, at a hundred and fifty thousand.

src: DUJYKENjZc5 · asset: template

C-2026-09-12-07_S03.png

C-2026-09-12-07 · S04 · 38.4-48.0s · DATA_VISUAL · demo · ->overlay_reveal

Ordinary for its author

...

[SCREEN]

The same twelve rows: this video sits seventh of the twelve, at the bottom of the upper cluster, with both middles still marked.

Neither number is wrong. Same twelve videos, both times. Nobody checked which one the decision needed, and the video is ordinary for this author.

src: DaGTeJNNiGR · asset: template

C-2026-09-12-07_S04.png


C-2026-09-12-07 · S05 · 48.0-60.4s · A_ROLL_MEDIUM · a_roll · ->hard_cut

The question travels with the number

...

[SHOT] medium

Max, back on camera, medium shot, same setup as the setup beat.



A number is not a fact. It is the answer to a question, and the question has to travel with it. One definition, written down once, carried everywhere the number goes.

src: DYC-x8DogOI · asset: template


C-2026-09-12-07_S05.png

C-2026-09-12-07 · S06 · 60.4-09.2s · A_ROLL_MEDIUM · a_roll · ->none

Definition sheet in the caption

[SHOT] medium

Max, same take, closing line to camera.



The one-page definition sheet is in the caption.
Take the figure you quote most, find the second place it is worked out.

src: DVtgbY8Av6A · asset: template

C-2026-09-12-07_S06.png

C-2026-09-12-07 · S00 · 0.0-7.6s · DATA_VISUAL · demo · ->none

14.4x and 1.87x. Same item.

...

[SCREEN]

Full frame, no presenter: the two figures from our own files side by side, 14.4x on the left and 1.87x on the right, both labelled with the same item and the same date.

S000.0-7.6s

C-2026-09-12-07 · S01 · 7.6-17.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Two receipts, one job

[SHOT] medium

Max, medium shot, plain wall, direct to camera.



S017.6-17.2s

C-2026-09-12-07 · S02 · 17.2-26.0s · SCREENSHOT · demo · ->hard_cut

Twelve videos. Nothing in the middle.

...

[SCREEN]

Full-frame capture of all twelve of this creator's rows in our own files, sorted by views: six between seven thousand and nineteen thousand, a blank stretch, then six between two hundred and eighty thousand and seven hundred and sixty-eight thousand.

S0217.2-26.0s

C-2026-09-12-07 · S03 · 26.0-38.4s · SCREENSHOT · demo · ->none

Out: 19,624. In: 151,570.

...

[SCREEN]

The same unbroken screen take: the middle of the same twelve rows marked twice - once with this video taken out, landing on 19,624, once with it left in, landing at 151,570 inside the empty stretch.

S0326.0-38.4s

C-2026-09-12-07 · S04 · 38.4-48.0s · DATA_VISUAL · demo · ->overlay_reveal

Ordinary for its author

...

[SCREEN]

The same twelve rows: this video sits seventh of the twelve, at the bottom of the upper cluster, with both middles still marked.

S0438.4-48.0s

C-2026-09-12-07 · S05 · 48.0-60.4s · A_ROLL_MEDIUM · a_roll · ->hard_cut

The question travels with the number

[SHOT] medium

Max, back on camera, medium shot, same setup as the setup beat.



S0548.0-60.4s

C-2026-09-12-07 · S06 · 60.4-69.2s · A_ROLL_MEDIUM · a_roll · ->none

Definition sheet in the caption

[SHOT] medium

Max, same take, closing line to camera.





The one-page definition sheet is in the caption.
Take the figure you quote most, find the second
place it is worked out.

src: DVtgbY8Av6A · asset: template

S06 60.4-69.2s

Contact sheet

Editing

- **Cut timing:** Three hard cuts, every one of them a handover between the screen and a presenter: screen to Max at 7.6 s, Max back to the screen at 17.2 s, and the screen to Max at 48.0 s. The opening frame is not a cut. The two screen blocks each run as a single take; the closing face block runs as a single take through payoff and CTA.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 600 at 40px, two lines maximum, bottom-anchored in Instagram zone A/B; one persistent top banner held for the whole reel.
- **Emphasis:** Every figure on screen is one we can open the file on: 14.4, 1.87, 19,624, 151,570. The words median, baseline and percentile appear nowhere, on screen or in speech.
- **Zooms:** None. The two records are shown at a size that can be read without a zoom.
- **Motion:** Linear easing only, no bounce, per design/rules/motion.md.
- **Assets required:**
 - Side-by-side capture of the two figures for the same item from our own files (S00)
 - Capture of all twelve of that creator's rows sorted by views, with the empty stretch between the two clusters visible (S02)
 - The same twelve rows with the middle marked twice - this video out, then this video in - showing 19,624 and 151,570 (S03)
 - The same twelve rows with this video marked seventh of twelve (S04)
 - Max A-roll, medium shot, one location, two takes
 - One-page definition sheet published in the caption, not gated

Claims

CLAIM

ST
AT
E

SOURCE

One of our 23 drafted planning cards carried a headline figure of 14.4 times normal; recomputed against all twelve of that author's reels in our files the same item is 1.87 times normal.

OB
SE
RV
ED

reports/analysis/07-existing-cards-verdict.md section 1, card #20 DdCv-FdHsnj1; reproduced from data/radar.db on 2026-09-12 (283 515 plays / 19 624 = 14.45; 283 515 / 151 569.5 = 1.87).

Both yardsticks come from the same twelve reels by that creator in the same snapshot. 19 624 is the middle of the eleven that are not the item itself; 151 569.5 is the middle of all twelve.

OB
SE
RV
ED

scores.baseline_n = 11 and scores.author_median_play = 19 624 for DdCvFdHsnj1; video_features.creator_median_play = 151 569.5 with creator_n = 12; reels holds exactly 12 distinct codes for soojintech, all in snapshot 3. score.py excludes a reel from its own baseline and narrows to an age band only when that band itself has enough reels - baseline_n = 11 = 12 minus 1 shows the age band was not applied, so there is no narrow window and no separate history: the two figures differ only by whether the item is inside its own comparison.

That creator's plays fall into two clumps with nothing between them - six reels from 7 402 to 19 624 and six from 283 515 to 768 383 - so the middle sits in the empty stretch and one removal moves it across the gap.

OB
SE
RV
ED

reels in data/radar.db, snapshot 3, username soojintech, 12 distinct codes: 7 402, 8 231, 9 725, 10 951, 13 057, 19 624, 283 515, 319 265, 325 781, 326 412, 652 304, 768 383. Middle of all twelve = (19 624 + 283 515)/2 = 151 569.5; middle of the eleven without 283 515 = 19 624.

Our own stats already label this creator HIGH_VARIANCE on twelve videos, but the label is not diagnostic on its own: 115 of 132 creators carry it.

OB
SE
RV
ED

creator_stats for soojintech: n_videos 12, reliability HIGH_VARIANCE (engine/SPEC.md section 7). The 115-of-132 share is I-23, quoted in reports/analysis/09-reference-selection.md R1. Recorded so the reel does not imply a warning was ignored that would have singled this creator out.

The two yardsticks differ by nearly eight times: 19,624 against 151,570.

OB
SE
RV
ED

Same two rows; 151 569.5 / 19 624 = 7.72.

The item is ordinary for its author: seventh of that author's twelve reels, 54th percentile, outlier_status NORMAL.

OB
SE
RV
ED

video_features for DdCvFdHsnj1: percentile_in_creator 54.17, robust_z 0.492, outlier_status NORMAL.

The card was prioritised on the strength of the larger figure and was never shot.

OB
SE
RV
ED

reports/audit/04-existing-cards-review.md row 20 (week 09-10, priority 6, 'Yes, 14.4x'); our_posts is empty, so nothing was published.

CLAIM	STATE	SOURCE
Every card's headline figure will be recomputed against the accumulated creator median, with the number of items behind it travelling alongside.	PLANED	reports/analysis/07-existing-cards-verdict.md §1, 'Rule this implies'; not yet enforced in the pipeline.
Carrying one written definition with the figure prevents the error from recurring.	TOMESURE	The definition sheet has not yet been used on a full weekly run.
Every business has the same figure calculated in two places.	MISING	Asserted as a parallel in the script and never claimed as measured - we have not looked inside any company's books. Flagged so the reel does not present the parallel as a finding.

Traceability

- **Insights:** I-23, I-13, I-08, I-09, I-01, I-29, I-15
- **Hypothesis:** H-20
- **Evidence summary:** One of our 23 drafted planning cards carried a headline figure of 14.4 times normal. The same item recomputed against all twelve of that author's reels in our files is 1.87 times normal, and it sits seventh of the twelve, at the 54th percentile, outlier_status NORMAL. Both figures come from the same twelve reels in the same snapshot and differ only by whether the item is inside its own comparison: 19 624 is the middle of the other eleven, 151 569.5 the middle of all twelve. It moves that far because the creator's plays fall into two clumps - six from 7 402 to 19 624, six from 283 515 to 768 383 - with the middle sitting in the empty stretch. Both figures reproduce from data/radar.db. The four references contribute shape only: an exact figure spoken over the thing it describes (DaGTeJNN1GR, 30.2 shares/1k, robust_z 3.41), one named thing per step (DUJYKENjZc5, 36.6 shares/1k), the screen carrying the payoff and the face only the interpretation (DYC-x8DogOI, 59.9 saves/1k) and, inverted, a save bought by withholding the artefact (DVTgbY8Av6A, 57.3 saves/1k).
- **Original synthesis:** No reel in the 266-reel corpus corrects its own number - 36 use numbers as proof and every one presents the number as settled. The wrong figure, the right figure, the twelve rows behind both and the decision taken on the larger one are ours and dated. The script deliberately contains no statistical vocabulary: the mechanism is delivered as two receipts for the same job with different totals, and the artefact we hand over is a one-page definition sheet published in the caption rather than gated behind a comment.
- **Why better than generic:** A generic AI script would warn that 'AI can hallucinate numbers' and move on. Nothing hallucinated here: both figures are arithmetically correct on the same twelve rows, which is the whole point and the reason reviewing harder does not catch it. This reel names the exact item, both middles, the gap in the data that makes them seven times apart, and the one-line rule that closes it - and it does so without using the words median, baseline or percentile once.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] HARD - the mechanism was wrong. The script said the first figure compared the item 'against eleven others from one narrow window' and the second 'against everything that person had ever posted'. Neither is true: reels holds exactly twelve reels for soojintech, all in snapshot 3, and both yardsticks are computed from those same twelve. 19 624 is the middle of the eleven that are not the item (score.py excludes a reel from its own baseline); 151 569.5 is the middle of all twelve. $\text{baseline_n} = 11 = 12 \text{ minus } 1$ proves the age band was never applied, because score.py narrows to an age band only when the band itself has enough reels. -> both explanation sections rewritten, claims[1] rewritten, S02/S03 visuals, overlays and assets rebuilt. Script version 2.
 - [evidence] the real cause is now in the reel: that creator's plays fall into two clumps - six from 7 402 to 19 624 and six from 283 515 to 768 383, nothing in between - so the middle sits in the empty stretch and removing one high value drops it across the gap. Added as an OBSERVED claim with all twelve play counts listed.
 - [evidence] 'We acted on the first one' implied action beyond what happened - reports/audit/04 row 20 records the card as drafted, week 09-10, priority 6, never shot -> 'The first one went into the plan', in the script and in hooks[0].
 - [evidence] added an OBSERVED claim that creator_stats labels this creator HIGH_VARIANCE, with the caveat that 115 of 132 creators carry the same label (I-23), so the reel does not imply a decisive warning was ignored. Deliberately kept out of the spoken script.
 - [hook] hooks[1] was hooks[0] reordered -> replaced with a distinct bold_claim on the two normals.
 - [visual] editing.cut_timing said 'four hard cuts, every one of them a handover'; three are handovers and the fourth was the opening frame's own transition -> S00 set to none and the count corrected to three.
 - [evidence] verified unchanged: 14.45 and 1.87, baseline_n 11, author_median_play 19 624, creator_median_play 151 569.5, $151\,569.5 / 19\,624 = 7.72$, percentile_in_creator 54.17, robust_z 0.492, outlier_status NORMAL - all reproduce from video_features and scores. The figures were never the problem; their explanation was.
 - [pacing] left as is: the CTA carries a caption pointer and an action against general-conclusions section G rule 14.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-07.json
- **Remotion validation:** {"card_id": "C-2026-09-12-07", "checked_at": "2026-09-12T08:53:27Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2076, "edl": "cards/edl/C-2026-09-12-07.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-08 — The paragraph that beat the agents

Strategy

Format: M2 Radar

Concept. The same job run twice on screen - one paragraph of written instructions, then a multi-step build - with the comparison happening in the footage rather than in an argument. Then the same lesson taken from our own work: a hundred and fifteen measurements across two hundred and sixty-eight videos separated nothing, and the one instruction that survived was a question we could have asked on day one.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company responsible for one repeated process. Specifically someone who has been told they need an agent.

Objective. Get one viewer to write their repeated task out as a single paragraph with three checks and run it by hand five times before commissioning anything larger.

Positioning. builder

Pain. dont_know_where_to_start

Promise. A first version that works this week and that the person doing the job can change without asking anybody.

Rationale. Agent building is the highest-lift topic in the corpus (n=27, median view_lift 6.63) and 23 reels carry an agent_build solution type; not one of them argues against the build (I-25). myth_bust carries the second-highest robust_z of any narrative at 1.93 on the smallest explanation share, which is this reel's shape. The register risk is real - contrarian hooks sit at 0.0076 share and 0.0139 save, the lowest of the large hook cells - so the reel opens on the demo, not on the disagreement.

FILT ANSWER
ER

Proce ss	A knowledge task the team repeats about five times a week: a quote, a call summary, a reply to a standard enquiry, a weekly status note - something with a clear input, a clear output and a person who owns it.
Frictio n	The advice on offer is to build a multi-step agent, and nobody in a company of nine can maintain one. The elaborate version fails quietly, at a step nobody in the room can read.
AI bound ary	A written paragraph plus three checks keeps the judgement with the person running the task; they can read every line and change any of it. A multi-step build moves that judgement into software that has no owner in the room.
Next action	Write the task out as one paragraph with three checks and run it by hand five times this week before automating anything.

Hooks

1. **[demo_first]** (statement) — **USED IN SCRIPT** Same job, twice. This one is a paragraph of instructions. That one is six agents. Watch where each stops.
1. **[result_first]** (statement) Here is the same job done two ways, and the version that finishes is the one you could have written in a text box.
1. **[question]** (question) Do you actually need an agent for a task you do five times a week? Watch both versions run.

Reference reels

REEL	CR EA TO R	FU NC TI ON	P L A R Y	CRE ATO R IAN	VI EW LI E TE	SH AR E RA TE	SA VE RA TE	REASON
Dbs2Jd-JKUDd (https://www.instagram.com/reel/Dbs2JdJKUDd/)	dai nika uto mat es	ex pl an ati on	1 6 2 9 0	3,6 19 x	3.5 3.0 4.3	% %		The only reel in the pool whose whole argument is that the heavy tool should be abandoned for one prompt, and its pain block is stated as acts rather than as a condition: think about it, adjust the layout, no colours, boring template. Pain first, demo second, objection third - the three things stopping a switch answered in the order they arise. Small numbers (16 290 plays) but a clean structure and robust_z 3.28.
DbVKVz0y8xo (https://www.instagram.com/reel/DbVKVz0y8xo/)	oliv erm erri ck_	b_ rol l 4	7 0 4 7	9,7 34 x	6.2 3.3 5.8	% %		A single unbroken B_ROLL_PROCESS state for the full 87 seconds - the highest median view_lift of any visual sequence with more than three reels behind it (30.90 on n=6, 4 creators; only a three-reel sequence, A_ROLL_CLOSE_UP>SCREENSHOT>A_ROLL_CLOSE_UP at 52.04, scores higher) - narrated over a live artefact. It proves that a long single take over a working screen holds attention, which is exactly what a two-runs comparison needs. Ungated, 33.0 shares and 57.6 saves per 1 000. The contact sheet shows a hand moving over a real monitor and no presenter face at all until the close.
DZQZOvhg-NvK (https://www.instagram.com/reel/DZQZOvhgNvK/)	jaso nco ope rso	pr oo f 4, 0 3 2	1 2 4, 3 2	24, 492 x	4.1 2.5 5.0	% %		Puts the output before the architecture and never lets the viewer hold an unanswered question: the result is named in full, then the phases. Proof inside the first fifteen seconds, numbered phases afterwards. 25.0 shares and 50.2 saves per 1 000, 4.06x lift, robust_z 2.18. Its contact sheet also shows the tell we are inverting - a running clock overlay from the midpoint on, dramatising speed that is never actually compared with anything.
DVaBP3w-DAFE (https://www.instagram.com/reel/DVaBP3wDAFE/)	tenf old mar c	to pic 3, 5 2 8	5 3 5 8	7,4 71 4x	70. 2.2 3.8	% %		The archetype of the thing this reel argues against, and therefore the reference that defines our angle: a six-agent command centre narrated as self-evidently good, two thirds of its runtime an unbroken UI_DEMO, demo before threat, gate last. 70.4x lift, robust_z 5.0 clipped, 22.4 shares per 1 000. Its contact sheet gives us the exact look of the elaborate run: a dense multi-panel dashboard filmed across the room off a real monitor.

Dbs2JdJKUDd (dainikautomates, explanation)

- Transcript (Dbs2JdJKUDd-B1): “If you make a carousel in Kenwa, it takes hours. Think about it, adjust the layout. There are no colors. And in the end, it becomes the boring template.”
- Frame (Dbs2JdJKUDd-S02): the replacement workflow is shown running full-frame, then cut back to the face only to interpret it
- Transformed: Borrowed: the pain-demo-objection ordering and the habit of naming the pain as a sequence of acts rather than a state. Changed: our comparison is run on screen in both directions rather than asserted, the task is an ordinary business task rather than a carousel, and 'Kenwa' is our own ASR's spelling of a product name - a standing reminder that we quote the structure and never the words.

DbVKVz0y8xo (olivermerrick___, b_roll)

- Transcript (DbVKVz0y8xo-B2): “build this thing. First you're going to scrape any short form account and get all of their clips.”
- Frame (DbVKVz0y8xo-S00): one continuous camera-filmed monitor take for the entire runtime, a live hand pointing at each part as it is named, word-by-word captions
- Transformed: Borrowed: one continuous process shot with narration over it and no cut to a presenter until the end. Changed: we run the shot twice - the small version, then the elaborate one - so the comparison happens in the footage rather than in a sentence; and the tech stack it recites in its CTA is replaced by a single paragraph plus three checks. Different function from this reel's use in H-04, which takes only the spoken boundary statement.

DZQZOVhgNvK (jasoncooperson, proof)

- Transcript (DZQZOVhgNvK-B1): “I get all the top-performing posts in my niche fully transcribed with a visual analysis. I get the light count, I get the hook, I get the comment count, I get an outlier score, I get everything that I need. Here is exactly how it works.”
- Frame (DZQZOVhgNvK-S00): the output is on screen before any mechanism is described, with a running clock overlay added from the midpoint
- Transformed: Borrowed: output-before-architecture ordering, so the viewer sees the result before any mechanism. Changed: we show two outputs rather than one, and the comparison is the point; we refuse its proof style, an outcome claim about views we are not allowed to make, and its paid-community CTA; and we do not use the clock overlay, because a speed comparison we have not measured is a claim.

DVaBP3wDAFE (tenfoldmarc, topic)

- Transcript (DVaBP3wDAFE-B1): “So I have my researcher my script writer my editor my scheduler and my analytics here”
- Frame (DVaBP3wDAFE-S00): dense multi-panel agent dashboard filmed across the room off a real monitor, badge-title held top-left, sentence-building captions at the bottom
- Transformed: Borrowed: nothing of its voice. What is borrowed is the subject as a foil - the dashboard-of-agents genre whose energy the corpus rewards (agent building, n=27, median view_lift 6.63) and whose shape our positioning refuses - plus the across-the-room filmed-dashboard framing for our elaborate run. Changed: we show the same kind of artefact stalling rather than triumphing, we name the person who owns the task, and we drop the fear register and the profanity of its last third entirely.

Script

Same job, twice. This one is a paragraph of instructions. That one is six agents. Watch where each stops. The task runs about five times a week. A quote, a summary, a reply. Whatever you build, somebody has to keep it running. This is the paragraph. What the task is, what good looks like, three checks at the bottom. Anyone doing the job can read every line. This is the other one. Six steps, four of them fine, and this one waiting on something nobody in the room can fix. We did the elaborate version to ourselves. A hundred and fifteen measurements across two hundred and sixty-eight videos. Not one reliably separated a strong reel from a weak one. What survived was a question we could have asked on day one. Three questions before you build. Who can change it. What happens when it is wrong. How long until it first runs. Write your task as one paragraph and run it by hand five times this week. Send this to whoever is scoping it.

174 words / 69.6s (2.5 wps target) — tone: unhurried, evidential, willing to be the one who overbuilt;
emotional effect: permission to start with the small thing, given by someone who did not

ROLE	WORDS	SECONDS	CUMULATIVE
hook	19	7.6	7.6s
problem	23	9.2	16.8s
explanation	25	10.0	26.8s
proof	23	9.2	36.0s
example	41	16.4	52.4s
payoff	21	8.4	60.8s
cta	22	8.8	69.6s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRA ME TYPE	LA YO UT	VISUAL	OVERL AY TEXT	TR AN SI TION	SOURCE	INSPIRATION	ASSET STATUS
0-07.6s	Same job, twice. This one is a paragraph of instructions. That one is six agents. Watch w...	hook	SCREEN_RECORDING	demo	Full-frame, filmed off a real monitor: the simple version already running, the written paragraph visible in the box above the output. A hand rests at the edge of frame. No presenter.	Same job. Two versions.	none	DbVKVz0y8xo (one continuous filmed-monitor take with a live hand, no presenter until the end); DZQZOVhgNvK (the output is on screen before any mechanism is described)	template	
7.6-16.8s	The task runs about five times a week. A quote, a summary, a reply. Whatever you build, s...	problem	SCREEN_RECORDING	demo	Same unbroken take. The output finishes and sits there while the narration names the kind of task this is.	Five times a week	none	Dbs2JdJKUDd (the pain is named as acts while the replacement is already on screen)	template	
16.8-26.8s	This is the paragraph. What the task is, what good looks like, three checks at the bottom...	explanation	SCREEN_RECORDING	demo	Same take. The hand scrolls up to the paragraph itself and stops on the three checks at the bottom of it.	One paragraph. Three checks.	ui_zoom	DbVKVz0y8xo (the hand physically points at each part as it is named)	template	
26.8-36.0s	This is the other one. Six steps, four of them fine, and this one waiting on something no...	proof	UI_DEMO	demo	Cut to the elaborate version, filmed across the room off a real monitor: a dense multi-panel build, six steps in a row, four green and one stalled with an error nobody in shot can read.	Six steps. One stuck.	hardcut	DVaBP3wDAFE (the across-the-room filmed-dashboard framing that the agent genre uses to look impressive)	template	
36.0-52.4s	We did the elaborate version to ourselves. A hundred and fifteen measurements across two...	example	DATA_VISUAL	demo	Full-frame screenshot of our own results: the strong-versus-weak table with a hundred and fifteen feature rows and a p-value column, none of them clearing the threshold, and one line at the bottom naming the single instruction that survived.	115 measurements. None separated.	hardcut	DZQZOVhgNvK (the result is put in front of the viewer before the method is explained)	template	
52.4-60.8s	Three questions before you build. Who can change it. What happens when it is wrong. How l...	payoff	A_ROLE_MEDIUM	actor	Michael, medium shot, one location, banner plate held above him. The three questions appear one at a time beside him as he says them and none of them leaves.	Change it. Wrong. How long.	hardcut	Dbs2JdJKUDd (the face is used only to interpret what the screen has already shown)	template	
60.8-69.6s	Write your task as one paragraph and run it by hand five times this week. Send this to wh...	cta	A_ROLE_MEDIUM	actor	Same take, no reframe. Michael holds still for the closing line.	One paragraph. Five runs. By hand.	none	—	template	

C-2026-09-12-08 · S00 · 0.0-7.6s · SCREEN_RECORDING · demo · ->none

...

Same job. Two versions.

[SCREEN]

Full-frame, filmed off a real monitor: the simple version already running, the written paragraph visible in the box above the output. A hand rests at the edge of frame. No presenter.

Same job, twice. This one is a paragraph of instructions. That one is six agents. Watch where each stops.

src: DbVKVz0y8xo · asset: template

C-2026-09-12-08_S00.png

C-2026-09-12-08 · S01 · 7.6-16.8s · SCREEN_RECORDING · demo · ->none

...

Five times a week

[SCREEN]

Same unbroken take. The output finishes and sits there while the narration names the kind of task this is.

The task runs about five times a week. A quote, a summary, a reply. Whatever you build, somebody has to keep it running.

src: Dbs23dJKUdd · asset: template

C-2026-09-12-08_S01.png

C-2026-09-12-08 · S02 · 16.8-26.8s · SCREEN_RECORDING · demo · ->ui_room

...

One paragraph. Three checks.

[SCREEN]

Same take. The hand scrolls up to the paragraph itself and stops on the three checks at the bottom of it.

This is the paragraph. What the task is, what good looks like, three checks at the bottom. Anyone doing the job can read every line.

src: DbVKVz0y8xo · asset: template

C-2026-09-12-08_S02.png

C-2026-09-12-08 · S03 · 26.8-36.0s · UI_DEMO · demo · ->hard_cut

...

Six steps. One stuck.

[SCREEN]

Cut to the elaborate version, filmed across the room off a real monitor: a dense multi-panel build, six steps in a row, four green and one stalled with an error nobody in shot can read.

This is the other one. Six steps, four of them fine, and this one waiting on something nobody in the room can fix.

src: DVaBP3mDAFE · asset: template

C-2026-09-12-08_S03.png

C-2026-09-12-08 · S04 · 36.0-52.4s · DATA_VISUAL · demo · ->hard_cut

...

115 measurements. None separated.

[SCREEN]

Full-frame screenshot of our own results: the strong-versus-weak table with a hundred and fifteen feature rows and a p-value column, none of them clearing the threshold, and one line at the bottom naming the single instruction that survived.

We did the elaborate version to ourselves. A hundred and fifteen measurements across two hundred and sixty-eight videos. Not one reliably separated a strong reel from a weak one. What survived was a question we could...

src: DZQZOVhgNvK · asset: template

C-2026-09-12-08_S04.png

C-2026-09-12-08 · S05 · 52.4-60.8s · A_ROLL_MEDIUM · a_roll · ->hard_cut

...

Change it. Wrong. How long.

[SHOT] medium

Michael, medium shot, one location, banner plate held above him. The three questions appear one at a time beside him as he says them and none of them leaves.

Three questions before you build. Who can change it. What happens when it is wrong. How long until it first runs.

src: Dbs23dJKUdd · asset: template


C-2026-09-12-08_S05.png

C-2026-09-12-08 · S06 · 60.8-69.6s · A_ROLL_MEDIUM · a_roll · ->none

One paragraph. Five runs. By hand.

[SHOT] medium

Same take, no reframe. Michael holds still for the closing line.



Write your task as one paragraph and run it by hand five times this week. Send this to whoever is scoping it.

src: - · asset: template

C-2026-09-12-08_S06.png

C-2026-09-12-08 · S00 · 0.0-7.6s · SCREEN_RECORDING · demo · ->none

Same job. Two versions.

...

[SCREEN]

Full-frame, filmed off a real monitor: the simple version already running, the written paragraph visible in the box above the output. A hand rests at the edge of frame. No presenter.

Same job, twice. This one is a paragraph of instructions. That one is six agents. Watch where each stops.

src: DbVKVz0y8xo, DZQZ0VhgNvK · asset: template

S00 0.0-7.6s

C-2026-09-12-08 · S01 · 7.6-16.8s · SCREEN_RECORDING · demo · ->none

Five times a week

...

[SCREEN]

Same unbroken take. The output finishes and sits there while the narration names the kind of task this is.

The task runs about five times a week. A quote, a summary, a reply. Whatever you build, somebody has to keep it running.

src: Dbs2Jd3KU0d · asset: template

S01 7.6-16.8s

C-2026-09-12-08 · S02 · 16.8-26.8s · SCREEN_RECORDING · demo · ->ui_zoom

One paragraph. Three checks.

...

[SCREEN]

Same take. The hand scrolls up to the paragraph itself and stops on the three checks at the bottom of it.

This is the paragraph. What the task is, what good looks like, three checks at the bottom. Anyone doing the job can read every line.

src: DbVKVz0y8xo · asset: template

S02 16.8-26.8s

C-2026-09-12-08 · S03 · 26.8-36.0s · UI_DEMO · demo · ->hard_cut

Six steps. One stuck.

...

[SCREEN]

Cut to the elaborate version, filmed across the room off a real monitor: a dense multi-panel build, six steps in a row, four green and one stalled with an error nobody in shot can read.

This is the other one. Six steps, four of them fine, and this one waiting on something nobody in the room can fix.

src: DVaBP3mDAFE · asset: template

S03 26.8-36.0s

C-2026-09-12-08 · S04 · 36.0-52.4s · DATA_VISUAL · demo · ->hard_cut

115 measurements. None separated.

...

[SCREEN]

Full-frame screenshot of our own results: the strong-versus-weak table with a hundred and fifteen feature rows and a p-value column, none of them clearing the threshold, and one line at the bottom naming the single instruction that survived.

We did the elaborate version to ourselves. A hundred and fifteen measurements across two hundred and sixty-eight videos. Not one reliably separated a strong reel from a weak one. What survived was a question we could...

src: DZQZ0VhgNvK · asset: template

S04 36.0-52.4s

C-2026-09-12-08 · S05 · 52.4-60.8s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Change it. Wrong. How long.

[SHOT] medium

Michael, medium shot, one location, banner plate held above him. The three questions appear one at a time beside him as he says them and none of them leaves.



Three questions before you build. Who can change it. What happens when it is wrong. How long until it first runs.

src: Dbs2Jd3KU0d · asset: template

S05 52.4-60.8s

C-2026-09-12-08 · S06 · 60.8-69.6s · A_ROLL_MEDIUM · a_roll · ->none

One paragraph. Five runs. By hand.

[SHOT] medium

Same take, no reframe. Michael holds still for the closing line.





Write your task as one paragraph and run it by hand five times this week. Send this to whoever is scoping it.

```
src: - · asset: template
```

S06 60.8-69.6s

Contact sheet

Editing

- **Cut timing:** Seven sections, three visible cuts. The first three sections are one unbroken filmed-monitor take, modelled on the corpus's highest-lift visual sequence; the presenter material at the end is a second unbroken take. There is no speaker change, which is the only cut PRODUCTION.md allows.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 40px, Paper on an Ink plate with 24px padding, bottom safe zone, two lines maximum. One persistent Ink banner plate in Saira SemiCondensed Black at the top for the whole reel. No emoji on screen.
- **Emphasis:** No colour on the elaborate run - it is not marked as bad, it is simply shown stalling. Oxide appears once, on the stalled step's own error state as the screen actually renders it, never as an added graphic.
- **Zooms:** One ui_zoom onto the three checks in the paragraph. No zoom on the elaborate run and none on any A-roll scene.
- **Motion:** Linear easing only, per design/rules/motion.md. No bounce, no gradients, no rounded corners, no centred text. No clock or timer overlay anywhere.
- **Assets required:**
 - One continuous filmed-monitor take of the simple version running end to end, with the written paragraph and its three checks visible
 - Across-the-room filmed-monitor shot of the multi-step version with one step stalled
 - Screenshot of our own strong-versus-weak feature table with the p-value column visible (real, from reports/data/associations.json and reports/analysis/04-strong-vs-weak.md)
 - Presenter A-roll, one continuous take, one location, banner plate fixed

Claims

CLAIM	STATE	SOURCE
We computed 115 numeric features across 268 analysis-ready videos, and none of them separates the top from the bottom performance quartile at $p < 0.01$.	OBSERVED	data/analysis/insights.json I-22; reports/analysis/01-general-conclusions.md §C - of 115 features, none reaches $p < 0.01$ and only seven reach $p < 0.05$.
Of 460 correlation pairs tested, 13 clear our reliability bar, eleven of them against save rate (ten distinct features) and none at all against robust_z, our creator-relative reach measure; one clears it against view_lift.	OBSERVED	reports/analysis/01-general-conclusions.md §C; data/analysis/insights.json I-22.
The one visual instruction that survived is 'put a real screen on camera'.	OBSERVED	data/analysis/insights.json I-08 - screen_share against share_rate, $\rho + 0.203$, $p = 0.00085$, creator-normalised $+0.129$, survives.
The two runs shown on screen - one paragraph of instructions and one multi-step build - are recorded on the shoot day on a task the presenters actually repeat.	PLANNED	No footage exists yet. The comparison must be in the footage or the reel becomes an opinion; the hypothesis names this as its main risk.
That a written prompt beats a multi-step build in general.	MISING	We tested nothing of the kind. The reel shows one comparison on one task and hands over three questions that decide it, never a general verdict.
How long each version takes to set up, and what each costs to maintain.	TOMEASURE	Not measured. The script names the three questions to ask rather than numbers we do not have, and no clock or timer appears on screen.
That agents are the wrong choice.	MISING	Not asserted. The closing hands over three questions; the case where the larger build is right is left open rather than argued away.
That our own measurement failure generalises to the viewer's task.	MISING	It is one experiment on one corpus, offered as a shape to recognise, not as evidence about anyone else's build.

Traceability

- **Insights:** I-22, I-25, I-08, I-09, I-27, I-16, I-05
- **Hypothesis:** H-09
- **Evidence summary:** Our own elaborate version lost: 115 numeric features over 268 analysis-ready videos, none separating the strong from the weak quartile at $p < 0.01$, only seven at $p < 0.05$, and of 460 correlation pairs just 13 clear the reliability bar with none against reach (I-22). The single instruction that survived was 'put a real screen on camera' (I-08, $\rho + 0.203$, $p = 0.00085$, survives

creator normalisation) - a question we could have asked on day one. The angle takes the energy of the corpus's highest-lift topic and refuses its shape: agent building runs at $n=27$, median view_lift 6.63, and not one of the 23 agent_build reels argues against the build (I-25; the nearest counter-example, DcbPb2yxrTx, argues against the dashboard and for the architecture underneath, not against building). The visual model is the single unbroken process shot, the highest median view_lift of any visual sequence with more than three reels behind it (30.90 on $n=6$, 4 creators; only a three-reel sequence scores higher), carried here by DbVKVz0y8xo's 87-second unbroken take at 33.0 shares and 57.6 saves per 1 000; DVaBP3wDAFE (70.4x lift) supplies the look of the elaborate run as a foil, DZQZOVhgNvK (50.2 saves/1k) the output-before-architecture ordering, Dbs2JdJKUDd (robust_z 3.28) the pain-demo-objection order.

- **Original synthesis:** The corpus contains 27 agent-build reels and 23 with an agent_build solution type, and not one of them argues against the build - the nearest thing to a counter-example, DcbPb2yxrTx ('Please stop building Jarvis dashboards', myth_bust / warning_dont), argues against the visible dashboard and for the architecture underneath it. A demonstrated 'the small version won', with the author's own overbuild shown as the evidence, is unrepresented. The three questions - who can change it, what happens when it is wrong, how long until the first version runs - are ours and appear in no reference.
- **Why better than generic:** A generic AI script would praise agents or, in the contrarian version, call them hype and show nothing. This one opens on the demo rather than the disagreement, because contrarian hooks carry the lowest share and save rates of the large hook cells (0.0076 and 0.0139), runs both versions in the footage so the comparison is watchable rather than arguable, and uses our own failed measurement as the proof - which means the reel's evidence is a thing we did badly, not a thing we are selling.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] The script said 'Not one separated a strong reel from a weak one'. reports/analysis/01-general-conclusions.md §C and I-22 say none of the 115 features reaches $p<0.01$ but seven reach $p<0.05$, so seven do separate them at the weaker bar → changed to 'Not one reliably separated a strong reel from a weak one', which is exactly the project's RELIABLE threshold. One word added; the section goes to 41 words / 16.4 s and the reel to 174 words / 69.6 s, still inside the band.
 - [evidence] claims[1] said 13 RELIABLE pairs with 'none at all against the reach measure'. Nine are save-rate and none is against robust_z, but one (lex_entities_distinct → view_lift, +0.161, cn +0.198) is against view_lift, which is also a reach measure → claim reworded to name robust_z and to state the one view_lift pair.
 - [evidence] Reference DbVKVz0y8xo was credited with 'the highest median view_lift of any visual sequence in the corpus (30.90 on $n=6$)'. reports/data/visual_patterns.json has A_ROLL_CLOSE_UP>SCREENSHOT>A_ROLL_CLOSE_UP at 52.04 on $n=3$ → qualified to 'the highest of any sequence with more than three reels behind it', in the reference reason and in traceability.
 - [evidence] The other headline numbers reproduce: 115 features, none at $p<0.01$ and seven at $p<0.05$ (I-22, §C); 460 pairs and 13 RELIABLE (§C); screen_share → share_rate +0.203, $p=0.00085$, survives creator normalisation (I-08); agent building $n=27$ at median view_lift 6.63 and 23 agent_build solution types (I-25, I-29); contrarian hooks at 0.0076 share and 0.0139

save, the lowest of the large hook cells (§D); myth_bust at robust_z 1.93 on the smallest explanation share (05-script-parts §6) → left.

- [originality] 'Not one of them argues against the build' is the hypothesis' own novelty note, but the corpus does contain DcbPb2yxrTx ('Please stop building Jarvis dashboards', myth_bust / warning_dont, agent-building topic), which argues against the dashboard while defending the architecture underneath → the counter-example is now named in traceability so the novelty claim is falsifiable rather than absolute. No script change: the claim is not spoken.
- [evidence] All four reference metric sets re-derived from data/radar.db and all four function tags match data/analysis/hypothesis_refs.json; the four quotes are verbatim from the beat ids given; the script shares no five-word sequence with any of the four transcripts → left.
- [evidence] The demo is PLANNED. If the two runs are not both recorded the card must not be shot - the hypothesis names 'sounds like an opinion' as its main risk → left.
- [hook] Demo-first rather than contrarian, per the hypothesis reviewer's note; the disagreement is shown and never spoken; three candidates, one question, first one used; 19 words / 7.6 s → left.
- [positioning] No clock or timer overlay anywhere, because a speed comparison we have not measured would be a claim (recorded TO_MEASURE); the elaborate run is filmed to look as impressive as the genre does and the stalled step is not dressed as a failure → left.
- [positioning] No split screen, no comment gate, no income or savings framing, no vendor named → left.
- [visual] Coarse sequence SCR>A: the first three sections are one unbroken filmed-monitor take and the presenter appears only for the verdict, which is pattern 4 in reports/analysis/03-visual-patterns.md. Opens on a running screen at 0.0 s (I-09) → left.
- [timing] Recomputed: 174 words, 69.6 s at 2.5 wps, storyboard contiguous 0.0-69.6 s, no scene under 1.5 s, every overlay ≤ 7 words → passes.
- [schema] python3 -m engine.cards_v2 validate → 0 errors, 0 warnings after the edits.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-08.json
- **Remotion validation:** {"card_id": "C-2026-09-12-08", "checked_at": "2026-09-12T09:13:54Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2088, "edl": "cards/edl/C-2026-09-12-08.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-09 — Supported is not the same as connected

Strategy

Format: M2 Radar

Concept. A tool is only as useful as the systems it can already reach, and on Thursday our own weekly run proved it twice: one source went silent, and one destination that had answered three days earlier stopped answering because it had been moved at the other end.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company who is responsible for one repeated process and has tried ChatGPT but has no repeatable way of using it at work. Specifically anyone about to trial a tool for a process that already lives in an inbox, a spreadsheet or a calendar.

Objective. Get the viewer to open the integrations page before the feature page, check three named systems against it before a trial starts, and check again that each one is still answering.

Positioning. builder

Pain. manual_repetition

Promise. A two-minute check, before you trial anything, that tells you whether the tool can touch the work you actually do.

Rationale. Naming specific things is the only lexical feature that predicts reach and survives creator normalisation (distinct entities vs view_lift +0.161 pooled, +0.198 within creators), and this reel is built entirely out of names. The references supply grammar: each destination introduced at the moment the workflow needs it (DNA7-d3o5sQ), two identically shaped steps each closing on a concrete count (DZU9x_1AY2t), a connector being attached on screen in real time (Db9MT-axxcO), and the everybody-wants-X-but-nobody-does-Y opening with the friction named as a concrete act (DVZGBgTk6zz).

FIL ANSWER
TER

Process	Any repeated task whose inputs and outputs already sit in an inbox, a spreadsheet and a calendar - taking an enquiry, recording it, and putting a slot in the diary for it.
Friction	A tool that cannot read and write those three systems turns the work into copying and pasting, which is more manual work than there was before. And 'supported' on a feature page is not the same as connected: a connection can be live on Monday, be removed at the other end, and be gone by Thursday without anything on the feature page changing.
AI boundary	The tool moves and drafts between named systems. The human owns the account access, decides what the tool may see, and stays responsible for whether a connection is actually open rather than merely listed.
Next action	Write down the three systems your process actually touches and check them, by name, against the tool's own integrations page before the trial starts.

Hooks

1. **[demo_first]** (statement) — **USED IN SCRIPT** Watch an enquiry go from the inbox to the sheet to the calendar. Each name is where a tool stops.

1. **[bold_claim]** (statement) The integrations page decides whether a tool touches your work. The feature page does not.

1. **[question]** (question) Can the tool you are about to trial even see your inbox?

Reference reels

REEL	CRE ATO R	FU NC TI ON	PL AY S	CRE ATO R MED IAN	VI EW LI FT	SHA RE RAT E	SAV E RAT E	REASON
DNA7-d3o5sQ (https://www.instagram.com/reel/DNA7-d3o5sQ/)	edhillai	expl	45 1,4 91	8,01 0	55. 4x	2.3 %	3.7 %	Rated the best overall template in the pool and the only ungated reel in it (cta_type none, so its 23.3 shares per 1 000 carry no keyword inflation). Its mechanism beats are a connector list read aloud in dependency order - two named ordinary systems, each introduced by the reason it is needed, which is exactly the grammar a connector reel needs. Screen-first opening held for 56 seconds.
DZU9x_1AY2t (https://www.instagram.com/reel/DZU9x_1AY2t/)	badarmunir_official	rhythm	1,2 82, 45 8	6,67 8	19 1.0 x	2.6 %	4.8 %	The densest reel in the pool - 37 seconds, no intro, no story, two identically shaped steps, five named artefacts - and the line that carries it is pure specificity: three exact numbers in one breath. 25.8 shares and 48.1 saves per 1 000.
Db9MT-axxcO (https://www.instagram.com/reel/Db9MT-axxcO/)	shrugman	screenplay	37, 79 1	14,4 48	1.6 x	4.8 %	6.9 %	Best combined share and save rates in the corpus (47.9 and 68.9 per 1 000) on a reel that barely beat its own author's median (1.62x, robust_z 1.08), which makes it the purest available example of resonance without reach - forwarding behaviour rather than algorithmic luck. Its mechanism beat is a connector literally being attached on screen.
DVZGBgTk6zz (https://www.instagram.com/reel/DVZGBgTk6zz/)	shareholder	hook	41 9,0 00	1,37 4	30 3.8 x	3.5 %	4.6 %	Names an objection the audience has never heard said out loud and then repeats it back solved: one sentence of friction, one sentence of promise, 6.8 seconds, inside the corpus's 5-10 s hook norm. The 303.8x rests on a 1 374-view author median, so it is read as a curiosity and the sentence shape is the finding.

DNA7-d3o5sQ (edhillai, explanation)

- Transcript (DNA7-d3o5sQ-B1): “This agent understands your goal but needs data. Therefore, it uses Google Calendar to check your schedule and Google Sheets to save any details.”
- Frame (DNA7-d3o5sQ-S00): frame one is the workflow canvas already listening, held unbroken for 56 seconds; the presenter is a small head at the bottom of frame and never takes the picture
- Transformed: Borrowed naming each destination at the moment the workflow needs it, and opening on the artefact doing its own first move rather than on a presenter. Changed: our subject is whether the tool can reach those systems at all, not how to assemble an agent; the owner of the process is named, which this reel never does; the systems shown are an inbox, a spreadsheet and a calendar rather than a voice stack; and our presenter is absent from the frame entirely until 50 s rather than sitting under it, because any split state runs at median robust_z 1.03 against 1.43.

DZU9x_1AY2t (badarmunir_official, rhythm)

- Transcript (DZU9x_1AY2t-B1): “This gives Claude a database of over 50 UI styles, 97 color palettes and 57 font pairings.”

- Frame (DZU9x_1AY2t-S04): the install log filling the frame with the success line highlighted, so the claim and its receipt are the same image
- Transformed: Borrowed the cadence - two steps of identical shape, each closing on a concrete count - which is how we name three destinations and two failures without the middle sagging. Changed: the numbers are our own and describe a process rather than a product's feature inventory, and we do not copy its visual sequence, which alternates into split screen four times.

Db9MT-axxcO (shrug.manny, screen_proof)

- Transcript (Db9MT-axxcO-B1): “First, hit the plus, add the Higgs Field Connector, upload your resume, any inputs and the dream job that you want to land.”
- Frame (Db9MT-axxcO-S01): full-frame capture of the connector menu itself, each connector listed with its toggle visibly on or off - the settings screen, not a marketing page
- Transformed: Borrowed the act of showing a connection being made or refused on screen in real time. Changed: the connectors are the viewer's existing inbox, sheet and calendar rather than a sponsor's product - our own analysis records this reel as an undisclosed tool partnership, which is exactly what we must not reproduce - and we claim no outcome number of any kind, where it claims six interviews in 24 hours.

DVZGBgTk6zz (sharmadhruveshh, hook)

- Transcript (DVZGBgTk6zz-B0): “Everyone using cloud code to build AI agent but nobody wants to touch terminal. Here's the easiest way to set this up under 5 minutes.”
- Frame (DVZGBgTk6zz-S01): cuts to the tool's own splash screen animating in on the second word, so the subject is on screen before the sentence finishes
- Transformed: Borrowed the shape - the friction named as a concrete act, immediately followed by the promise - and the habit of putting the subject on screen before the sentence ends. Changed: our friction is an ordinary business fact (the tool cannot see your inbox) rather than a developer's discomfort with a terminal, we name no tool in the hook, and the 1 374-view author median means we take the sentence shape and none of the performance claim.

Script

Watch an enquiry go from the inbox to the sheet to the calendar. Each name is where a tool stops. Nothing clever happens here. It reads the inbox, writes a row in the sheet, puts a slot in the calendar. Three names to check before you pay. Our own weekly run is a chain like that. Six named places. On Thursday two of them refused us. One source went quiet five times in a row, so none of the hundred and thirty profiles updated. And the notes tool said the same thing three times: cannot find the database. Three days earlier that same connection worked. Someone moved the destination and the tool lost it. The one place nobody moved saved everything, first time. Supported is not the same as connected, and connected is not the same as still connected. Write down the three systems your work lives in and check them by name. Where a name is missing, that step stays manual. Send this to whoever is picking the tool.

170 words / 68.0s (2.5 wps target) — tone: plain, concrete, unhurried - a chain named link by link; emotional effect: the relief of having one checkable question to ask before committing to a tool

ROLE	WORDS	SECONDS	CUMULATIVE
hook	20	8.0	8.0s
explanation	27	10.8	18.8s
proof	19	7.6	26.4s
proof	32	12.8	39.2s
explanation	25	10.0	49.2s
payoff	30	12.0	61.2s
cta	17	6.8	68.0s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYERS	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE	INSPIRATION	ASSET STATUS
00:08:00s	Watch an enquiry go from the inbox to the sheet to the calendar. Each name is where a too...	hook	UI_D EMO	demo	Full frame, no presenter: one enquiry arrives in an inbox, a row appears in a spreadsheet, a slot appears in a calendar. Real screens, one continuous move.	Inbox. Sheet. Calendar.	none	DVZGBgTk6zz (the subject is on screen before the opening sentence finishes); DNA7-d3o5sQ (frame one is the artefact already doing its first move, no establishing shot of a person)	template	
00:08:08.8s	Nothing clever happens here. It reads the inbox, writes a row in the sheet, puts a slot i...	exposition	UI_D EMO	demo	The same take continues; each of the three destinations gets its name added to the top banner as the item lands in it.	Three names to check	none	DNA7-d3o5sQ (each destination named at the moment the workflow needs it, in dependency order)	template	
00:18:26.4s	Our own weekly run is a chain like that. Six named places. On Thursday two of them refuse...	proof	SCREEN_R ECOR DING	demo	Full-frame capture of our own run log, scrolled to the top, the six named steps visible in order.	Six named places. Two refused.	hardcut	DZU9x_1AY2t (the claim and its receipt are the same image - the log itself, not a description of it)	template	
00:26:39.2s	One source went quiet five times in a row, so none of the hundred and thirty profiles upd...	proof	SCREEN_R ECOR DING	demo	The log continues: the five silent attempts and the zero-of-130 line, then the same 404 message repeated at three different steps.	0 of 130 updated. 404, three times.	none	DZU9x_1AY2t (two identically shaped steps, each closing on a concrete count)	template	
00:39:24.9s	Three days earlier that same connection worked. Someone moved the destination and the too...	exposition	SCREENSHOT	demo	Full-frame: the same step in two runs three days apart - Monday's log line where the destination answered, and Thursday's 404 in its place, beside the one destination that answered in both.	Worked Monday. Gone Thursday.	hardcut	Db9MT-axxcO (the settings screen where a connection is actually on or off, shown rather than described)	template	
00:49:26.1s	Supported is not the same as connected, and connected is not the same as still connected....	payoff	A_ROLL M EDIUM	roll	Max, medium shot, plain wall, first appearance in the reel.	Supported. Connected. Still connected.	hardcut	DNA7-d3o5sQ (the reel hands back to a human only once the artefact has finished making the argument)	template	
00:61:28.0s	Where a name is missing, that step stays manual. Send this to whoever is picking the tool.	cta	A_ROLL M EDIUM	roll	Max, same take, closing line to camera.	Missing name, manual step	none	Db9MT-axxcO (returns to the face for the closing line after the screen has carried the mechanism)	template	

C-2026-09-12-09 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

Inbox. Sheet. Calendar.

...

[SCREEN]

Full frame, no presenter: one enquiry arrives in an inbox, a row appears in a spreadsheet, a slot appears in a calendar. Real screens, one continuous move.

Watch an enquiry go from the inbox to the sheet to the calendar. Each name is where a tool stops.

src: DVZ68gTk6zz, DNA7-d3o5sQ · asset: template

C-2026-09-12-09_S00.png

C-2026-09-12-09 · S01 · 8.0-18.8s · UI_DEMO · demo · ->none

Three names to check

...

[SCREEN]

The same take continues; each of the three destinations gets its name added to the top banner as the item lands in it.

Nothing clever happens here. It reads the inbox, writes a row in the sheet, puts a slot in the calendar. Three names to check before you pay.

src: DNA7-d3o5sQ · asset: template

C-2026-09-12-09_S01.png

C-2026-09-12-09 · S02 · 18.8-26.4s · SCREEN_RECORDING · demo · ->hard_cut

Six named places. Two refused.

...

[SCREEN]

Full-frame capture of our own run log, scrolled to the top, the six named steps visible in order.

Our own weekly run is a chain like that. Six named places. On Thursday two of them refused us.

src: DZU9x_1AY2t · asset: template

C-2026-09-12-09_S02.png

C-2026-09-12-09 · S03 · 26.4-39.2s · SCREEN_RECORDING · demo · ->none

0 of 130 updated. 404, three times.

...

[SCREEN]

The log continues: the five silent attempts and the zero-of-130 line, then the same 404 message repeated at three different steps.

One source went quiet five times in a row, so none of the hundred and thirty profiles updated. And the notes tool said the same thing three times: cannot find the database.

src: DZU9x_1AY2t · asset: template

C-2026-09-12-09_S03.png

C-2026-09-12-09 · S04 · 39.2-49.2s · SCREENSHOT · demo · ->hard_cut

Worked Monday. Gone Thursday.

...

[SCREEN]

Full-frame: the same step in two runs three days apart - Monday's log line where the destination answered, and Thursday's 404 in its place, beside the one destination that answered in both.

Three days earlier that same connection worked. Someone moved the destination and the tool lost it. The one place nobody moved saved everything, first time.

src: Db9MT-axxc0 · asset: template


C-2026-09-12-09_S04.png

C-2026-09-12-09 · S05 · 49.2-61.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Supported. Connected. Still connected.

[SHOT] medium

Max, medium shot, plain wall, first appearance in the reel.



Supported is not the same as connected, and connected is not the same as still connected. Write down the three systems your work lives in and check them by name.

src: DNA7-d3o5sQ · asset: template


C-2026-09-12-09_S05.png

C-2026-09-12-09 · S06 · 61.2-68.0s · A_ROLL_MEDIUM · a_roll · ->none

Missing name, manual step

[SHOT] medium

Max, same take, closing line to camera.



Where a name is missing, that step stays manual. Send this to whoever is picking the tool.

src: Db9MT-axxc0 · asset: template

C-2026-09-12-09_S06.png

C-2026-09-12-09 · S00 · 0.0-8.0s · UI_DEMO · demo · ->none

Inbox. Sheet. Calendar.

...

[SCREEN]

Full frame, no presenter: one enquiry arrives in an inbox, a row appears in a spreadsheet, a slot appears in a calendar. Real screens, one continuous move.

src: DVZGBgTk6zz, DNA7-d3o5sQ · asset: template

S000.0-8.0s

C-2026-09-12-09 · S01 · 8.0-18.8s · UI_DEMO · demo · ->none

Three names to check

...

[SCREEN]

The same take continues; each of the three destinations gets its name added to the top banner as the item lands in it.

src: DNA7-d3o5sQ · asset: template

S018.0-18.8s

C-2026-09-12-09 · S02 · 18.8-26.4s · SCREEN_RECORDING · demo · ->hard_cut

Six named places. Two refused.

...

[SCREEN]

Full-frame capture of our own run log, scrolled to the top, the six named steps visible in order.

src: DZU9x_1AY2t · asset: template

S0218.8-26.4s

C-2026-09-12-09 · S03 · 26.4-39.2s · SCREEN_RECORDING · demo · ->none

0 of 130 updated. 404, three times.

...

[SCREEN]

The log continues: the five silent attempts and the zero-of-130 line, then the same 404 message repeated at three different steps.

src: DZU9x_1AY2t · asset: template

S0326.4-39.2s

C-2026-09-12-09 · S04 · 39.2-49.2s · SCREENSHOT · demo · ->hard_cut

Worked Monday. Gone Thursday.

...

[SCREEN]

Full-frame: the same step in two runs three days apart - Monday's log line where the destination answered, and Thursday's 404 in its place, beside the one destination that answered in both.

src: Db9MT-axxcQ · asset: template

S0439.2-49.2s

C-2026-09-12-09 · S05 · 49.2-61.2s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Supported. Connected. Still connected.

[SHOT] medium

Max, medium shot, plain wall, first appearance in the reel.



src: DNA7-d3o5sQ · asset: template

S0549.2-61.2s

C-2026-09-12-09 · S06 · 61.2-68.0s · A_ROLL_MEDIUM · a_roll · ->none

Missing name, manual step

[SHOT] medium

Max, same take, closing line to camera.





Where a name is missing, that step stays manual. Send this to whoever is picking the tool.

src: Db9MT-axxc0 · asset: template

S06 61.2-68.0s

Contact sheet

Editing

- **Cut timing:** Three hard cuts: into our own run log at 18.8 s, into the two runs side by side at 39.2 s, and to the presenter at 49.2 s. The opening frame is not a cut. Everything before 18.8 s is one capture, the log block is one capture, and the closing face block is one take.
- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 600 at 40px, two lines maximum, bottom-anchored in Instagram zone A/B; one persistent top banner that accumulates the three destination names and then holds.
- **Emphasis:** Every named system on the banner is one the viewer already has. The two failures are quoted from the log exactly as the log wrote them, and no tool is named in the first three seconds.
- **Zooms:** None. The log is captured at a size that reads without a zoom.
- **Motion:** Linear easing only, no bounce, per design/rules/motion.md.
- **Assets required:**
 - Continuous screen capture of one enquiry moving inbox to spreadsheet to calendar, our own accounts (S00-S01)
 - Capture of the 2026-09-10 run log showing the six named steps (S02)
 - Capture of the same log at the five silent attempts, the 0-of-130 line, and the three repeated 404 messages (S03)
 - Capture of the 2026-09-07 and 2026-09-10 logs side by side: the same step answering on Monday and 404-ing on Thursday, beside the destination that answered in both (S04)
 - Max A-roll, medium shot, one location, one take

Claims

CLAIM

STATUS

SOURCE

Our weekly run is a chain of six named systems: a paid data source, a free public-profile source, local transcription, local frame extraction, a database file and a notes tool.

OB
SE
RV
ED

engine/SPEC.md section 0.3 production default path plus run.py's Notion and follower steps; each step of the 2026-09-10 07:00 run is quoted in reports/audit/03-server-runtime-hiker.md section 3, which quotes the run line by line; the log itself is on the server at /opt/radar/data/runs/2026-09-10_0700.log and is not in this repository (data/runs/ does not exist locally).

In the run of 2026-09-10 the free public-profile source returned nothing five times in a row and 0 of 130 profiles were updated.

OB
SE
RV
ED

Step 5 of that run, quoted in reports/audit/03-server-runtime-hiker.md section 3, which quotes the run line by line; the log itself is on the server at /opt/radar/data/runs/2026-09-10_0700.log and is not in this repository (data/runs/ does not exist locally).

Three steps of the same run failed with the same Notion 404 message asking that the destination be shared with the integration - steps 7 and 10 on the Cards database, step 12 on the Reels database.

OB
SE
RV
ED

Steps 7, 10 and 12, quoted in reports/audit/03-server-runtime-hiker.md section 3, which quotes the run line by line; the log itself is on the server at /opt/radar/data/runs/2026-09-10_0700.log and is not in this repository (data/runs/ does not exist locally). Steps 7 and 10 name d5b6ab42-...; step 12 names a different id, 679fa6d6-..., so it is the same message on two destinations, not one.

In the same run the step that wrote to a page rather than to a database succeeded.

OB
SE
RV
ED

Step 11 of that run ('the stats block was updated'), quoted in reports/audit/03-server-runtime-hiker.md section 3, which quotes the run line by line; the log itself is on the server at /opt/radar/data/runs/2026-09-10_0700.log and is not in this repository (data/runs/ does not exist locally). It writes to a page rather than to a database.

The destination had been reachable three days earlier - steps 7 and 10 of the 2026-09-07 run both succeeded - and was gone by 2026-09-10. Our audit attributes that to the databases being moved to trash around 7 September, not to a connection that was listed and never made.

OB
SE
RV
ED

reports/audit/03-server-runtime-hiker.md section 3 (09-07 run: step 7 'statuses updated: 5', step 10 '9 of 9 uploaded') and section 4, which reads the 404 text as Notion's generic 'the integration lost access' message and matches it to the three databases being moved to trash around 7 September. The 404 body names a remedy, not a cause.

What the integrations page said at any point is not something we checked; the reel makes no claim about it and names no vendor.

MI
SSI
NG

No integrations page was opened during the audit. Recorded so 'supported is not the same as connected' is read as the lesson from a connection that stopped, not as a finding about a published list.

We will publish the two-minute pre-trial check - name the three systems, find them on the integrations page, mark the missing ones as manual - as a written list.

PL
AN
NE
D

Not written yet; the reel does not offer it as an existing download.

CLAIM	STATE	SOURCE
Running that check before a trial reduces how often a tool turns out to be unusable for the process.	TO_M EA SU RE	We have not run a trial with and without the check.
A named commercial tool's integrations page is accurate about what it can reach in a real small business.	MI SSI NG	We have only tested our own chain. No commercial tool's integrations page has been checked against a real company's inbox, spreadsheet and calendar, and the reel names no tool.

Traceability

- **Insights:** I-09, I-13, I-08, I-15, I-02, I-05
- **Hypothesis:** H-05
- **Evidence summary:** Our own weekly run is a chain of six named systems and on 2026-09-10 two of the links failed for two different reasons. The free public-profile source returned nothing five times in a row and 0 of 130 profiles were updated. Three separate steps failed with the same 404 from the notes tool - two on one database, one on another - while the one step in the same run that wrote to a page rather than a database succeeded. The same destination had answered three days earlier, on 2026-09-07, when steps 7 and 10 both went through: the connection existed and was removed at the other end. The four references contribute grammar only: destinations named in dependency order over a 56-second unbroken screen take (DNA7-d3o5sQ, ungated, 23.3 shares/1k), two identically shaped steps each closing on a count (DZU9x_1AY2t, 25.8 shares/1k), a connection shown being made on the settings screen (Db9MT-axxcO, best combined 47.9 and 68.9 per 1k on only 1.62x lift) and a friction-then-promise opening (DVZGBgTk6zz, 34.8 shares/1k).
- **Original synthesis:** The corpus names tools constantly and integrations almost never; the connector list is treated as plumbing rather than as the buying decision. The failure is ours, dated and quotable from two run logs three days apart, and it is the more interesting failure mode because the connection was live on Monday and gone by Thursday with nothing on the tool's side changing. The reel names no vendor, makes no time or cost claim, and hands over a check rather than a tool recommendation.
- **Why better than generic:** A generic AI script would list five tools that integrate with Gmail. This one shows a single item crossing three ordinary systems, then produces two real logs three days apart in which the same supported connection worked and then did not, for a reason that has nothing to do with the model - and it separates supported, connected and still connected, which is the sentence the viewer can use in a sales call. Naming specific things is the only lexical feature that predicts reach and survives creator normalisation (+0.161 pooled, +0.198 within creators), and every load-bearing noun in this script is a name.

Review

- **Version:** 1
- **Revised:** yes

- **Findings:**

- [evidence] HARD - the reel's central claim contradicted our own audit. The script said 'It was on the list of things that tool supports. The database was simply not shared with it', and claims[4] said the 404 body 'names the cause'. reports/audit/03 section 4 reads it the other way: the three databases were moved to trash around 7 September and the 404 text is Notion's generic 'the integration lost access' message. Section 3 shows the same destination answering on 2026-09-07 (step 7 'statuses updated: 5', step 10 '9 of 9 uploaded'), so the connection existed and was removed - it was never listed-but-never-made. -> explanation section rewritten to 'Three days earlier that same connection worked. Someone moved the destination and the tool lost it.', claims[4] rewritten, strategy.concept and filter.friction corrected, S04's visual and overlay rebuilt around the two runs side by side. Script version 2; total_s 68.8 -> 68.0.
- [evidence] the payoff carried the whole lesson on one distinction that no longer fits -> 'Supported is not the same as connected, and connected is not the same as still connected', which is what the two logs actually show. Title kept.
- [evidence] 'Five named places' counted the follower source and the paid data source as one link, but they are different systems: step 1 (Hiker, 130 accounts, \$2.60) succeeded and step 5 (freeprofile.py, the free public-profile source) is the one that went silent -> six named places, in the script, in claims[0], in S02's visual and overlay, and in assets_required.
- [evidence] claims[2] said 'the same 404' on one database. Steps 7 and 10 are the Cards database (d5b6ab42-...); step 12 is the Reels database (679fa6d6-...) -> claim reworded to the same message on two destinations. The spoken line, 'the notes tool said the same thing three times', is accurate and stays.
- [traceability] every source cited data/runs/2026-09-10_0700.log, which is not in this repository - data/runs/ does not exist locally and the log is on the server -> all four sources now cite reports/audit/03 section 3, which quotes the run, and name the server path.
- [evidence] added a MISSING claim: no integrations page was ever opened, so the reel makes no claim about what any published list says.
- [visual] the opening frame carried transition hard_cut against the card's own 'three hard cuts' description -> set to none.
- [structure] left as is, noted: the script has no problem or friction section, against general-conclusions section G rule 11 and POSITIONING section 8's second filter item. The friction is carried by the hook's second sentence and by S04, and adding a section would push the reel past 70 s.
- [evidence] verified unchanged: five silent attempts and 0 of 130 profiles, three 404s, the page-destination step succeeding - all quoted verbatim in reports/audit/03 section 3. 2026-09-10 is a Thursday and 2026-09-07 a Monday.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-09.json
- **Remotion validation:** {"card_id": "C-2026-09-12-09", "checked_at": "2026-09-12T08:53:28Z", "contract": "m2.remotion-edl.v1", "duration_frames": 2040, "edl": "cards/edl/C-2026-09-12-09.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}

C-2026-09-12-10 — Twenty-three plans, nothing made

Strategy

Format: M2 Teardown

Concept. We take our own weekly planning ritual apart step by step on screen, with the artefact each step produces and the real count beside it, and show that the leak is not in any of the four automated steps but in the fifth - the handover, which has no name, no owner and, in three weeks, nothing recorded against it at all.

Audience. Non-technical small-business leader, AI level 0-6 of 10 (POSITIONING.md §4): an owner, operations lead or team lead in a 5-50 person company responsible for one repeated process. Specifically whoever runs the weekly planning ritual.

Objective. Get one viewer to take last week's list, write an owner's name next to every line, and count how many lines come back blank.

Positioning. builder

Pain. manual_repetition

Promise. A five-step teardown of a weekly ritual, with the real counts, and the exact step where a plan stops being anybody's job.

Rationale. The corpus states a problem in only 64 of 264 reels and never takes apart a process the author runs (I-17). RULES.md §2 names our own processes as the only Teardown source available from day one, and the format's second pass criterion - a reason to forward - is carried by the count itself: twenty-three plans, nothing produced. The reel opens on the artefact rather than on a face, because openings spent on the audience carry median robust_z 0.67 against 1.36 (I-06).

FILT ANSWER
ER

Proce ss	The recurring weekly planning meeting or run that produces a list of things to do - the Monday list, the sprint board, the shortlist of jobs for the week - in any business where the same ritual repeats and produces the same kind of list.
Frictio n	All the effort goes into producing the list, and the list is not where the process fails. It fails at the handover, where no name is attached, so nothing downstream can start and nobody notices that nothing did.
AI bound ary	The machine can collect, score, shortlist and draft. Deciding which items are real and who owns each one stays with a person - in our week that is one of us on Monday - and no step downstream may begin without that name attached.
Next action	Take last week's list, write the owner's name next to every item, and count how many are blank.

Hooks

1. **[demo_first]** (statement) — **USED IN SCRIPT** This is what our Monday produces. Twenty-three of these in three weeks. Not one of them was ever made.
1. **[result_first]** (statement) Our weekly planning run has produced twenty-three plans in three weeks and the table that records a finished one is empty.
1. **[question]** (question) Where does your Monday list actually stop? Ours stops at the step nobody has a name for.

Reference reels

REEL	CRE ATO R	FU NC TI ON	P L A Y S	CRE ATO R IAN	VI EW LI FT	SH AR E RA TE	SA VE RA TE	REASON
DYuivWz-Sj5k (https://www.instagram.com/reel/DYuivWz-Sj5k/)	aifor	busi ness 47	ho ok 69 7	11 3,42 3, 1	32. 2x %	5.1 %	4.7 %	50.8 shares per 1 000 - fourth-highest in the 40-reel pool, where 58.4, 53.4 and 52.9 are higher - from 59 spoken words, no tool, no step and no number, which makes it the purest available test of an opening that is nothing but a recognisable morning. 32.2x lift on a 3 421 median, so the multiplier is a curiosity and the forwarding rate is the finding. Its contact sheet is a single held DATA_VISUAL: an animated pipeline with live counters accumulating under each node.
DZLDk9ySi-7 (https://www.instagram.com/reel/DZLDk9ySi-7/)	aifor	busi ness 47	pai n 09 3	17 3,42 9, 1	51. 4x %	4.2 %	4.5 %	The corpus's best example of pain named as acts rather than as a condition, delivered as a four-line list at 26-33s: every line is a thing a person does with their hands, which is the framing the demand analysis measures as the strong one (manual_repetition median view_lift 6.94 against chaos_no_process 0.87). 41.9 shares and 45.4 saves per 1 000. Its contact sheet gives us the exact stacked-step grammar: a vertical list of named steps revealed one at a time, each with its own count on the right.
Db9MT-axxcO (https://www.instagram.com/reel/Db9MT-axxcO/)	shru g.m anny	rhy th m	37 14,4 7 91	1.6 x %	4.8 %	6.9 %		Result at 3 seconds, three steps in the next thirty, ask at forty-three - simultaneously a story and a procedure, which is what a process teardown has to be. Third-best combined share and save rates in the analysis-ready corpus (47.9 and 68.9 per 1 000; Db04zvJR1k8 at 140.3 and Dbj-yxNgrhG at 117.3 are higher) on a lift of only 1.62x, so the forwarding is the signal and the reach is not. Its contact sheet shows the artefact held up in hand as a phone screen, with the presenter's hand still in frame, rather than cut to as a clean insert.
DbbA7ww-TIz8 (https://www.instagram.com/reel/DbbA7wwTIz8/)	heys teve tan	cta 7	29 13,0 4, 09 19 7	21. 6x %	2.5 %	3.1 %		Narrows an enormous market claim down to one industry, one job, one business and one month, which is the best next-action shape in the pool: the ask occupies the final fourteen seconds and bounds the action on four axes at once. 24.6 shares per 1 000, 21.6x lift, robust_z 5.0 clipped. Its contact sheet shows the closing delivered over plain black cards carrying one bold word at a time rather than over the face.

DYuivWzSj5k (aiforbusinesses247, hook)

- Transcript (DYuivWzSj5k-B0): “Think, when you wake up in the morning, the clients are already interested in your service.”
- Frame (DYuivWzSj5k-S00): one held pipeline diagram for the entire runtime, six named nodes in a row with live counters climbing underneath them
- Transformed: Borrowed: the recognisable-weekday frame as the entry point to a process teardown, and the node-with-a-counter treatment. Changed: the direction is reversed - we open on the artefact our Monday actually produced rather than on an imagined better one, because openings spent on the audience rather than the subject carry median robust_z 0.67 against 1.36 (p=0.002); and our last counter stays at zero, where theirs climbs. Its single held state and its comment gate are not borrowed.

DZLDk9ySi-7 (aiforbusinesses247, pain)

- Transcript (DZLDk9ySi-7-B3): “No manual data scraping No cold outreach headache No manually demo websites No hours waste to find leads”
- Frame (DZLDk9ySi-7-S01): vertical stack of named workflow steps revealed one row at a time, each row carrying its own live count on the right-hand side
- Transformed: Borrowed: the pain list written as acts, and the stacked-step-with-a-count layout. Changed: our list is the five steps of our own Monday with the real counts attached, it sits in the first quarter as a friction statement rather than late as a benefit stack, and the counts are measurements rather than an animation. The author median is 3 421, so its 51x multiplier is not a finding and is not cited as one. Different function from this reel's use in H-11.

Db9MT-axxcO (shrug.manny, rhythm)

- Transcript (Db9MT-axxcO-B0): “What happens when you ask Claw to apply to 500 jobs for you? I tried this and I landed six interviews in 24 hours. Here's what I would do if I was looking for a job.”
- Frame (Db9MT-axxcO-S00): the outcome artefact held up in hand as a real screen with the hand still in frame, alternating with the face as the steps are named
- Transformed: Borrowed: the beat clock - the outcome inside the first five seconds, the steps in a single uninterrupted run, the ask in the last tenth - and the artefact-in-hand framing. Changed: our outcome is the unflattering one, twenty-three plans and nothing produced, and we ask no question in the hook, because question hooks sit at 0.0054 share and 0.0198 save against the corpus medians. Different function from this reel's use in H-05.

DbbA7wwTIZ8 (heystevetan, cta)

- Transcript (DbbA7wwTIZ8-B4): “So pick one industry you already understand. Find that one job nobody wants to do and build that for one local business this month.”
- Frame (DbbA7wwTIZ8-S04): the next action is delivered over plain black text cards carrying one bold word each, not over stock footage or a face
- Transformed: Borrowed: bounding the next action on several axes at once so it is small enough to do this week. Changed: ours is one action on one artefact the viewer already has - write the owner's name next to each item on last week's list and count the blanks - and the comment keyword that follows it in the original is dropped. We also do not borrow its stock b-roll or its third-party claim, which our own analysis records as contradicted by the reel's own caption.

Script

This is what our Monday produces. Twenty-three of these in three weeks. Not one of them was ever made. Five steps. Collect, score, shortlist, strike out, hand over. Four of them run without anyone touching anything. The whole machine run took sixty-five minutes. Collecting, scoring, shortlisting, all of it. Striking out is the one step that waits for one of us, and it removed nine of the twenty-three. Then the fifth step. Hand over. The table that would record a made item has nothing in it at all. And three of the twenty-three are the same source proposed again a week later, because nobody had marked the first one as used. The process is not broken where the money went. It is broken at the step where a plan is supposed to become somebody's job, and that step has no name and no owner. Take last week's list, write a name next to every line, and count the blanks. Send this to whoever runs your Monday.

166 words / 66.4s (2.5 wps target) — tone: plain, exact, willing to be the example; a count read out rather than a confession; emotional effect: the discomfort of recognising your own Monday, and the relief of being handed one column to add

ROLE	WORDS	SECONDS	CUMULATIVE
hook	19	7.6	7.6s
setup	17	6.8	14.4s
explanation	32	12.8	27.2s
proof	43	17.2	44.4s
payoff	33	13.2	57.6s
cta	22	8.8	66.4s

Storyboard

TIME	SCRIPT (TRUNCATED)	ROLE	FRAME TYPE	LAYOUT	VISUAL	OVERLAY TEXT	TRANSITION	SOURCE INSPIRATION	ASSET STATUS
00:07.6s	This is what our Monday produces. Twenty-three of these in three weeks. Not one of them w...	hook	SCREENSHOT	demo	Full-frame, filmed off a real screen and scrolling: our own plan table, twenty-three rows across three weeks, every row filled in and the produced column empty on every one of them. No presenter.	23 plans. Three weeks.	none	Db9MT-axxcO (the outcome artefact is shown in the first seconds as a real screen, hand still in frame); DYuivWzSj5k (the process artefact carries the opening instead of a presenter)	template
00:14.4s	Five steps. Collect, score, shortlist, strike out, hand over. Four of them run without an...	setup	MOTION GRAPHIC	motion graphic	Full-frame vertical stack of the five step names, revealed one row at a time, left-aligned on the grid. The first four rows are set in Graphite; the fifth stays empty for now.	Collect. Score. Shortlist. Strike. Hand over.	hardcut	DZLDk9ySi-7 (a vertical stack of named steps revealed one row at a time, each with room for its own count)	template
00:27.2s	The whole machine run took sixty-five minutes. Collecting, scoring, shortlisting, all of...	explanation	MOTION GRAPHIC	motion graphic	The same stack. The four automated rows are bracketed together against one figure - sixty-five minutes for the whole machine run - and the strike-out row carries nine struck of twenty-three. The fifth row's figure column is still blank.	65 minutes. Nine struck.	overlay	DYuivWzSj5k (a named node with its own counter underneath it, read in sequence); DZLDk9ySi-7 (each step carries its count on the right rather than in a separate summary)	template
00:44.4s	Then the fifth step. Hand over. The table that would record a made item has nothing in it...	proof	SCREENSHOT	demo	Cut to the real thing: the produced-items table, open and empty, zero rows, and then the plan table filtered to the three source codes that appear twice with their two week dates side by side in the same column.	Zero rows. Three repeats.	hardcut	Db9MT-axxcO (the alternation between the artefact and the interpretation carries a procedure without losing the story)	template
00:57.6s	The process is not broken where the money went. It is broken at the step where a plan is...	payoff	A_R OLL_IUM	a_roll	Michael, medium shot, one location, banner plate held above him. The five-row stack returns as a small inset with the fifth row highlighted in Oxide and its figure column still blank.	Broken at the handover	hardcut	Db9MT-axxcO (the reel returns to the face to say what the artefact meant, in the last fifth)	template
01:06.6s	Take last week's list, write a name next to every line, and count the blanks. Send this t...	cta	A_R OLL_IUM	a_roll	Same take, no reframe. Michael holds still for the closing line.	A name on every line	none	DbbA7wwTIZ8 (the next action is bounded tightly enough to be done this week, and delivered plainly)	template

C-2026-09-12-10 · S00 · 0.0-7.6s · SCREENSHOT · demo · ->none

23 plans. Three weeks.

...

[SCREEN]

Full-frame, filmed off a real screen and scrolling: our own plan table, twenty-three rows across three weeks, every row filled in and the produced column empty on every one of them. No presenter.

This is what our Monday produces. Twenty-three of these in three weeks. Not one of them was ever made.

src: Db9MT-axxc0, DYuivWzSj5k · asset: template

C-2026-09-12-10_S00.png

C-2026-09-12-10 · S01 · 7.6-14.4s · NOTION_GRAPHIC · motion_graphic · ->hard_cut

Collect. Score. Shortlist. Strike. Hand over.

[TEXT CARD]

Collect. Score. Shortlist. Strike. Hand over.

Full-frame vertical stack of the five step names, revealed one row at a time, left-aligned on the grid. The first four rows are set in Graphite; the fifth stays empty for now.

Five steps. Collect, score, shortlist, strike out, hand over. Four of them run without anyone touching anything.

src: DZLDk9ySi-7 · asset: template

C-2026-09-12-10_S01.png

C-2026-09-12-10 · S02 · 14.4-27.2s · NOTION_GRAPHIC · motion_graphic · ->overlay_reveal

65 minutes. Nine struck.

[TEXT CARD]

65 minutes. Nine struck.

The same stack. The four automated rows are bracketed together against one figure - sixty-five minutes for the whole machine run - and the strike-out row carries nine struck of twenty-three. The fifth row's figure column is still blank.

The whole machine run took sixty-five minutes. Collecting, scoring, shortlisting, all of it. Striking out is the one step that waits for one of us, and it removed nine of the twenty-three.

src: DYuivWzSj5k, DZLDk9ySi-7 · asset: template

C-2026-09-12-10_S02.png

C-2026-09-12-10 · S03 · 27.2-44.4s · SCREEN_RECORDING · demo · ->hard_cut

Zero rows. Three repeats.

...

[SCREEN]

Cut to the real thing: the produced-items table, open and empty, zero rows, and then the plan table filtered to the three source codes that appear twice with their two week dates side by side in the same column.

Then the fifth step. Hand over. The table that would record a made item has nothing in it at all. And three of the twenty-three are the same source proposed again a week later, because nobody had marked the first one as used.

src: Db9MT-axxc0 · asset: template

C-2026-09-12-10_S03.png

C-2026-09-12-10 · S04 · 44.4-57.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Broken at the handover

[SHOT] medium

Michael, medium shot, one location, banner plate held above him. The five-row stack returns as a small inset with the fifth row highlighted in Oxide and its figure column still blank.

The process is not broken where the money went. It is broken at the step where a plan is supposed to become somebody's job, and that step has no name and no owner.

src: Db9MT-axxc0 · asset: template

C-2026-09-12-10_S04.png

C-2026-09-12-10 · S05 · 57.6-66.4s · A_ROLL_MEDIUM · a_roll · ->none

A name on every line

[SHOT] medium

Same take, no reframe. Michael holds still for the closing line.

Take last week's list, write a name next to every line, and count the blanks. Send this to whoever runs your Monday.

src: Db9MT-axxc0 · asset: template

C-2026-09-12-10_S05.png

C-2026-09-12-10 · S00 · 0.0-7.6s · SCREENSHOT · demo · ->none

23 plans. Three weeks.
...

[SCREEN]
Full-frame, filmed off a real screen and scrolling: our own plan table, twenty-three rows across three weeks, every row filled in and the produced column empty on every one of them. No presenter.

This is what our Monday produces. Twenty-three of these in three weeks. Not one of them was ever made.

src: Db9MT-axxc0, DYuivWzSj5k · asset: template

S00 0.0-7.6s

C-2026-09-12-10 · S01 · 7.6-14.4s · MOTION_GRAPHIC · motion_graphic · ->hard_cut

Collect. Score. Shortlist. Strike. Hand over.

[TEXT CARD]

Collect. Score. Shortlist. Strike. Hand over.

Full-frame vertical stack of the five step names, revealed one row at a time, left-aligned on the grid. The first four rows are set in Graphite; the fifth stays empty for now.

Five steps. Collect, score, shortlist, strike out, hand over. Four of them run without anyone touching anything.

src: DZLDk9ySi-7 · asset: template

S01 7.6-14.4s

C-2026-09-12-10 · S02 · 14.4-27.2s · MOTION_GRAPHIC · motion_graphic · ->reveal_text

65 minutes. Nine struck.

[TEXT CARD]

65 minutes. Nine struck.

The same stack. The four automated rows are bracketed together against one figure - sixty-five minutes for the whole machine run - and the strike-out row carries nine struck of twenty-three. The fifth row's figure column is still blank.

The whole machine run took sixty-five minutes. Collecting, scoring, shortlisting, all of it. Striking out is the one step that waits for one of us, and it removed nine of the twenty-three.

src: DYuivWzSj5k, DZLDk9ySi-7 · asset: template

S02 14.4-27.2s

C-2026-09-12-10 · S03 · 27.2-44.4s · SCREEN_RECORDING · demo · ->hard_cut

Zero rows. Three repeats.
...

[SCREEN]
Cut to the real thing: the produced-items table, open and empty, zero rows, and then the plan table filtered to the three source codes that appear twice with their two week dates side by side in the same column.

Then the fifth step. Hand over. The table that would record a made item has nothing in it at all. And three of the twenty-three are the same source proposed again a week later, because nobody had marked the first one as used.

src: Db9MT-axxc0 · asset: template

S03 27.2-44.4s

C-2026-09-12-10 · S04 · 44.4-57.6s · A_ROLL_MEDIUM · a_roll · ->hard_cut

Broken at the handover

[SHOT] medium
Michael, medium shot, one location, banner plate held above him. The five-row stack returns as a small inset with the fifth row highlighted in Oxide and its figure column still blank.



The process is not broken where the money went. It is broken at the step where a plan is supposed to become somebody's job, and that step has no name and no owner.

src: Db9MT-axxc0 · asset: template

S04 44.4-57.6s

C-2026-09-12-10 · S05 · 57.6-66.4s · A_ROLL_MEDIUM · a_roll · ->none

A name on every line

[SHOT] medium
Same take, no reframe. Michael holds still for the closing line.



Take last week's list, write a name next to every line, and count the blanks. Send this to whoever runs your Monday.

src: DbbA7weTIZ8 · asset: template

S05 57.6-66.4s

Contact sheet

Editing

- **Cut timing:** Six sections, three visible cuts. The two graphic scenes are one continuous build of the same stack; the presenter material at the end is one unbroken take. There is no speaker change, which is the only cut PRODUCTION.md allows.

- **Captions:** Burned-in speech captions, IBM Plex Sans SemiBold 40px, Paper on an Ink plate with 24px padding, bottom safe zone, two lines maximum. One persistent Ink banner plate in Saira SemiCondensed Black at the top for the whole reel. Step names and counts are set in IBM Plex Mono Medium as labels. No emoji on screen.
- **Emphasis:** Oxide appears exactly once, on the fifth row in the payoff scene. The four working steps are Graphite; nothing in the reel is coloured to look like a success.
- **Zooms:** None. The stack is read, not zoomed, and the empty table is held rather than pushed into.
- **Motion:** Linear easing only, per design/rules/motion.md. Rows reveal, figures do not count up. No bounce, no gradients, no rounded corners, no centred text.
- **Assets required:**
 - Filmed-screen scroll of our own plan table, twenty-three rows over three weeks with the produced column empty
 - Screen recording of the produced-items table showing zero rows, then the plan table filtered to the three repeated source codes
 - Five-row step stack rendered from design/tokens/tokens.json, with a figure column
 - Presenter A-roll, one continuous take, one location, banner plate fixed

Claims

CLAIM	STATE	SOURCE
Our weekly run produced 23 planning cards across three weeks: five on 3 September, nine on 7 September and nine on 10 September 2026.	OBSERVED	data/radar.db <code>cards</code> , grouped by week; reports/audit/04-existing-cards-review.md §A confirms 23 unique cards after deduplication across three databases.
None of the 23 has been shot or published: the table that would record a produced reel has zero rows.	OBSERVED	data/radar.db <code>our_posts</code> - 0 rows, and <code>our_metrics</code> - 0 rows; reports/audit/04-existing-cards-review.md - 'Zero of these cards have ever been shot or published'.
A person struck out 9 of the 23; 14 were developed into a full angle and hook.	OBSERVED	data/radar.db <code>cards</code> status counts: 9 struck out, 5 proposed, 9 draft.
Three source reels were proposed a second time in a later week, because the first proposal was never marked as used.	OBSERVED	data/radar.db <code>cards</code> - three codes appear twice across weeks; reports/audit/04-existing-cards-review.md classifies two of them as 'recycled, then cut without a new angle'.
The whole weekly machine run took sixty-five minutes on 10 September - collection, scoring, shortlisting and card generation together. No individual step has a recorded duration.	OBSERVED	reports/audit/03-server-runtime-hiker.md, run 2026-09-10 07:00-08:08: 'прогон занял 65 мин, потрачено сегодня \$2.60'. The log records a duration for the run as a whole and for no single step.
How long the strike-out step takes a person.	MISING	Never timed. The script says it is the one step that waits for a person and gives its outcome (9 of 23), not its duration; no minutes figure for it appears on screen.
That the same failure shape exists in the viewer's own weekly ritual.	MISING	We take apart one process, ours. The generalisation is offered as something to test on last week's list, never as a finding about anyone else.
That naming an owner on every line makes the handover work.	TO_MEASURE	We are adding the owner column to our own list this week and have no result yet; the reel says what we are changing, not what it fixed.
That the four automated steps are working correctly.	MISING	Not asserted. They run without a person, which is a different statement from running well; their output quality is audited elsewhere and is not this reel's subject.

Traceability

- **Insights:** I-17, I-09, I-08, I-15, I-13, I-29, I-05
- **Hypothesis:** H-19

- **Evidence summary:** Every number is our own and reproducible from data/radar.db: 23 planning cards across three weeks (5 / 9 / 9 on 3, 7 and 10 September 2026), a person struck out 9 of them, three source reels came back a second time in a later week, the whole machine run took 65 minutes, and the table that would record a produced reel has zero rows. That last figure is the reason to forward it, which is the Teardown format's second pass criterion. The corpus states a problem in only 64 of 264 reels and never takes apart a process the author runs (I-17), so a step-by-step of the author's own week with the failure point named is unrepresented. The references contribute shape only: DYuivWzSj5k (50.8 shares/1k, fourth in the 40-reel pool) the recognisable-weekday entry and the node-with-a-counter treatment, DZLDk9ySi-7 (41.9 shares, 45.4 saves/1k) the pain-as-acts list and the stacked-step layout, Db9MT-axxcO (47.9 shares and 68.9 saves per 1 000, third-best combined pair in the analysis-ready corpus, on a lift of only 1.62x) the beat clock and the artefact-in-hand framing, DbbA7wwTIZ8 the tightly bounded next action.
- **Original synthesis:** The process is ours, we may show it without anyone's permission - which RULES.md §2 names as the only Teardown source available from day one - and the finding is unflattering: the four steps we automated all work, and the one that does not is the step nobody built anything for. The five-step map with a real count against each step, ending on a count of zero, is ours and appears in no reference.
- **Why better than generic:** A generic process teardown would find the leak in the manual step and sell automation for it. This one automates four steps successfully and still produces nothing, which locates the failure in the handover instead - the part no tool is being sold for. It also refuses the two numbers it does not have: how long the human step takes, and whether naming an owner fixes anything.

Review

- **Version:** 1
- **Revised:** yes
- **Findings:**
 - [evidence] The script said 'Collecting took the machine sixty-five minutes'. reports/audit/03-server-runtime-hiker.md times the run 2026-09-10 07:00-08:08 as 65 minutes for all thirteen steps together and records no per-step duration; 'Scoring is instant' had no source at all → the line is now 'The whole machine run took sixty-five minutes. Collecting, scoring, shortlisting, all of it.' The claim and the S02 overlay, visual and editing note were changed with it, so the figure is never attached to the collect row alone. Script goes to version 2; 166 words, 66.4 s.
 - [evidence] Reference DYuivWzSj5k was called 'the second-highest share rate in the pool'. It is fourth: Dbi-yxNgrhG 58.4, Dbd7VDgP9PR 53.4 and DbO4zvJR1k8 52.9 per 1 000 are higher in the same 40-reel pool → corrected in the reference reason and in traceability.
 - [evidence] Reference Db9MT-axxcO was called 'best combined share and save rates in the corpus'. It is third, and the card's own source (06-creators-and-references §6, row 32) already names the two higher rows → corrected to 'third-best', with the two higher reels named.
 - [evidence] '1.62x its author's median' misreads the lift convention: view_lift is play/median – 1, so 1.62 lift is 2.62x the author's median → reworded to 'on a lift of only 1.62x' here and 'its 51.4x lift' for DZLDk9ySi-7. The point being made - that the reach multiplier is not the finding - is unchanged.
 - [evidence] The counting that carries the reel reproduces exactly from data/radar.db: 23 cards across three weeks (5 on 2026-09-03, 9 on 09-07, 9 on 09-10), 9 struck out and 14 developed,

our_posts and our_metrics both empty, and three source codes proposed twice (DcvBtiNtbYO, Dco1mOJzZ4L, DcwCCgivcIQ) → left.

- [evidence] The writer's decision to use 23 rather than the hypothesis' 33 is confirmed: reports/audit/04-existing-cards-review.md reaches 33 only by adding a separate ten-card lane to our 23 → left.
- [evidence] All four reference metric sets re-derived from data/radar.db and all four function tags match data/analysis/hypothesis_refs.json; the four quotes are verbatim from the beat ids given; the script shares no five-word sequence with any of the four transcripts → left.
- [positioning] The Teardown pass criteria are met: a named process with a named owner in shot ('the one step that waits for one of us') and a reason to forward (the count of zero). Our own process is the only Teardown source available today per RULES.md §2 → left.
- [hook] Opens on the artefact rather than on a recognisable morning, reversing the hook reference's direction, because openings spent on the audience carry median robust_z 0.67 against 1.36, p=0.002 (I-06). Three candidates, one question, first one used, 19 words / 7.6 s → left.
- [positioning] No split screen, no comment gate, no stock footage, no duration invented for the human step (recorded MISSING), no claim that naming an owner fixes anything (recorded TO_MEASURE) → left.
- [timing] Recomputed after the edit: 166 words, 66.4 s at 2.5 wps, storyboard contiguous 0.0-66.4 s, no scene under 1.5 s, every overlay ≤ 7 words → passes.
- [scheduling] Runs on the same Monday as C-2026-09-12-06, which describes the same pipeline as twenty-five named stages while this card describes it as five steps. Both are honest at their own level of description, but they should not be published in the same week.
- [schema] python3 -m engine.cards_v2 validate → 0 errors, 0 warnings after the edits.

Production state

- **Status:** REVIEWED
- **EDL path:** cards/edl/C-2026-09-12-10.json
- **Remotion validation:** {"card_id": "C-2026-09-12-10", "checked_at": "2026-09-12T08:53:28Z", "contract": "m2.remotion-edl.v1", "duration_frames": 1992, "edl": "cards/edl/C-2026-09-12-10.json", "fps": 30, "render_mode": "PREVIS", "result": {"ok": true, "returncode": 0, "skipped": false, "stderr": "", "stdout": "EDL_OK"}}